



Investor Presentation

October 2023

Except for the historical information contained herein, certain matters in this presentation including, but not limited to, statements as to: our financial position; our markets, market opportunity, demand and growth drivers; the benefits, impact, performance, features and availability of our products and technologies; the benefits, impact, features and timing of our collaborations or partnerships; NVIDIA accelerated computing being broadly recognized as the way to advance computing as Moore's law ends; data centers making a platform shift from general purpose to accelerated computing; trillion dollars of installed global data center infrastructure transitioning to accelerated computing; AI driving a platform shift in computing and enabling new, never-before-possible applications; broader enterprises driving the next wave of computing, followed by autonomous machines and industrial digitalization; accelerated computing being needed to tackle the most impactful opportunities of our time; NVIDIA's value to every stakeholder in the ecosystem; the ROI of high compute performance; enterprise as the next big generative AI opportunity; NVIDIA's expanding accelerated computing ecosystem; AI as the greatest technology force of our time; data centers becoming AI factories; generative AI unlocking new opportunities; the next wave of AI being robotics and industrial digitalization; NVIDIA's acceleration stacks and ecosystems helping to bring AI to the world's largest industries; NVIDIA's AI expertise and scale helping to revolutionize businesses; generative AI being the most important computing platform of our generation; full-stack and data center scale acceleration driving significant cost savings and workload scaling; our dividend program plan; and our Automotive design win pipeline and ramp expectations are forward-looking statements.

These forward-looking statements and any other forward-looking statements that go beyond historical facts that are made in this presentation are subject to risks and uncertainties that may cause actual results to differ materially. Important factors that could cause actual results to differ materially include: global economic conditions; our reliance on third parties to manufacture, assemble, package and test our products; the impact of technological development and competition; development of new products and technologies or enhancements to our existing product and technologies; market acceptance of our products or our partners' products; design, manufacturing or software defects; changes in consumer preferences and demands; changes in industry standards and interfaces; unexpected loss of performance of our products or technologies when integrated into systems and other factors.

NVIDIA has based these forward-looking statements largely on its current expectations and projections about future events and trends that it believes may affect its financial condition, results of operations, business strategy, short-term and long-term business operations and objectives, and financial needs. These forward-looking statements are subject to a number of risks and uncertainties, and you should not rely upon the forward-looking statements as predictions of future events. The future events and trends discussed in this presentation may not occur and actual results could differ materially and adversely from those anticipated or implied in the forward-looking statements. Although NVIDIA believes that the expectations reflected in the forward-looking statements are reasonable, the company cannot guarantee that future results, levels of activity, performance, achievements or events and circumstances reflected in the forward-looking statements will occur. Except as required by law, NVIDIA disclaims any obligation to update these forward-looking statements to reflect future events or circumstances. For a complete discussion of factors that could materially affect our financial results and operations, please refer to the reports we file from time to time with the SEC, including our most recent Annual Report on Form 10-K, Quarterly Reports on Form 10-Q, and Current Reports on Form 8-K. Copies of reports we file with the SEC are posted on our website and are available from NVIDIA without charge.

Many of the products and features described herein remain in various stages and will be offered on a when-and-if-available basis. The statements within are not intended to be, and should not be interpreted as a commitment, promise, or legal obligation, and the development, release, and timing of any features or functionalities described for our products is subject to change and remains at the sole discretion of NVIDIA. NVIDIA will have no liability for failure to deliver or delay in the delivery of any of the products, features or functions set forth herein.

NVIDIA uses certain non-GAAP measures in this presentation including non-GAAP gross profit, non-GAAP gross margin, non-GAAP operating income, non-GAAP operating margin, and free cash flow. NVIDIA believes the presentation of its non-GAAP financial measures enhances investors' overall understanding of the company's historical financial performance. The presentation of the company's non-GAAP financial measures is not meant to be considered in isolation or as a substitute for the company's financial results prepared in accordance with GAAP, and the company's non-GAAP measures may be different from non-GAAP measures used by other companies. Further information relevant to the interpretation of non-GAAP financial measures, and reconciliations of these non-GAAP financial measures to the most comparable GAAP measures, may be found in the slide titled "Reconciliation of Non-GAAP to GAAP Financial Measures".

Headquarters: Santa Clara, CA

NVIDIA pioneered accelerated computing to help solve impactful challenges classical computers cannot. A quarter of a century in the making, NVIDIA accelerated computing is broadly recognized as the way to advance computing as Moore's law ends and AI lifts off.

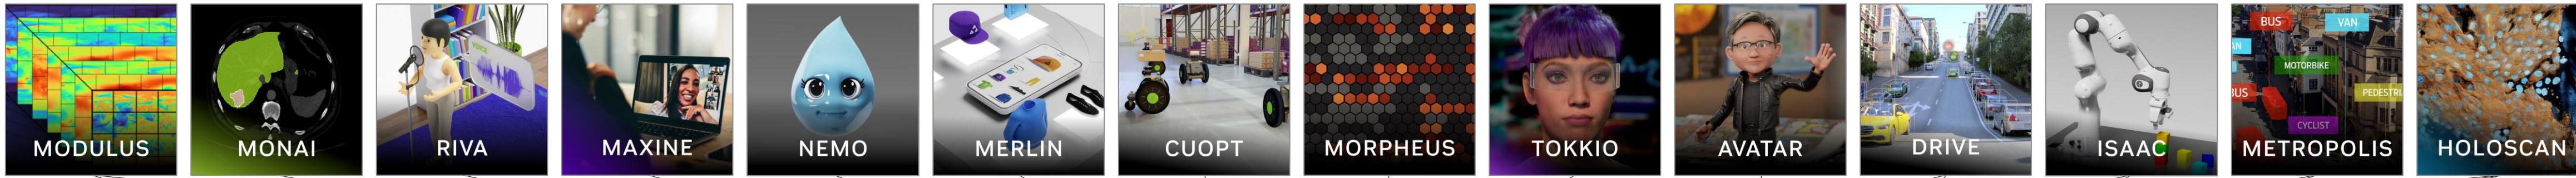
NVIDIA's platform is installed in several hundred million computers, is available in every cloud and from every server maker, powers 74% of the TOP500 supercomputers, and boasts over 4 million developers.



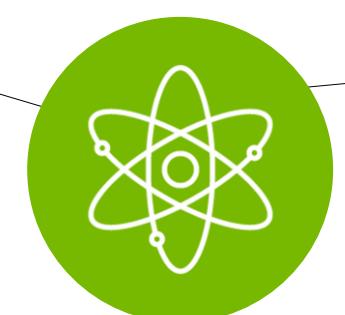
NVIDIA's Accelerated Computing Platform

Full-stack innovation across silicon, systems and software

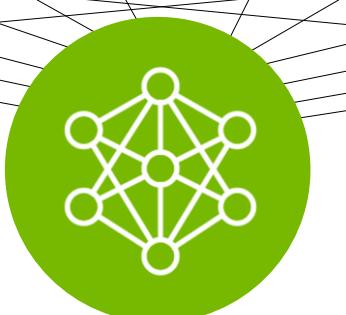
AI APPLICATION FRAMEWORK



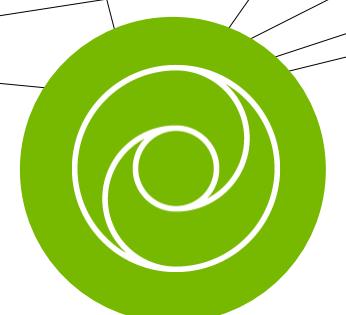
PLATFORMS



NVIDIA HPC



NVIDIA AI



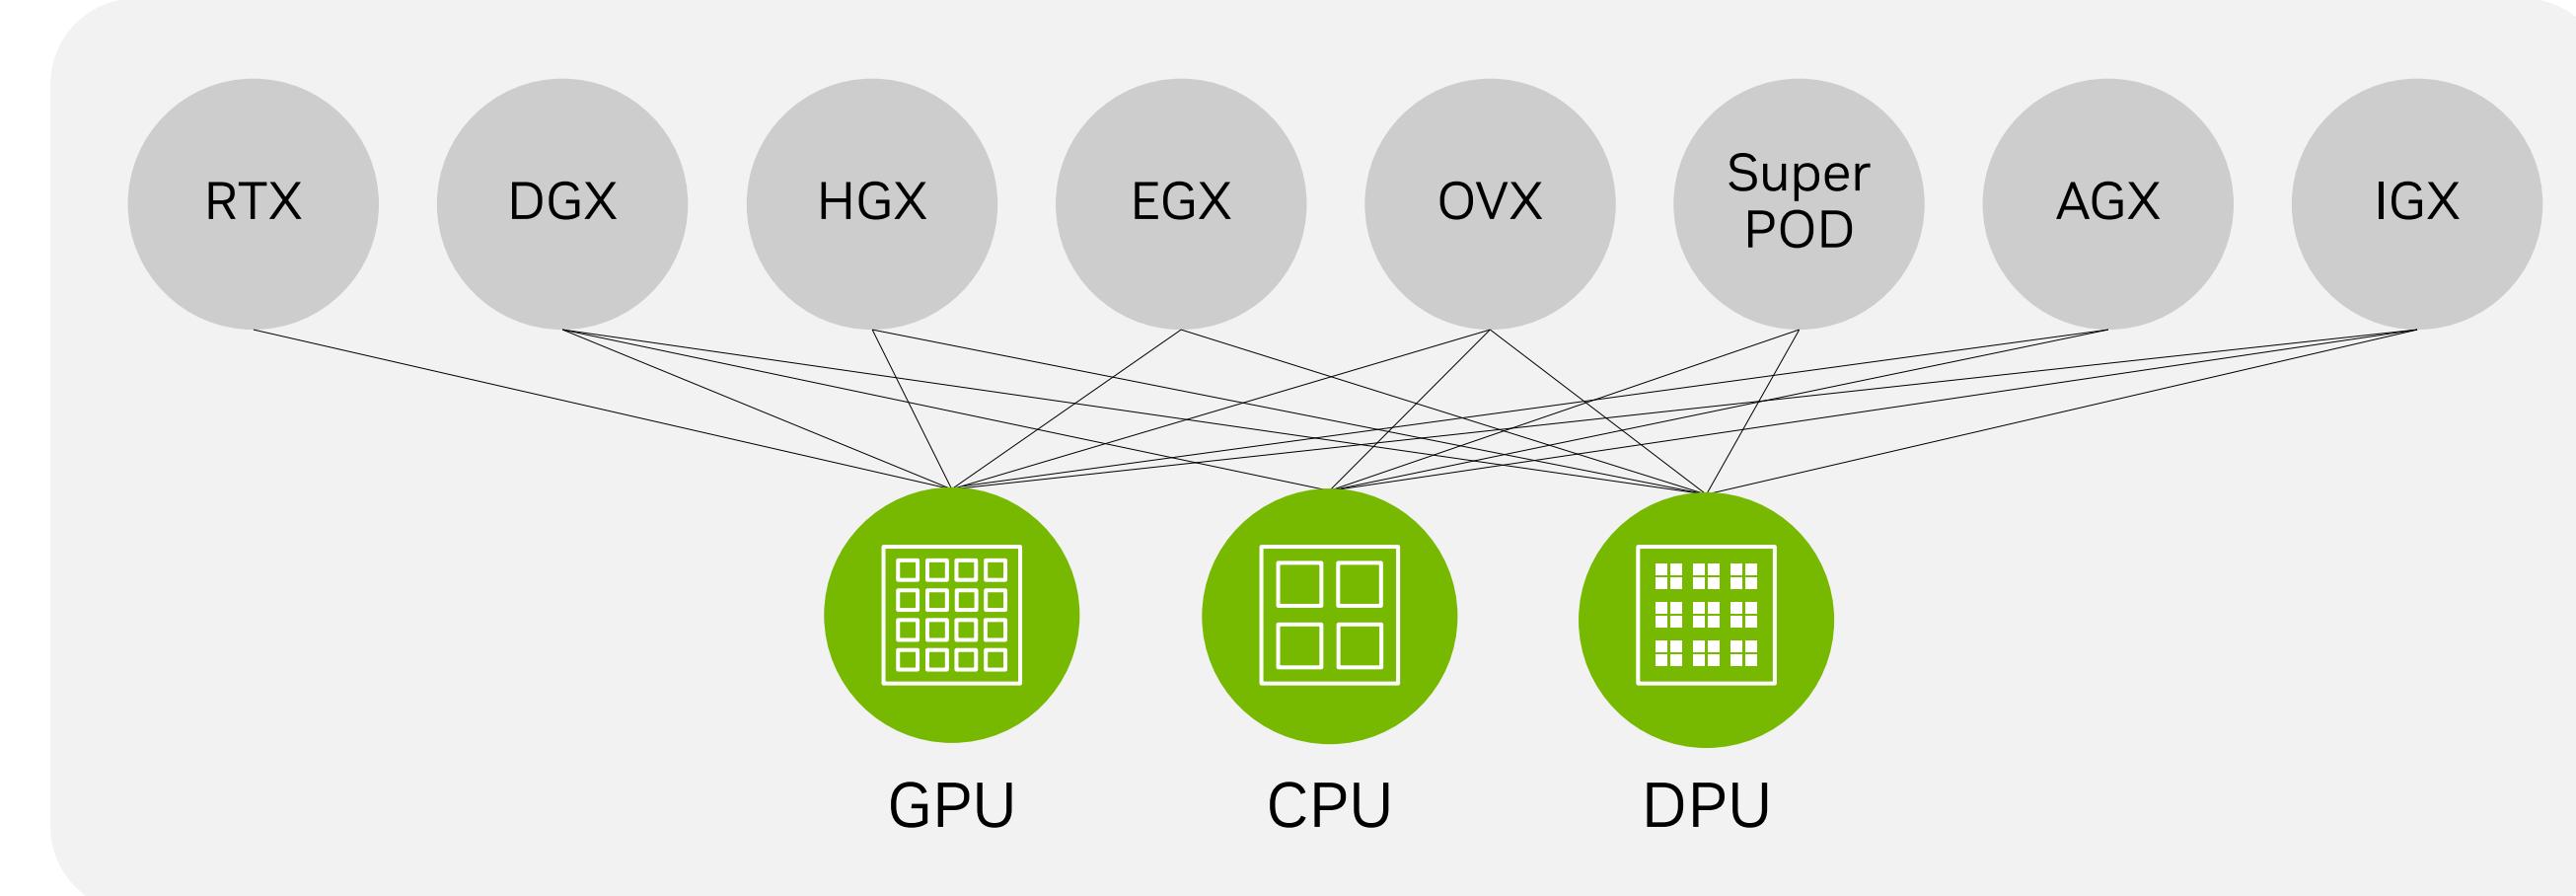
NVIDIA Omniverse

ACCELERATION LIBRARIES



CLOUD-TO-EDGE DATACENTER-TO- ROBOTIC SYSTEMS

3-CHIPS



With nearly three decades of a singular focus, NVIDIA is expert at accelerating software and scaling compute by a **Million-X**, going well beyond Moore's law.

Accelerated computing requires **full-stack** innovation – optimizing across every layer of computing – from silicon and systems to software and algorithms, demanding deep understanding of the problem domain.

Our full-stack platforms – NVIDIA HPC, NVIDIA AI, and NVIDIA Omniverse – accelerate high performance computing, AI, and industrial digitalization workloads.

We accelerate workloads at **data center scale**, across thousands of compute nodes, treating the network and storage as part of the computing fabric.

Our platform extends from the cloud and enterprise data centers to supercomputing centers, edge computing and PCs.

Why Accelerated Computing?

Advancing computing in the post-Moore's Law era

Accelerated computing is needed to tackle the most impactful opportunities of our time—like AI, climate simulation, drug discovery, ray tracing, and robotics.

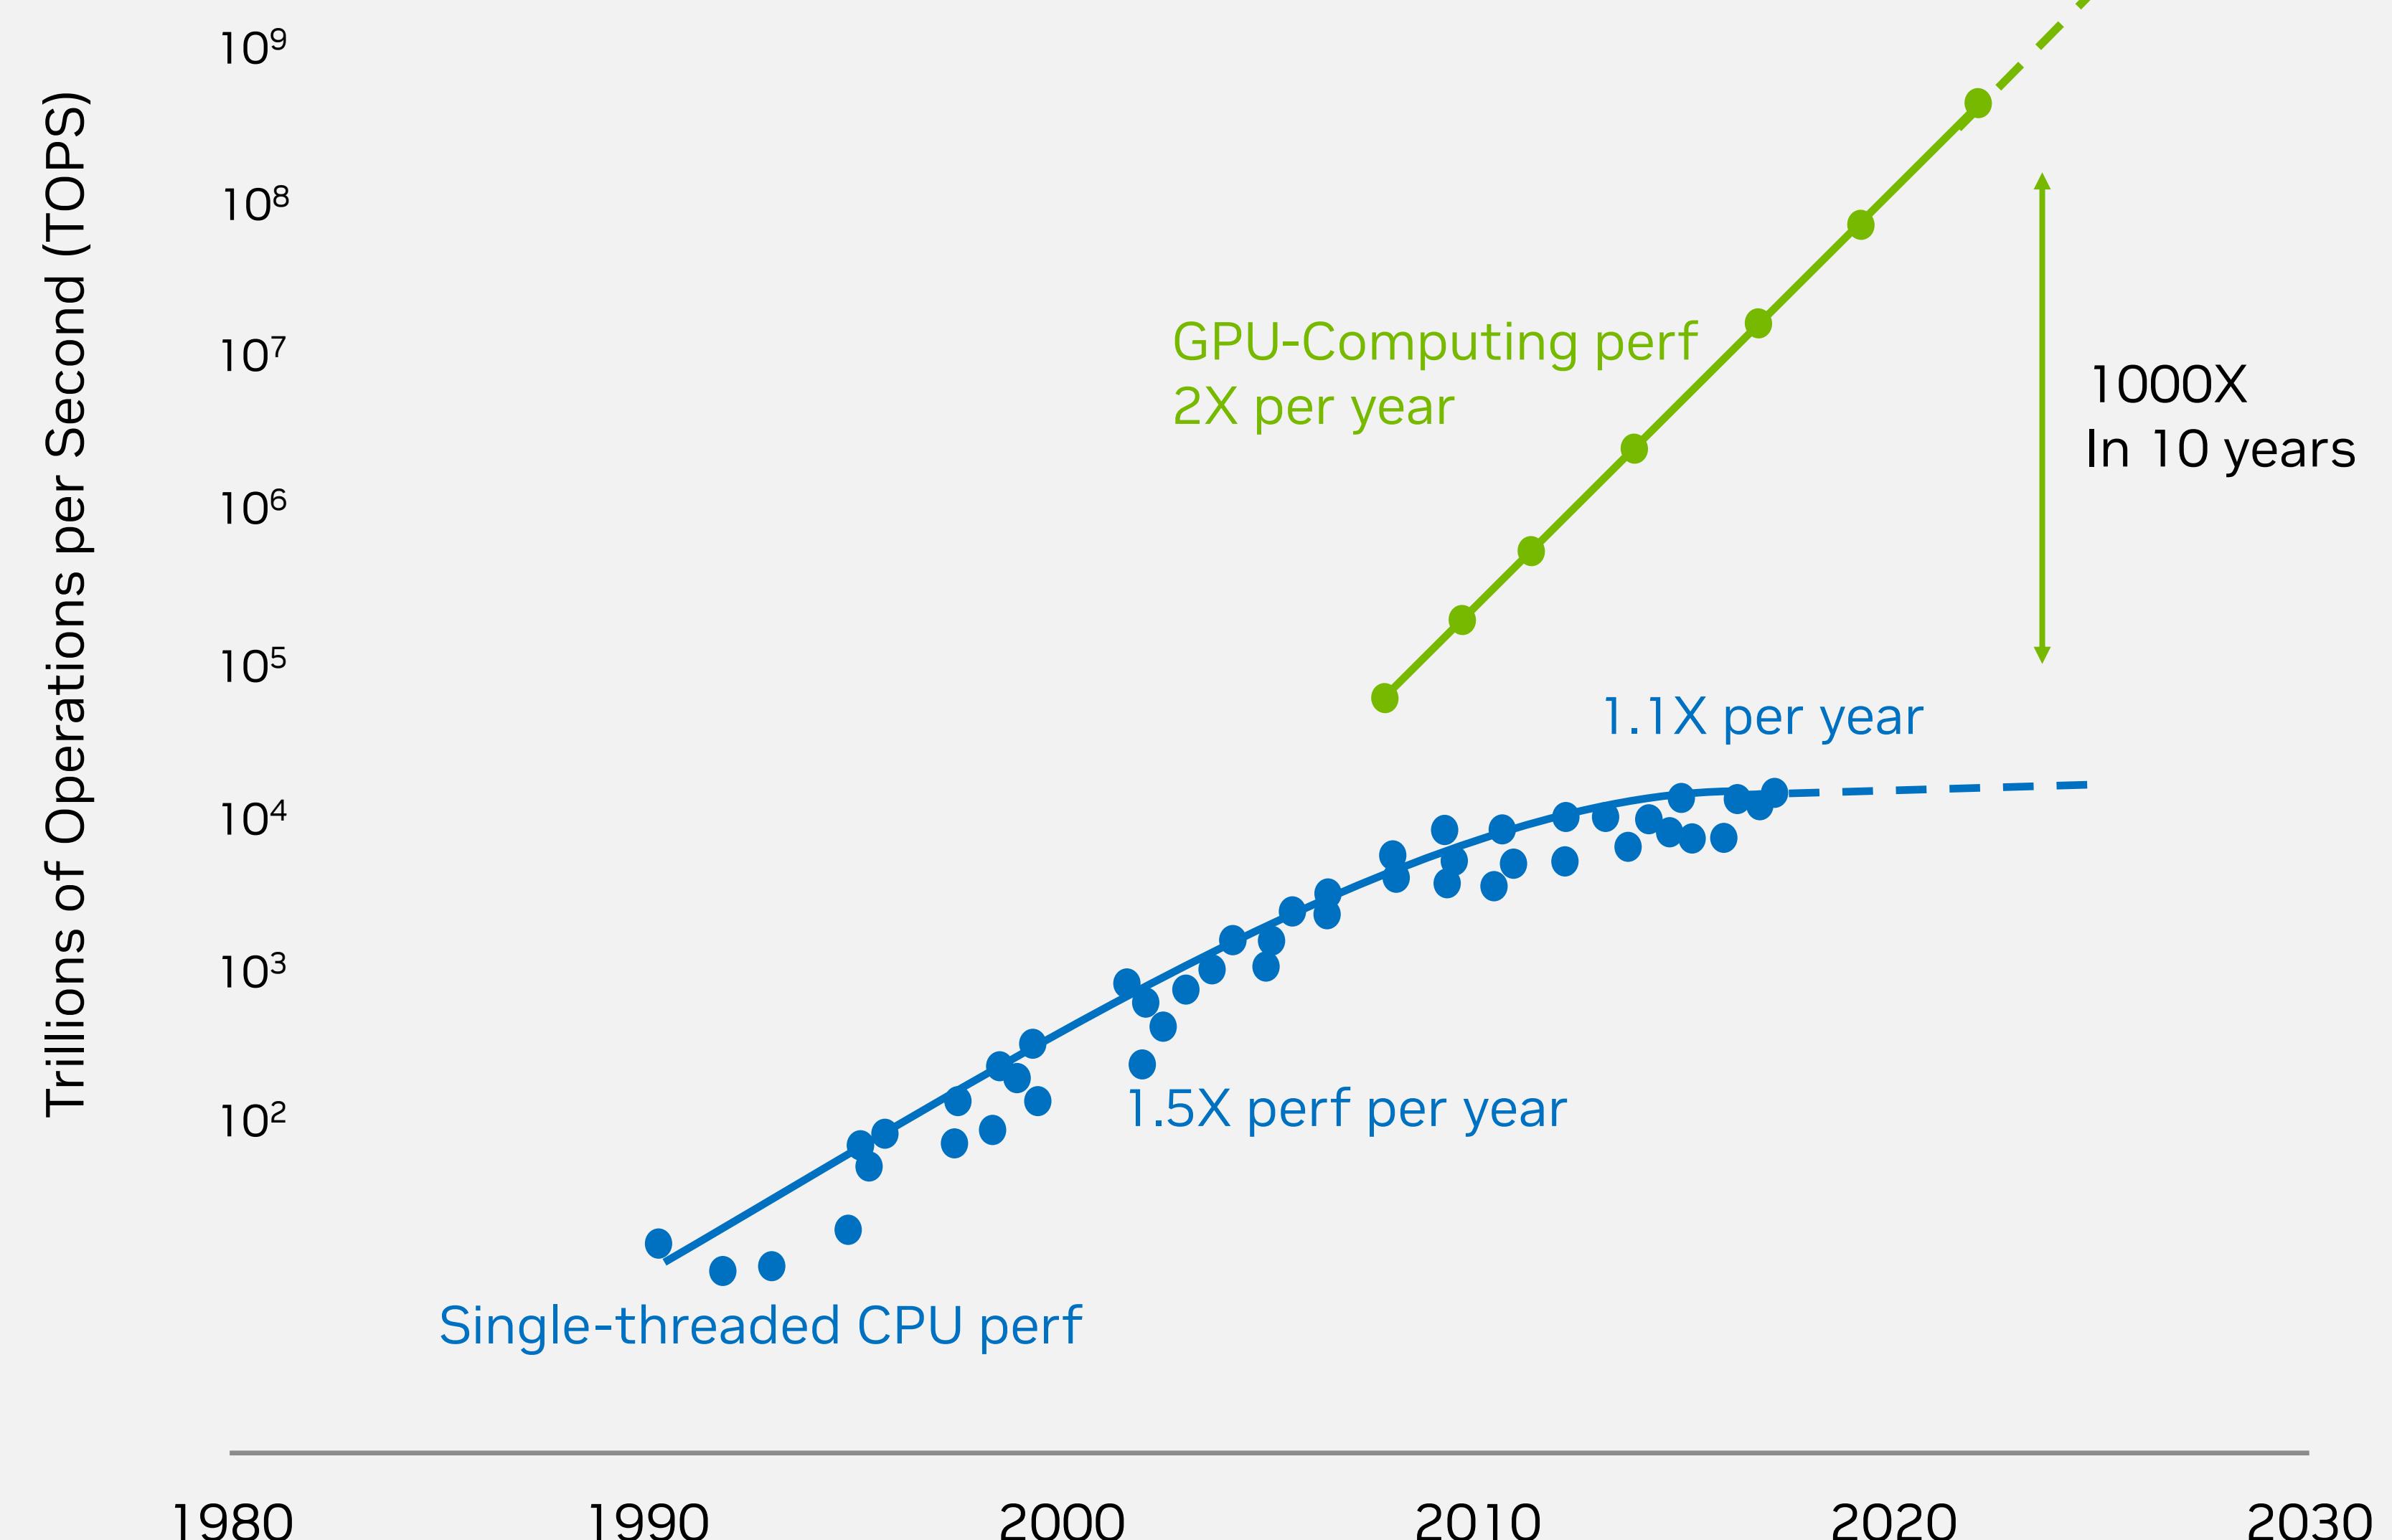
NVIDIA is uniquely dedicated to accelerated computing—working top-to-bottom—refactoring applications and creating new algorithms, and bottom-to-top—inventing new specialized processors, like RT Core and Tensor Core.

“It’s the end of Moore’s Law as we know it.”

- John Hennessy, Oct 2018

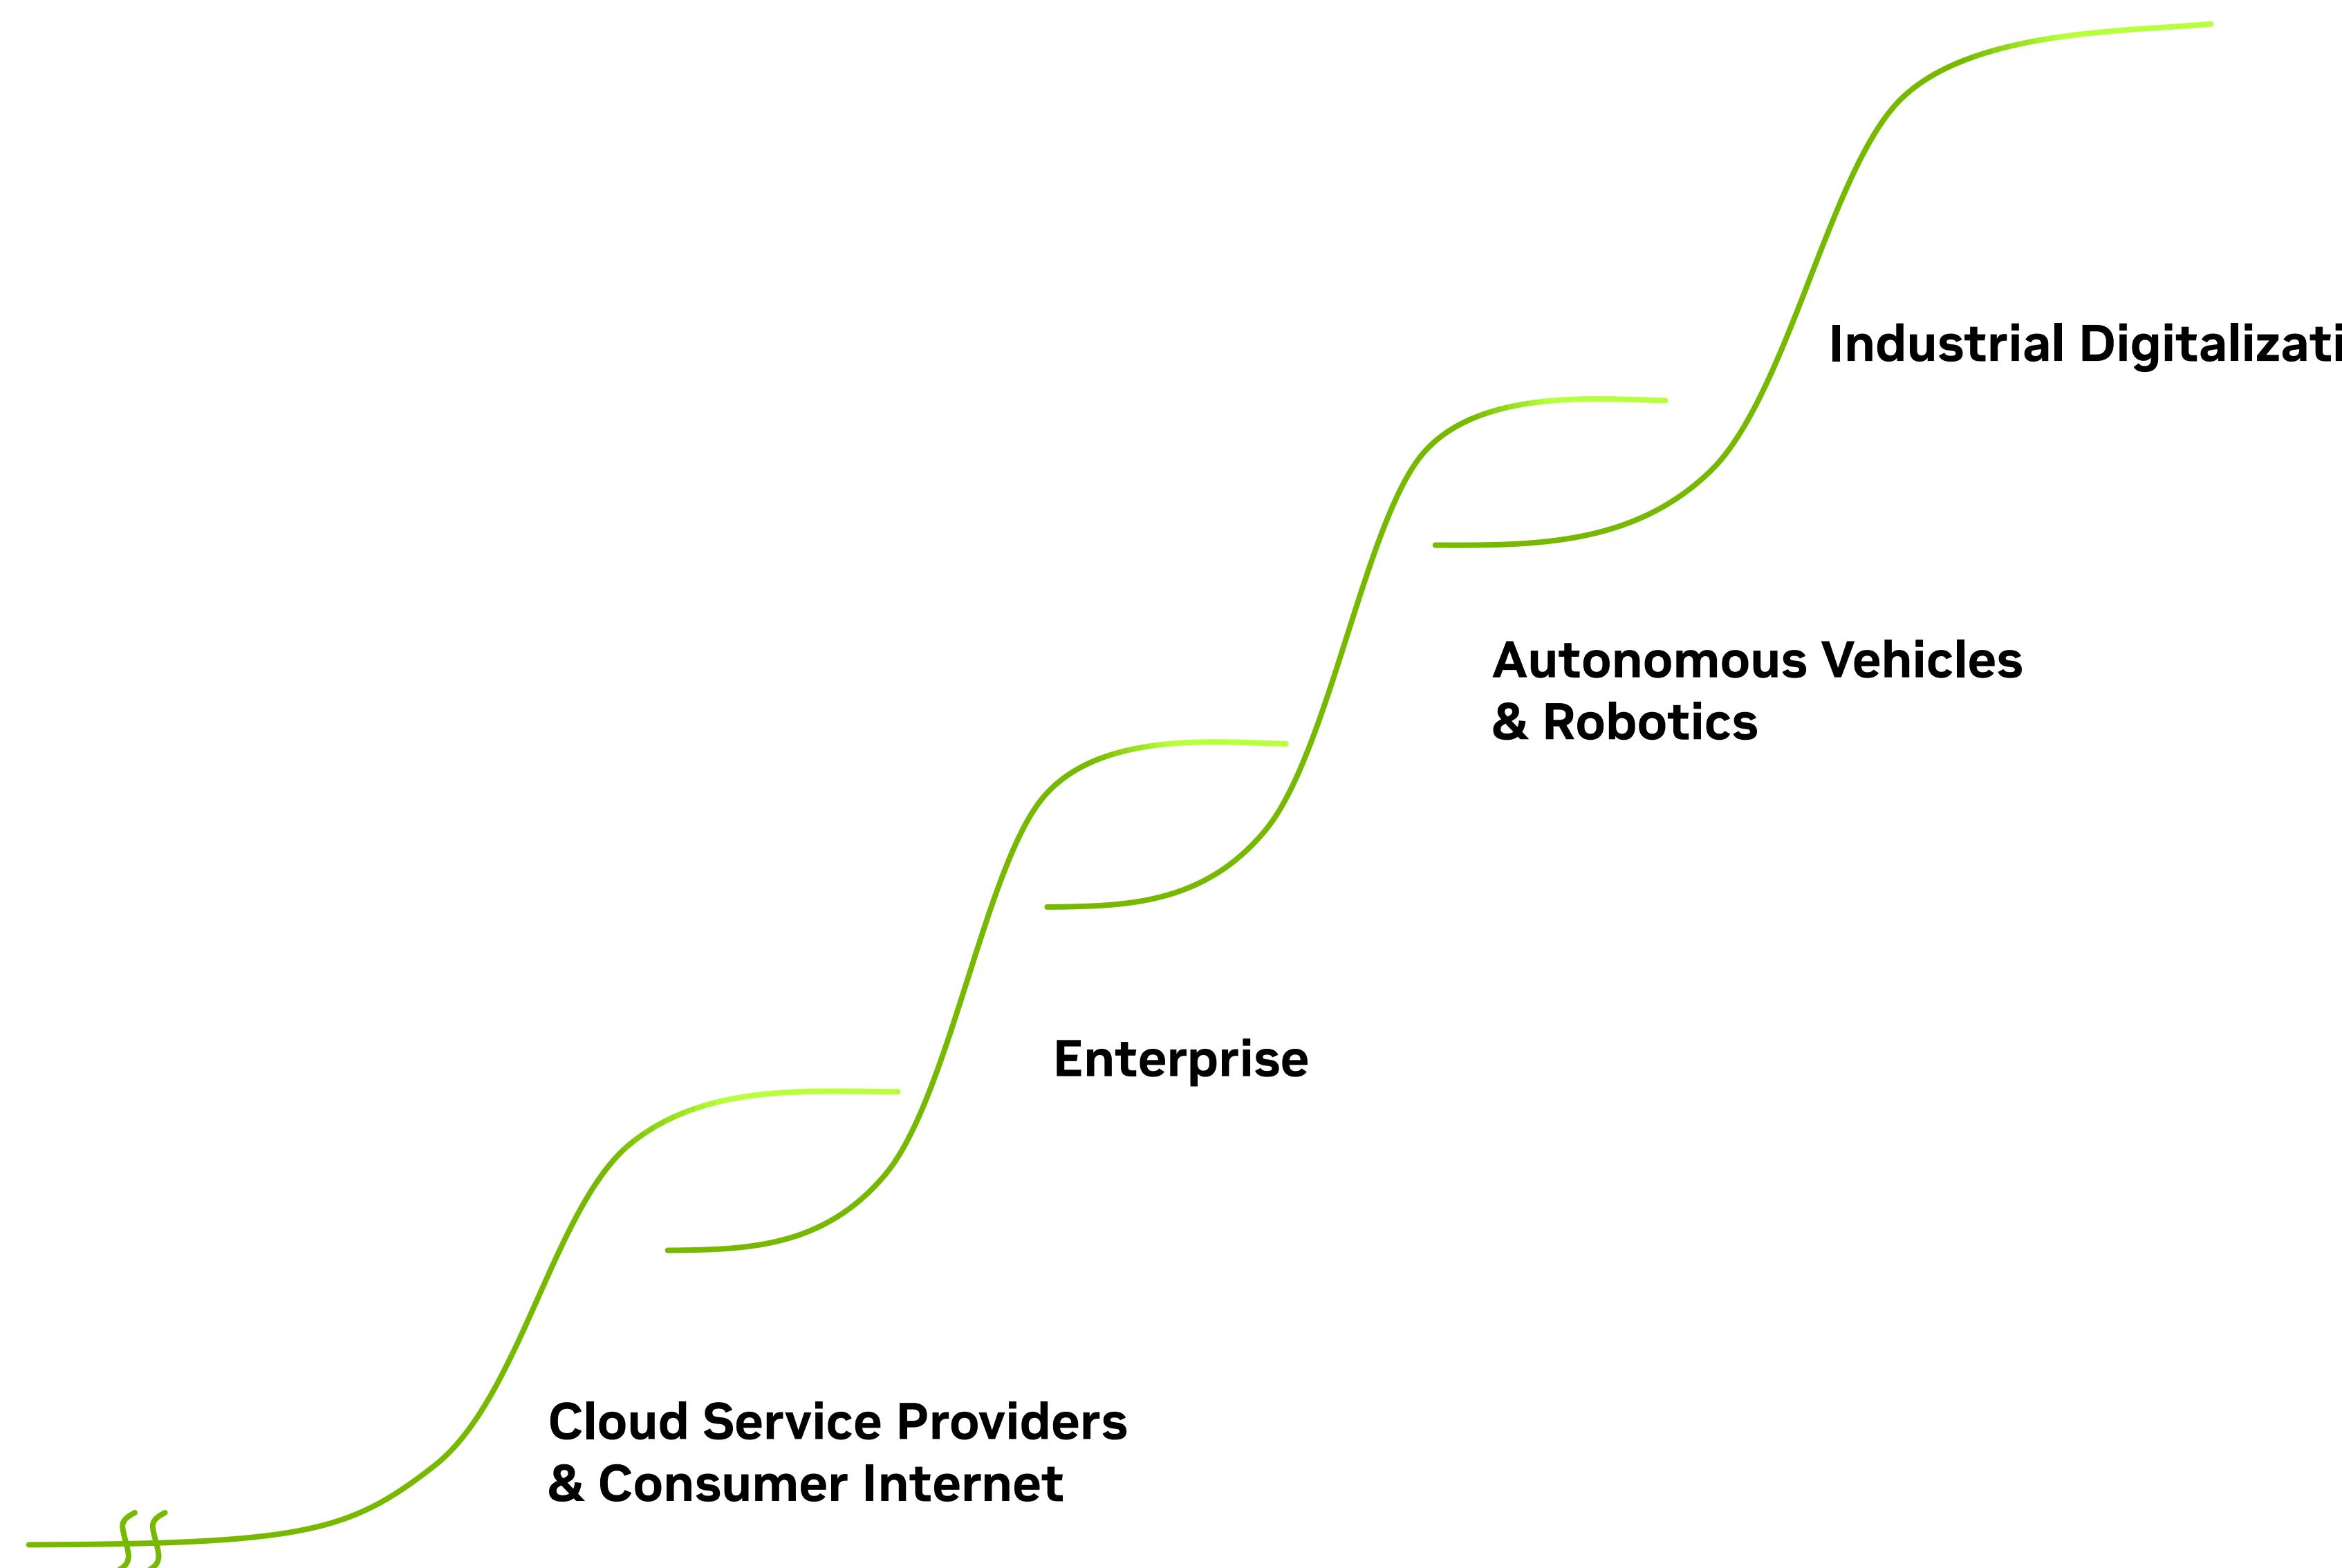
“Moore’s Law is dead.”

- Jensen Huang, GTC 2013



Waves of Adoption of Accelerated Computing

A generational computing platform shift



A new computing era has begun.

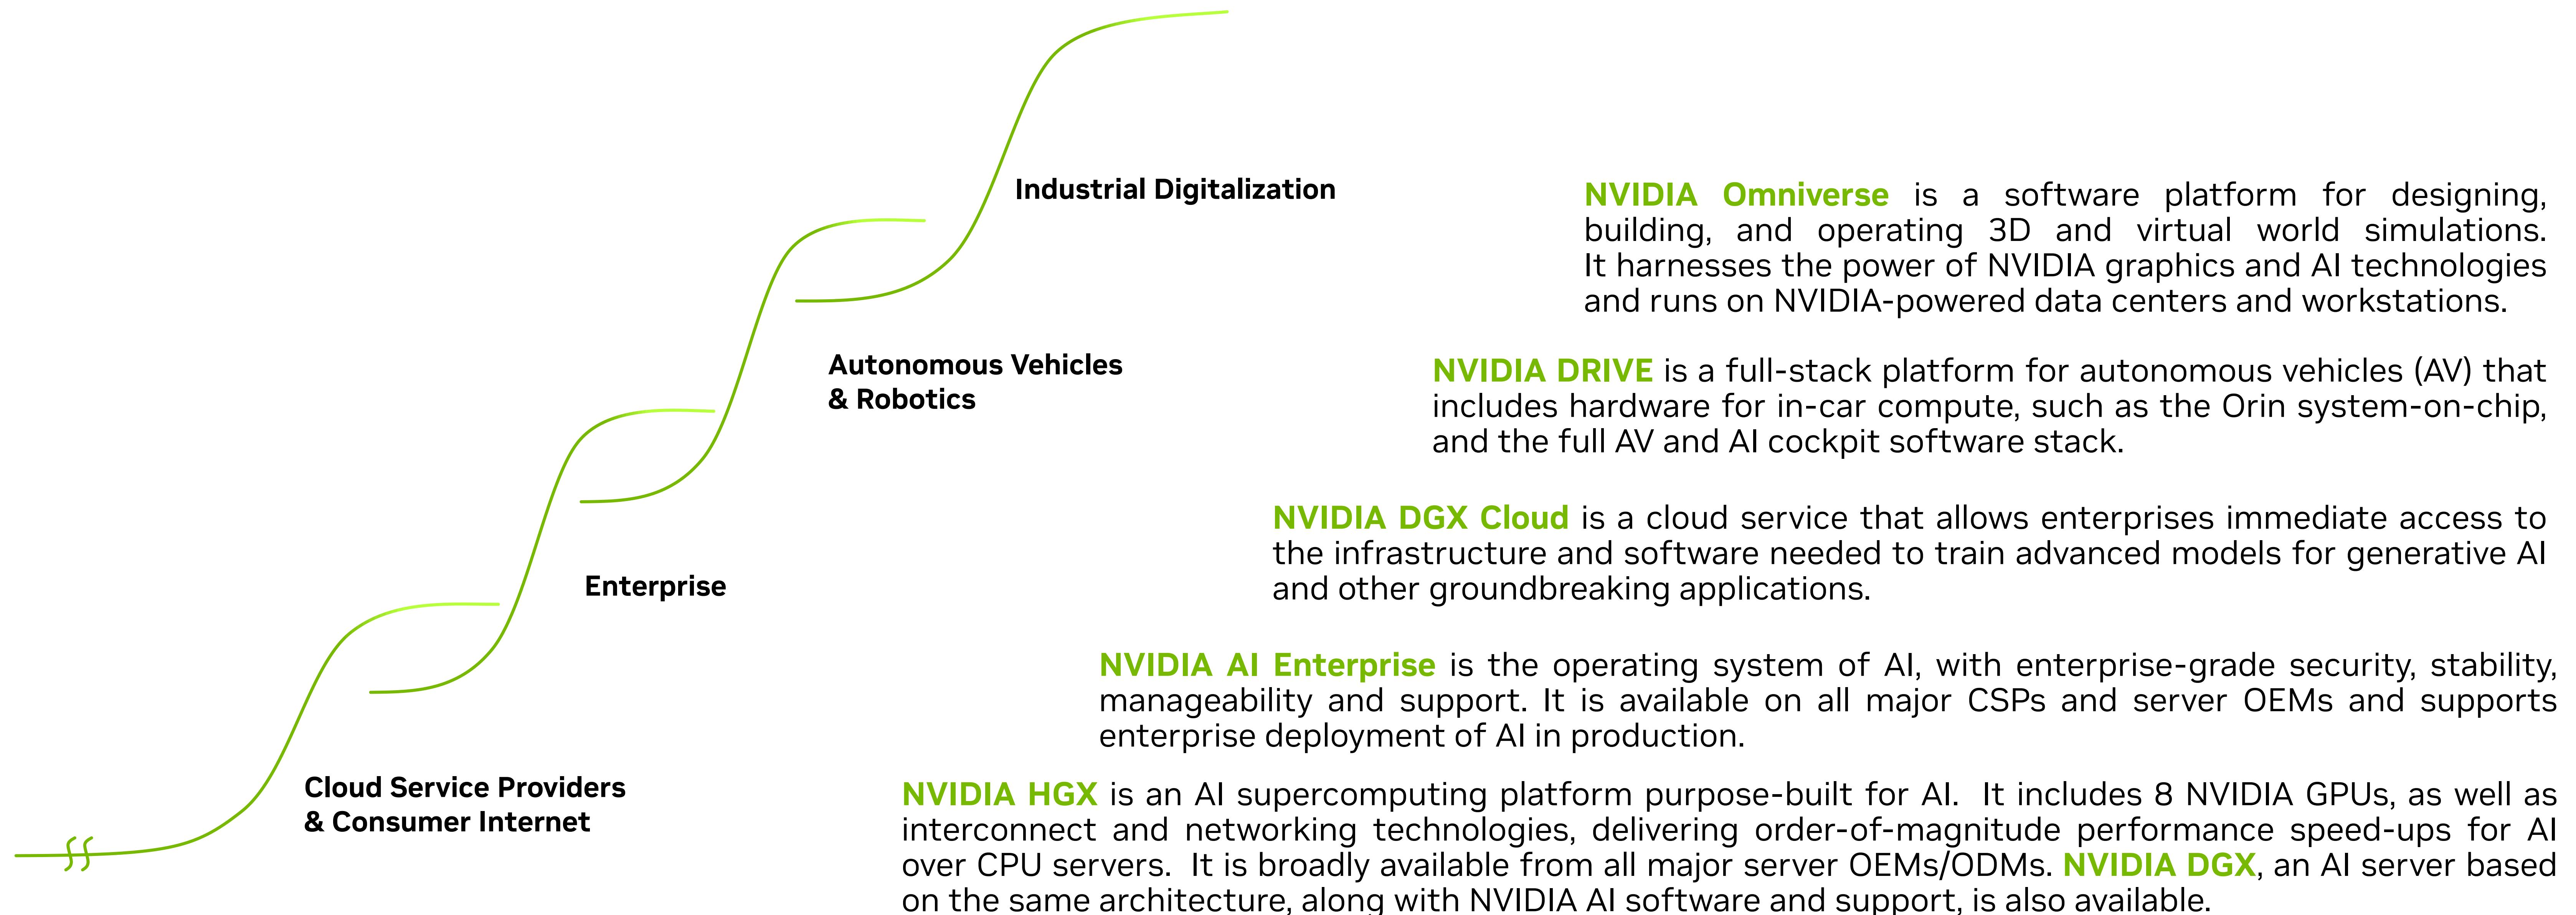
Accelerated computing enabled the rise of AI, which is driving a platform shift from general purpose to accelerated computing, and enabling new, never-before-possible applications.

The trillion dollars of installed global data center infrastructure will transition to accelerated computing to achieve an order of magnitude better performance, energy-efficiency and cost.

Hyperscale cloud service providers and consumer internet companies have been the early adopters of AI and accelerated computing, with broader enterprise adoption now under way.

AI and accelerated computing will also make possible the next big waves – autonomous machines and industrial digitalization.

NVIDIA Accelerated Computing for Every Wave



NVIDIA's Accelerated Computing Ecosystem

The NVIDIA accelerated computing platform has attracted the largest ecosystem of developers, supporting a rapidly growing universe of applications and industry innovation.

Developers can engage with NVIDIA through **CUDA** – our parallel computing programming model introduced in 2006 – or at higher layers of the stack, including libraries, pre-trained AI models, SDKs and other development tools.

300 Libraries

400 AI Models

100 Updated in the Last Year

Developers

1.8M

2020

4M

2023

CUDA Downloads*

20M

2020

45M

2023

*Cumulative

AI Startups

6K

2020

15K

2023

GPU-Accelerated Applications

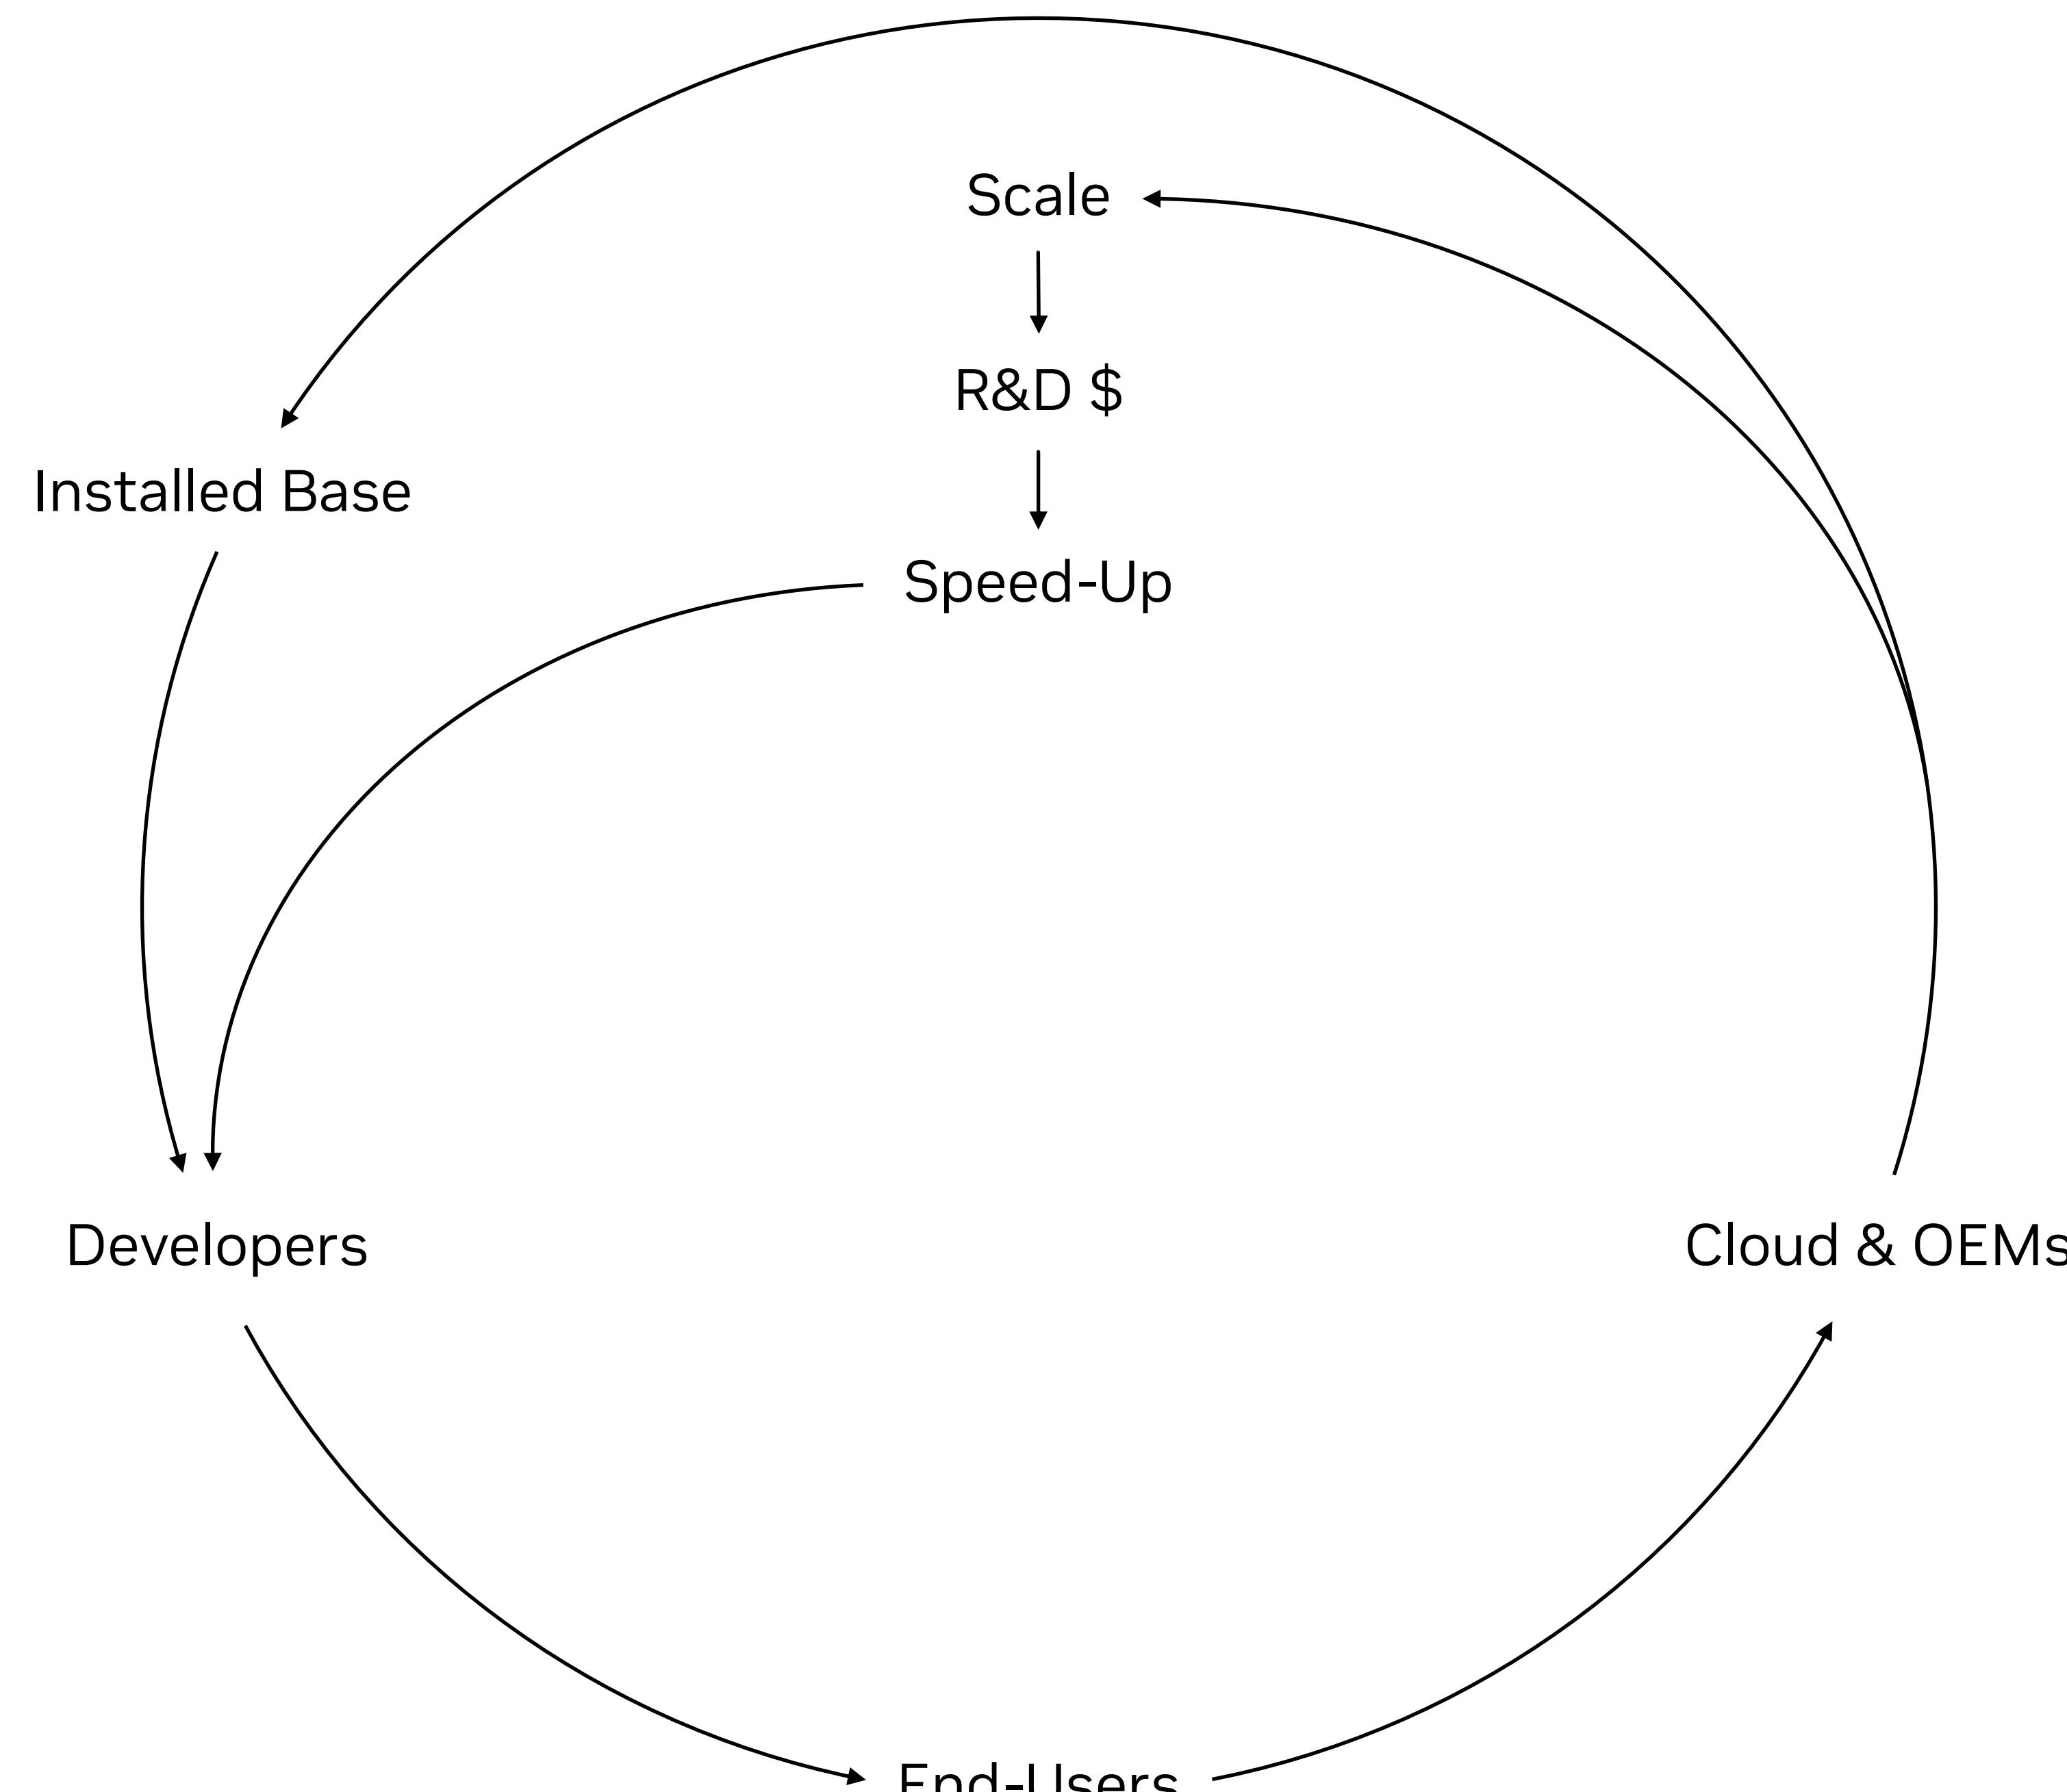
700

2020

3,200

2023

NVIDIA's Multi-Sided Platform and Flywheel



NVIDIA Accelerated Computing
Virtuous Cycle

The virtuous cycle of NVIDIA's accelerated computing starts with an installed base of several hundred million GPUs, all compatible with the CUDA programming model.

- **For developers** – NVIDIA's one architecture and large installed base give developer's software the best performance and greatest reach
- **For end users** – NVIDIA is offered by virtually every computing provider and accelerates the most impactful applications from cloud to edge
- **For cloud providers and OEMs** – NVIDIA's rich suite of Acceleration Platforms lets partners build one offering to address large markets including media & entertainment, healthcare, transportation, energy, financial services, manufacturing, retail, and more
- **For NVIDIA** – Deep engagement with developers, computing providers, and customers in diverse industries enables unmatched expertise, scale, and speed of innovation across the entire accelerated computing stack – propelling the flywheel

Huge ROI from AI Driving a Powerful New Investment Cycle

AI can augment creativity and productivity by orders of magnitude across industries

Knowledge workers will use copilots based on large language models to generate documents, answer questions, or summarize missed meetings, emails and chats – adding hours of productivity per week.

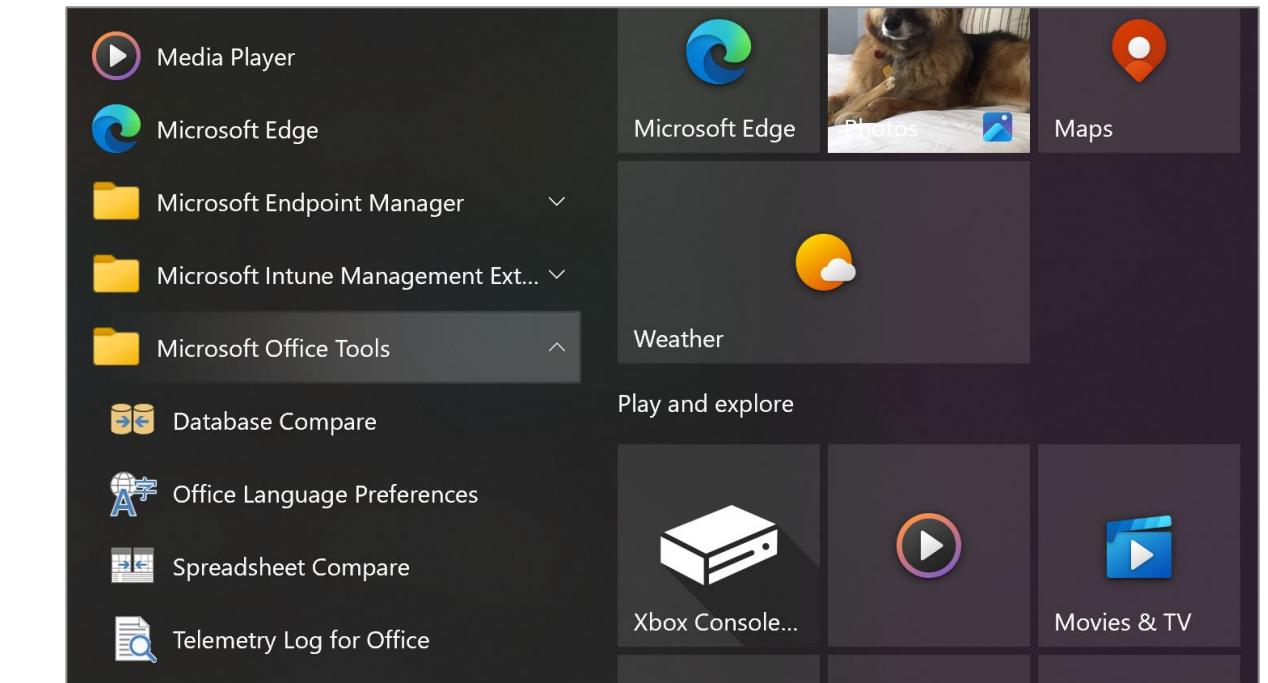
Copilots specialized for fields such as software development, legal services or education can boost productivity by as much as 50%.

Social media, search and e-commerce apps are using deep recommenders to offer more relevant content and ads to their customers, increasing engagement and monetization.

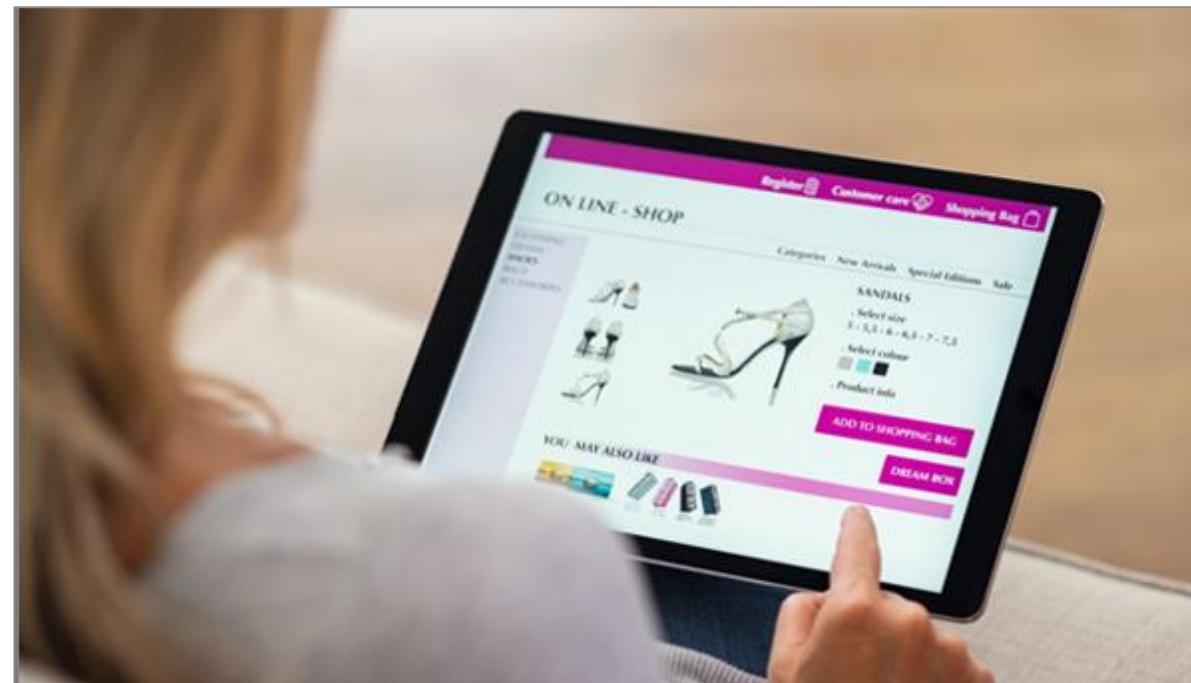
Creators can generate stunning, photorealistic images with a single text prompt – compressing workflows that take days or weeks into minutes in industries from advertising to game development.

Call center agents augmented with AI chatbots can dramatically increase productivity and customer satisfaction.

Drug discovery, financial services, agriculture and food services and climate forecasting are seeing order-of-magnitude workflow acceleration from AI.



Office AI Copilots
Over 1B knowledge workers



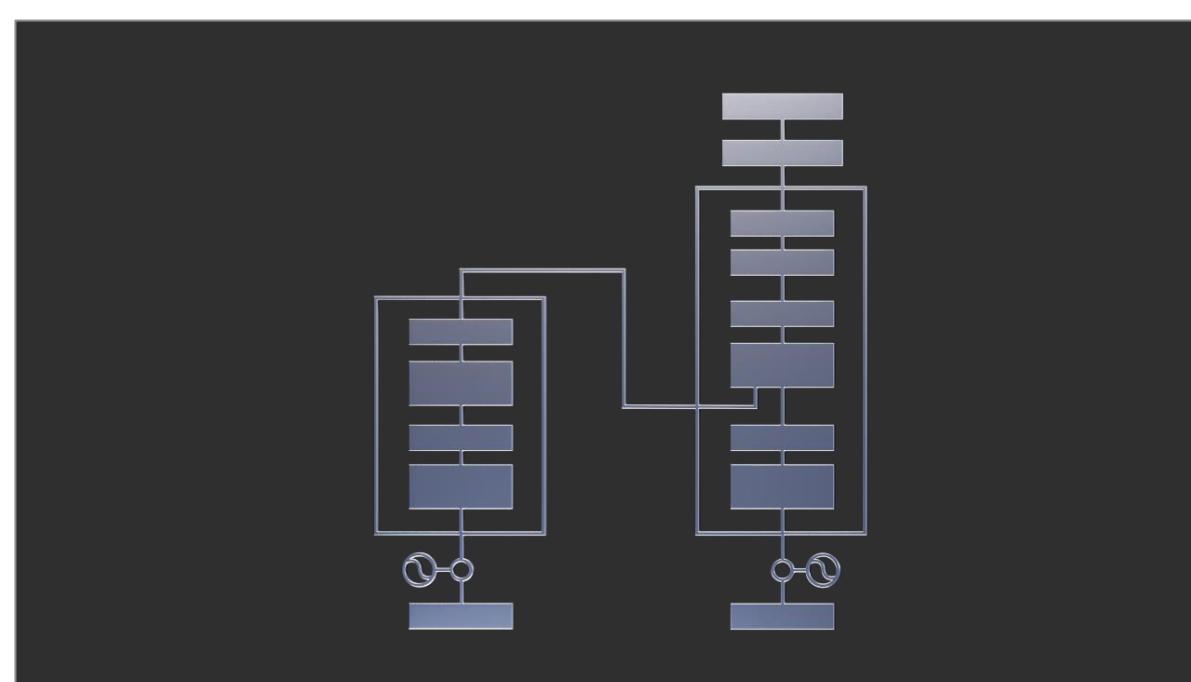
Search & Social Media
\$700B in digital advertising annually



AI Content Creation
50M creators globally



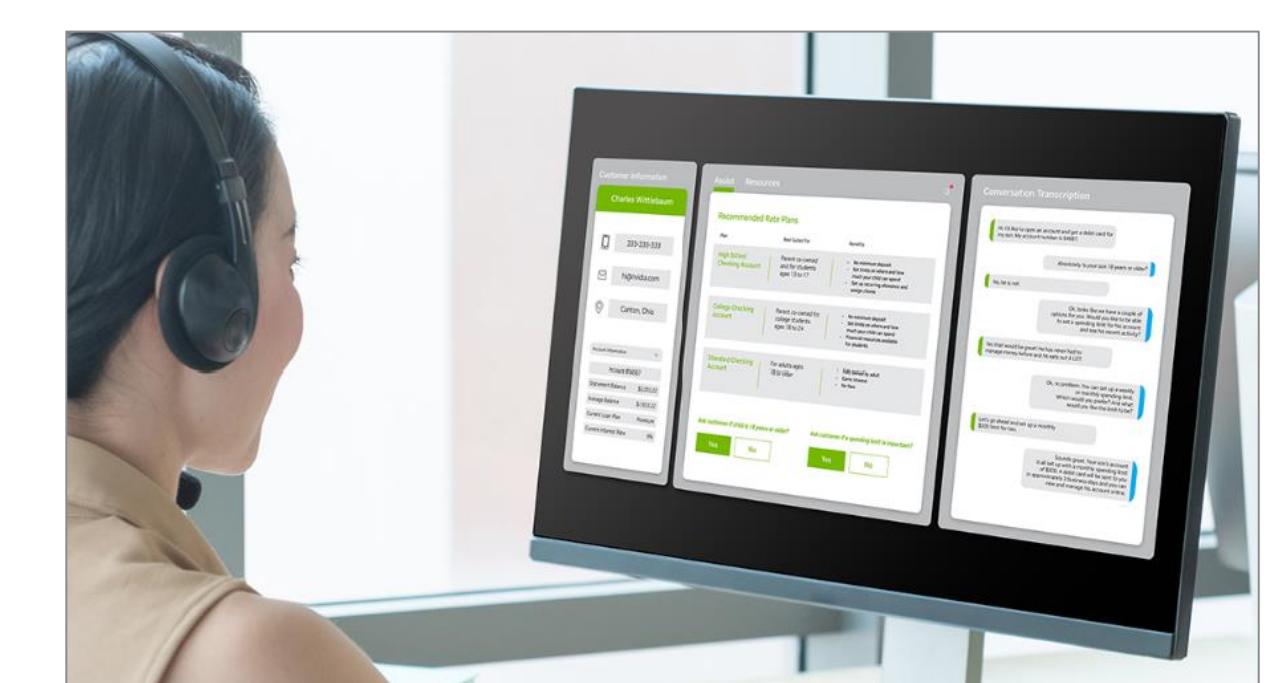
Legal Services, Education ...
1M legal professionals in the US
9M educators in the US



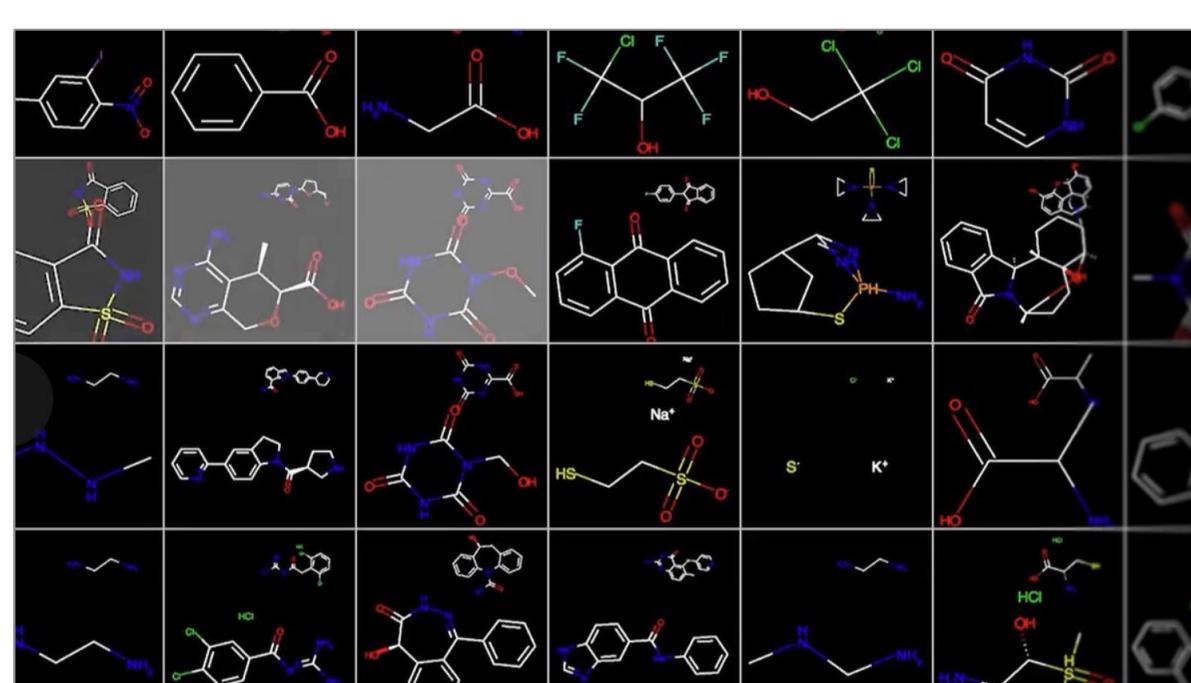
AI Software Development
30M software developers globally



Financial Services
678B annual credit card transactions



Customer Service with AI
15M call center agents globally



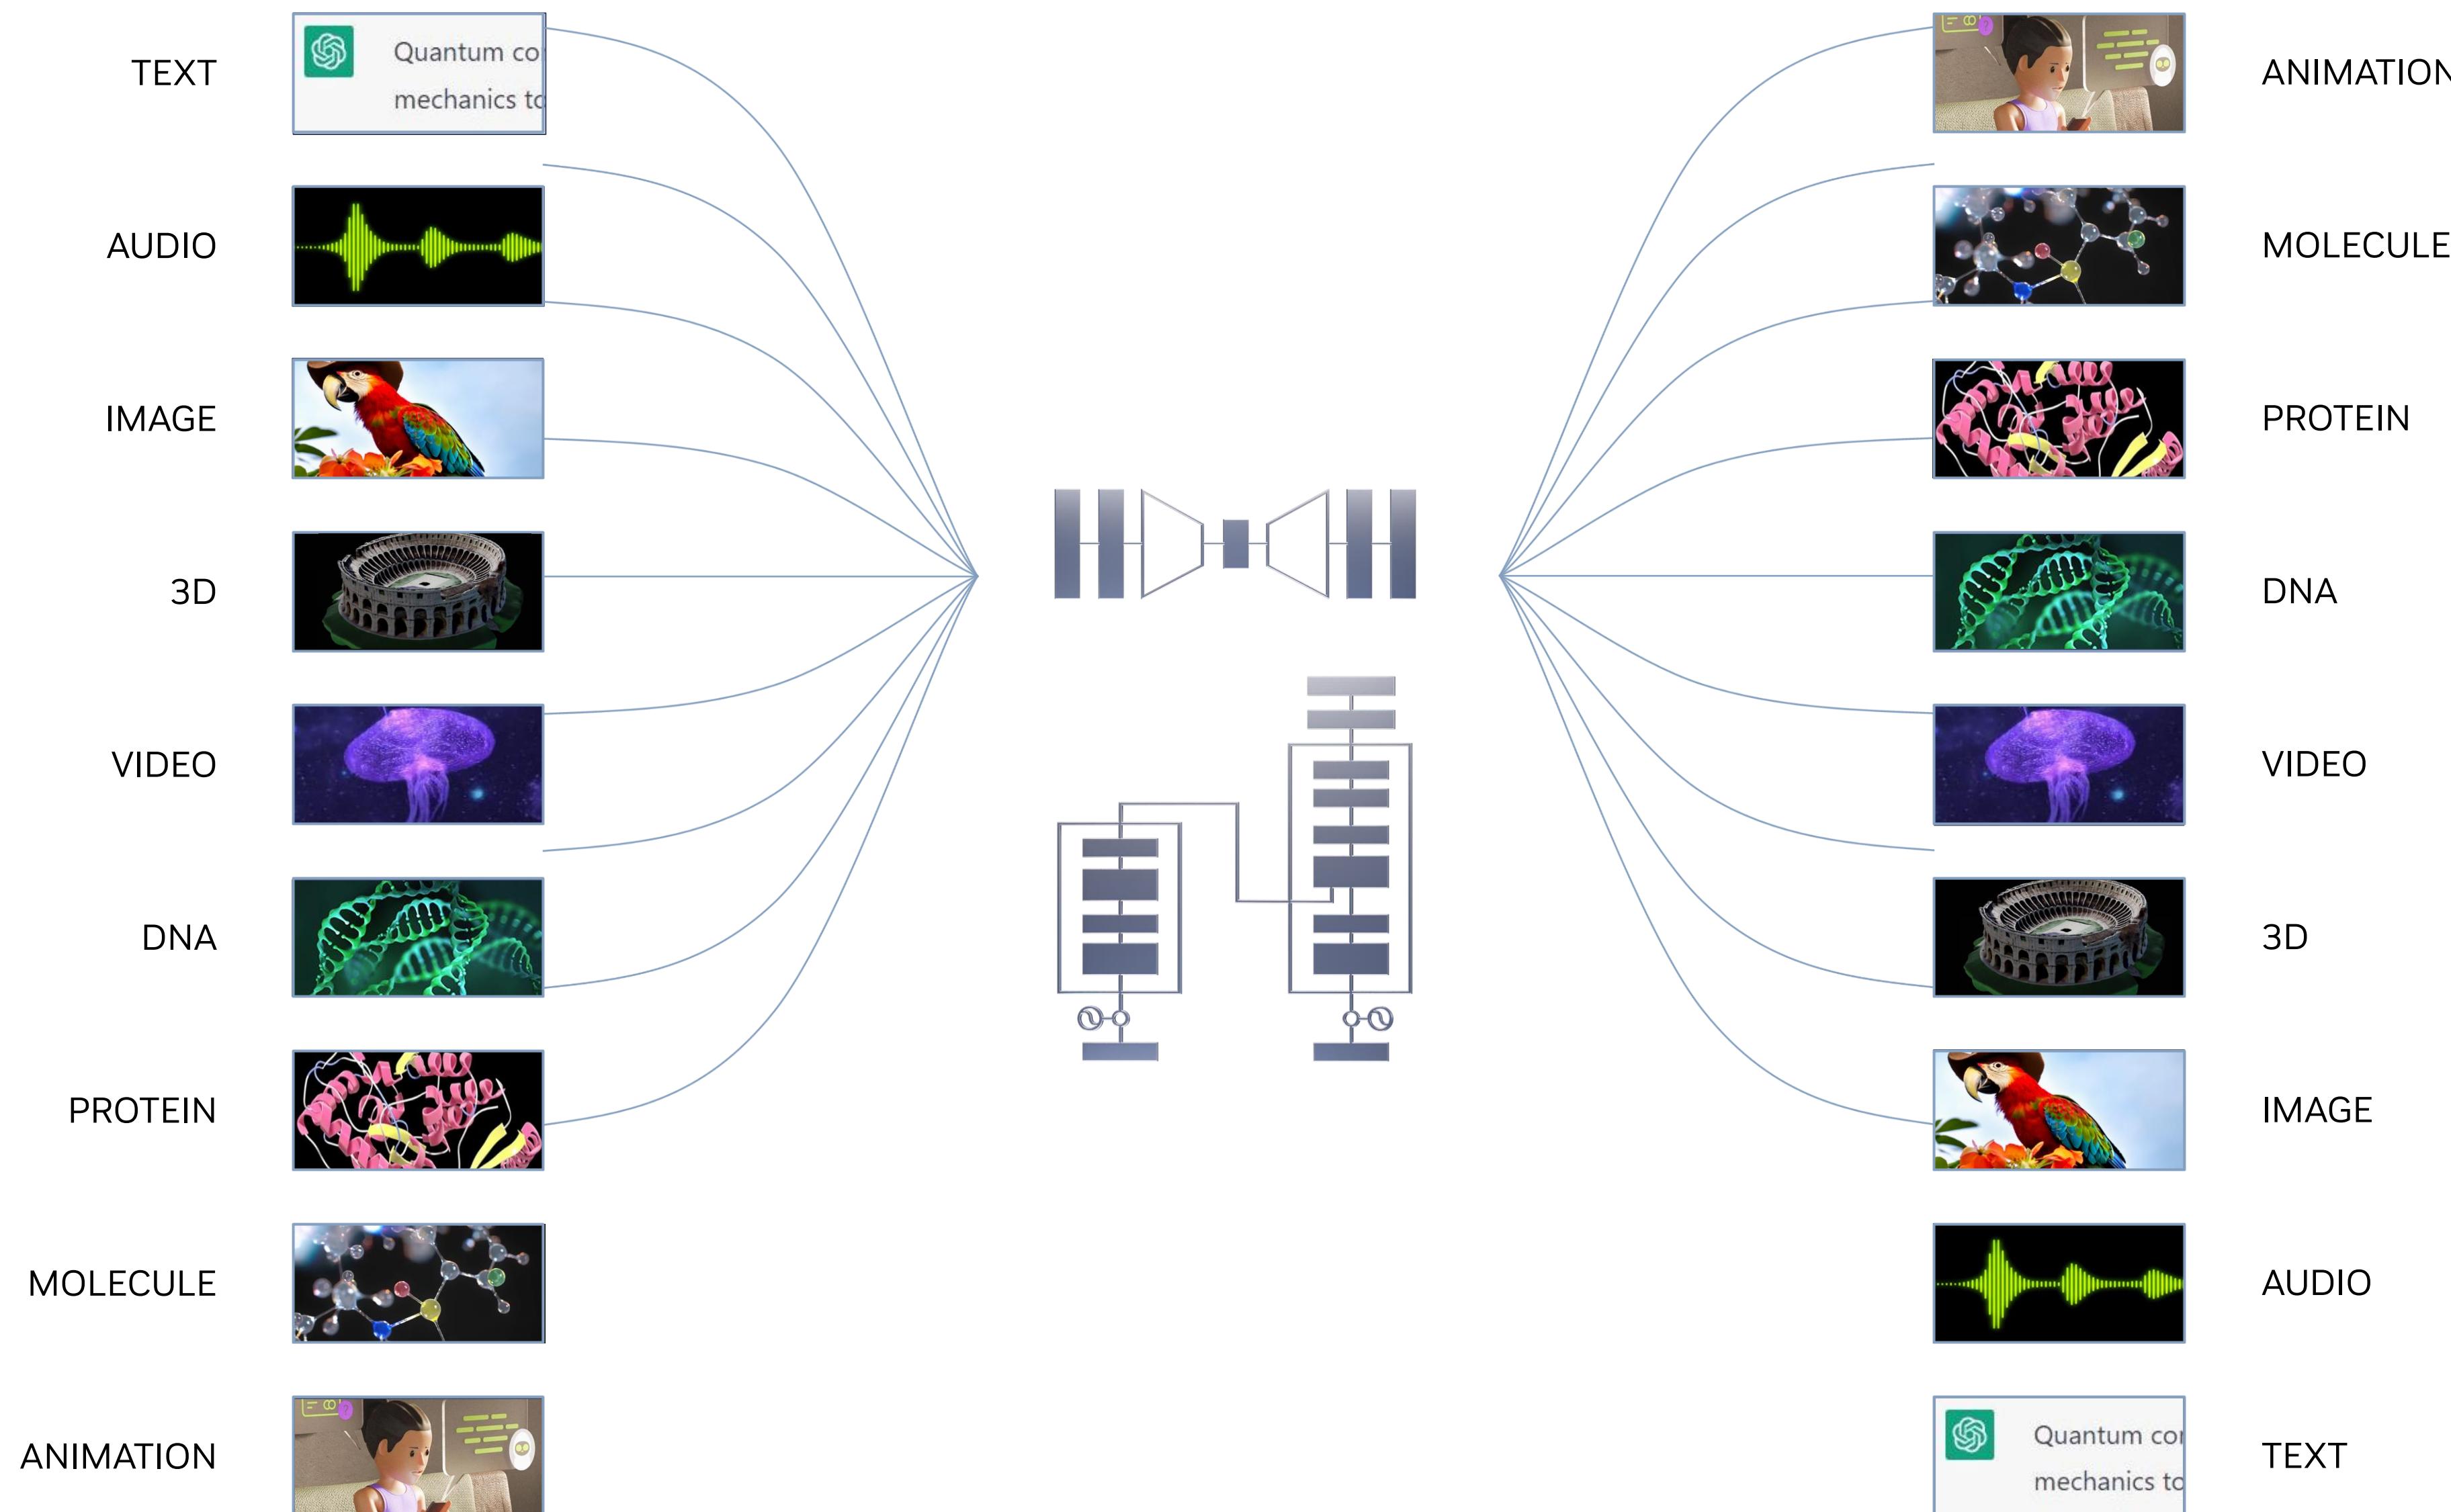
Drug Discovery
10¹⁸ molecules in chemical space
40 exabytes of genome data



Agri-Food | Climate
1B people in agri-food worldwide
Earth-2 for km-scale simulation

Generative AI

The most important computing platform of our generation



The era of generative AI has arrived, unlocking new opportunities for AI across many different applications.

Generative AI is trained on large amounts of data to find patterns and relationships, learning the representation of almost anything with structure.

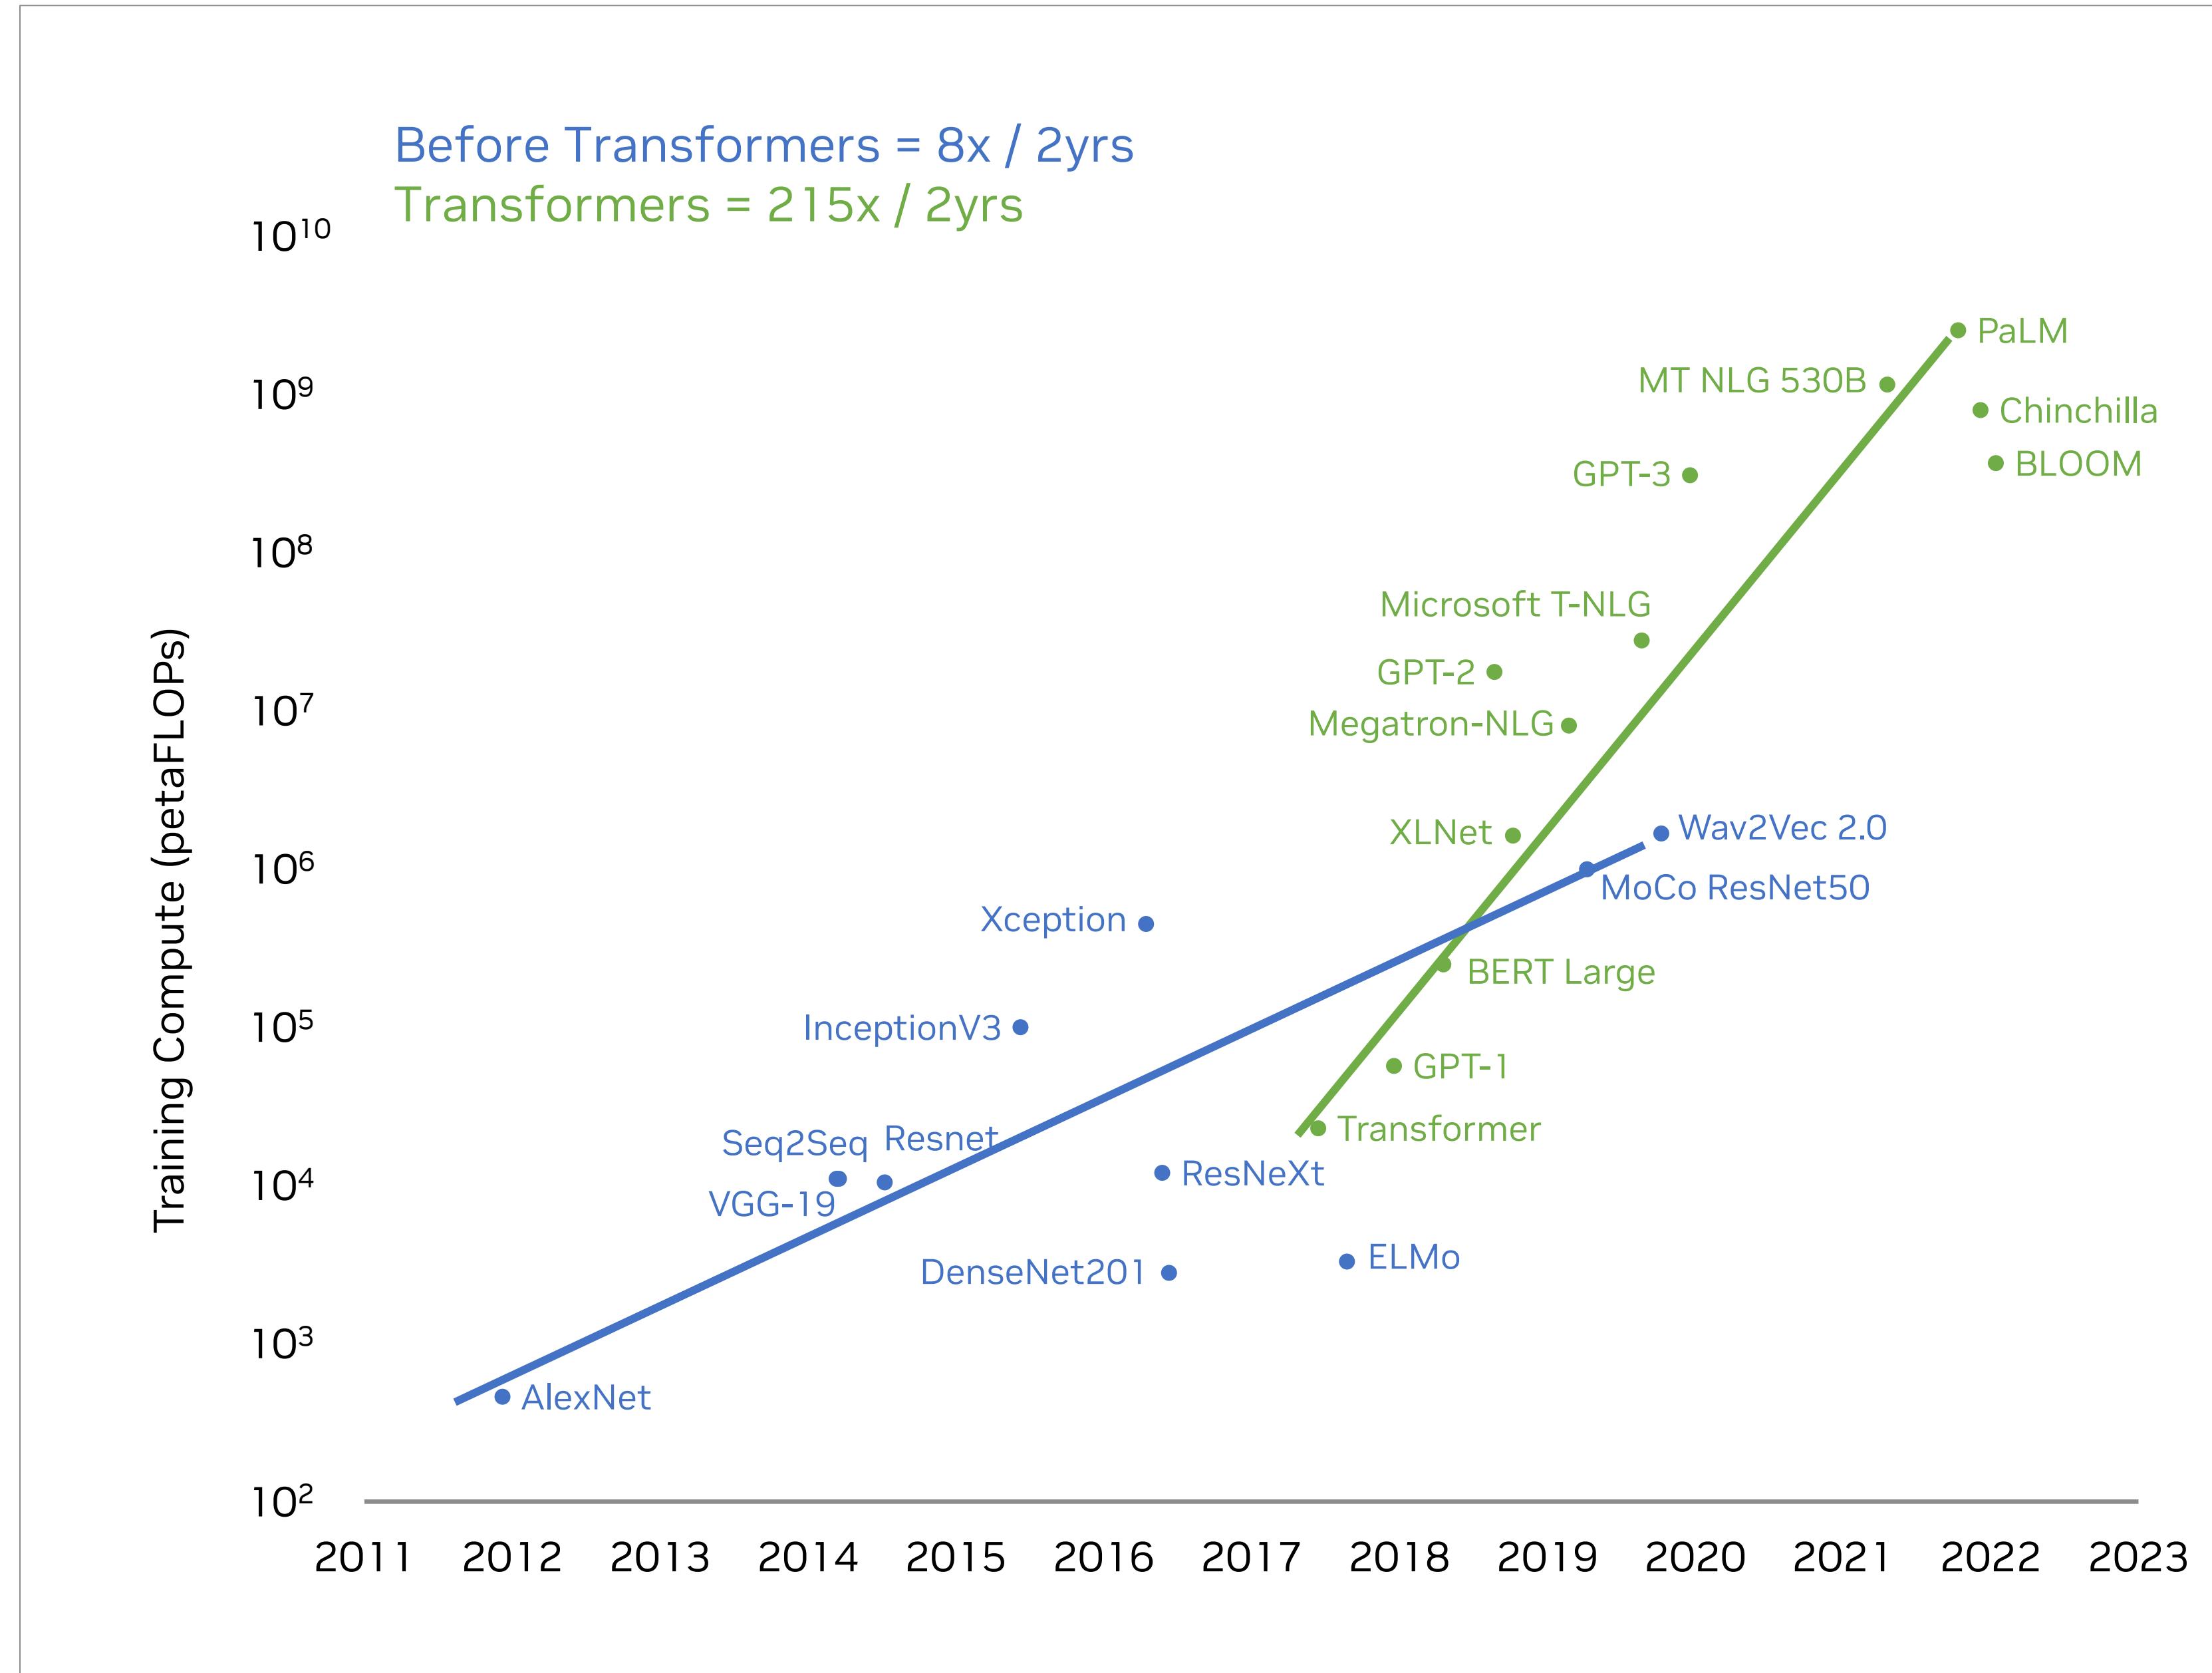
It can then be prompted to generate text, images, video, code, or even proteins.

For the very first time, computers can augment the human ability to generate information and create.

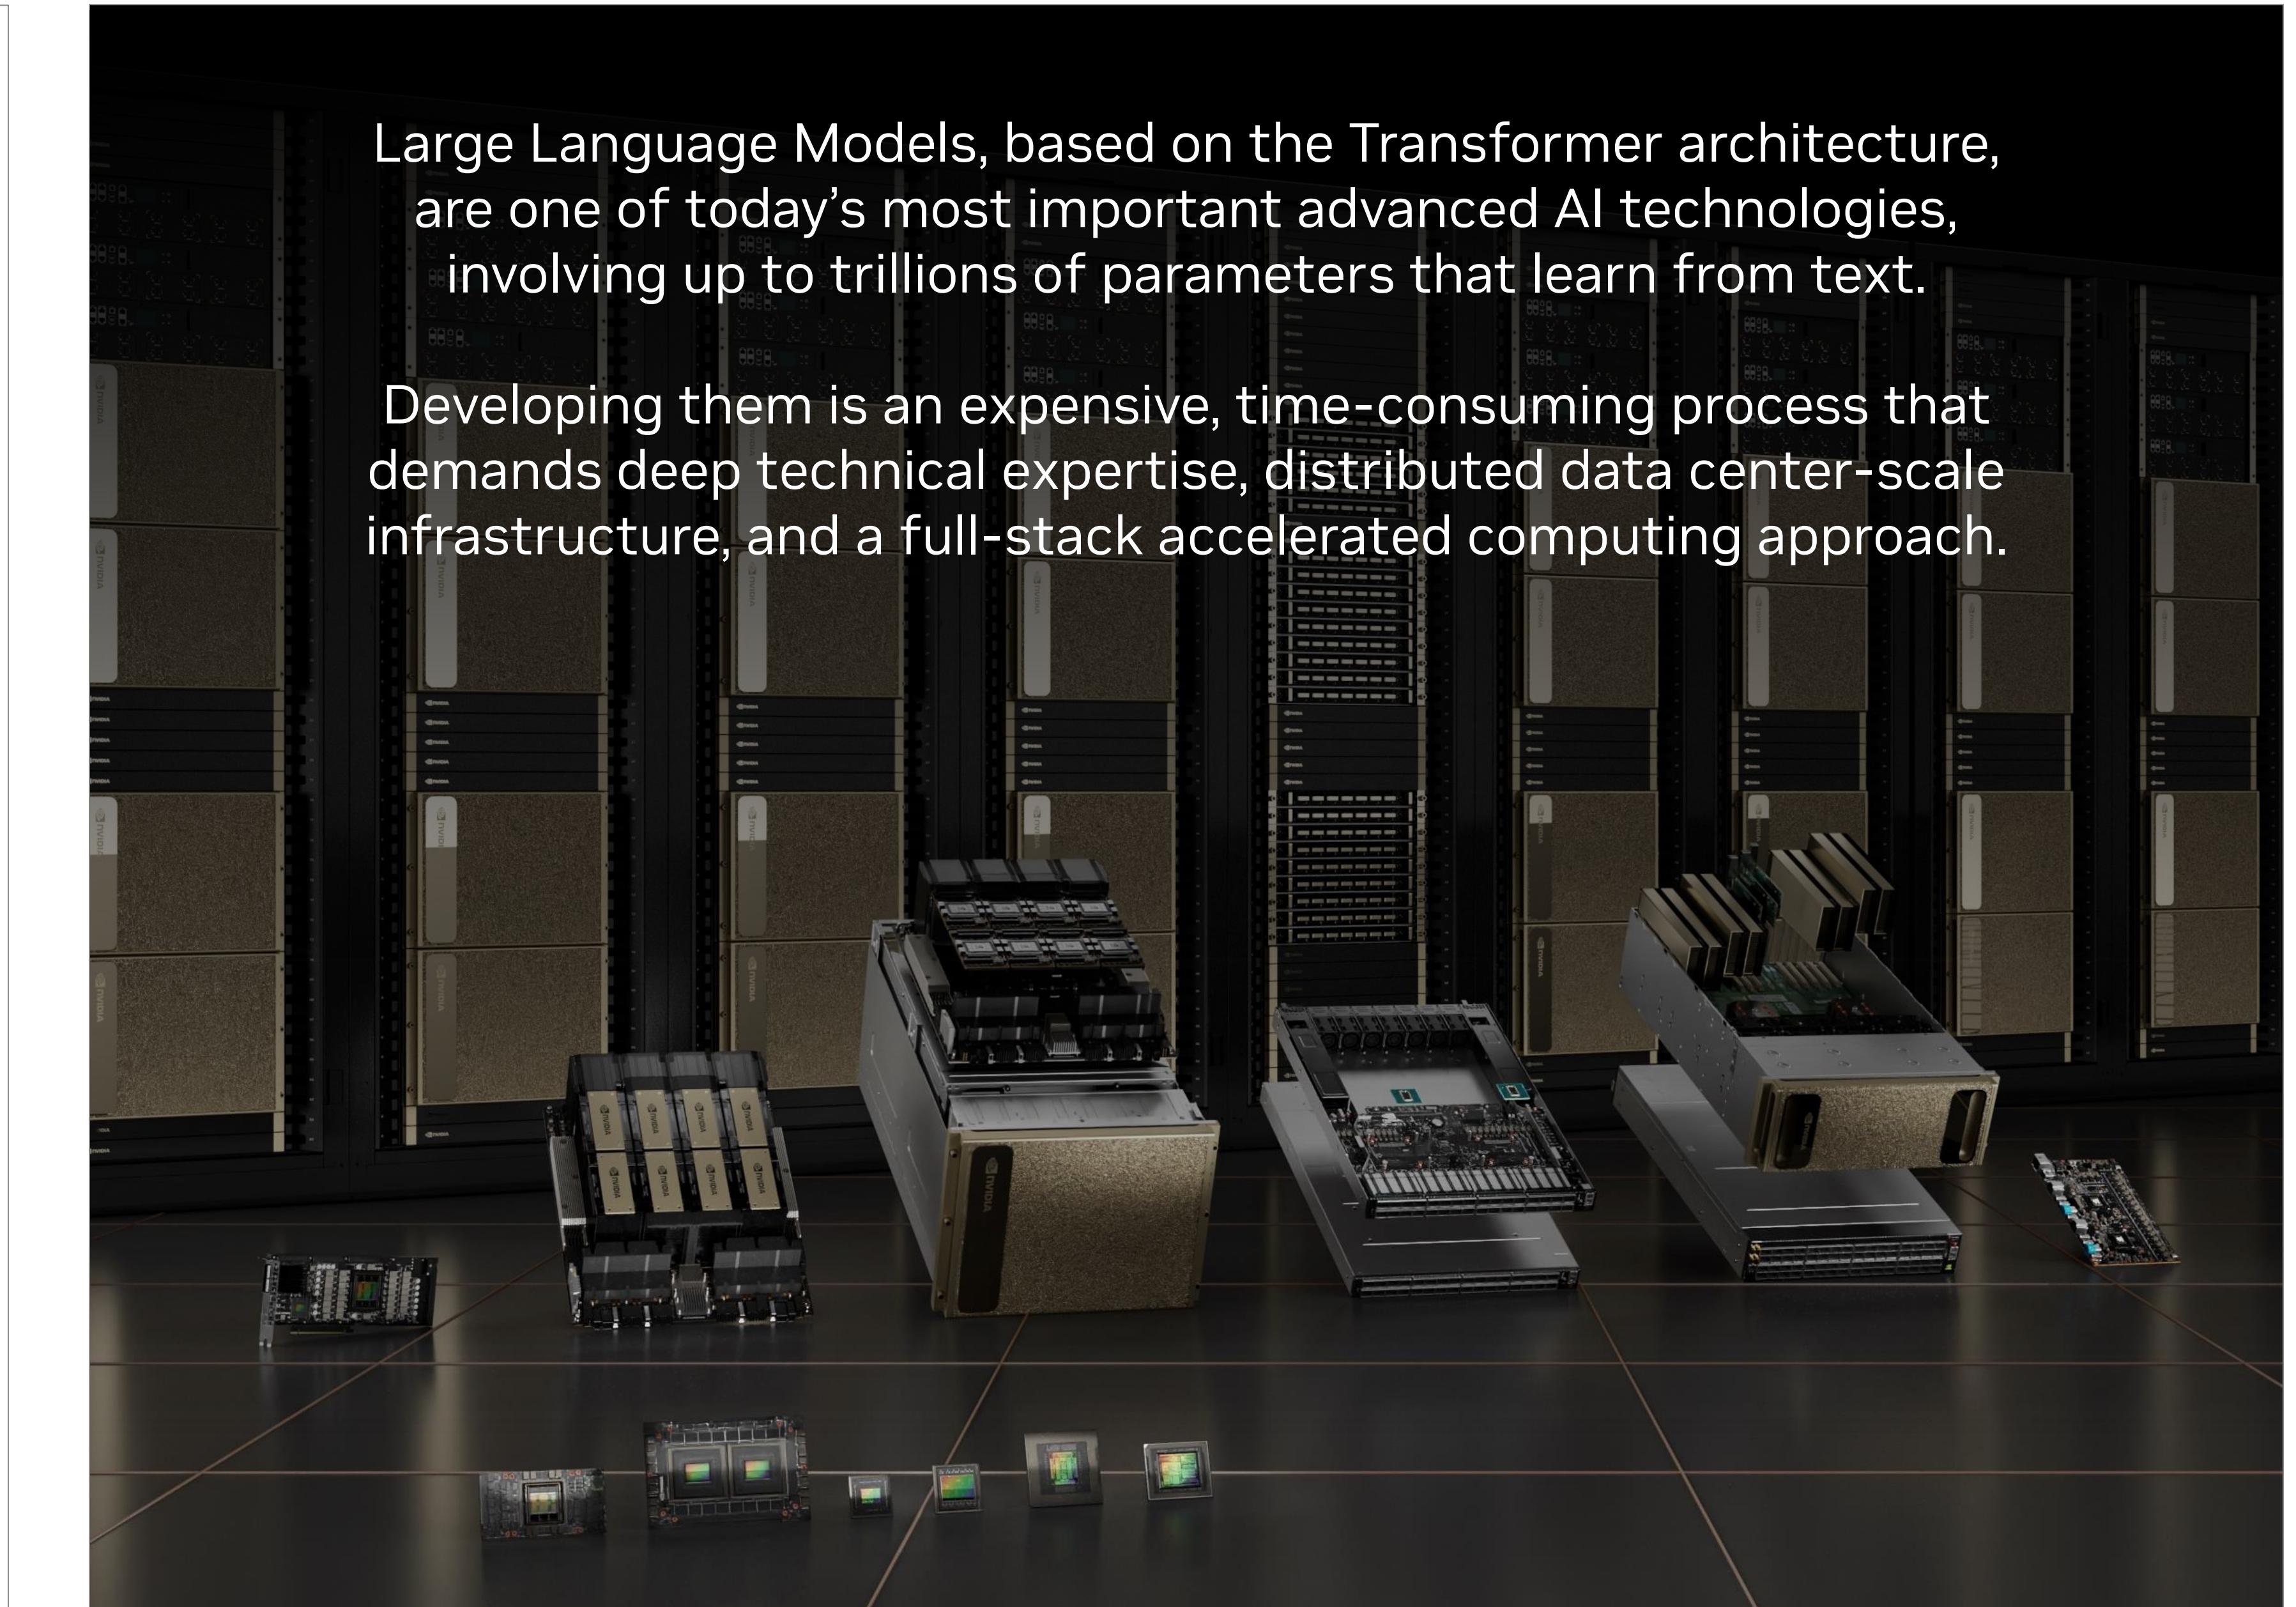
1,600+ Generative AI companies are building on NVIDIA.

Modern AI is a Data Center Scale Computing Workload

Data centers are becoming AI factories: Data as input, intelligence as output



AI Training Computational Requirements



Fueling Giant-Scale AI Infrastructure
NVIDIA compute & networking GPU | DPU | CPU

Full-Stack & Data Center Scale Acceleration

Drive significant cost savings and workload scaling

Classical Computing — 960 CPU-only servers

Application

CPU server racks

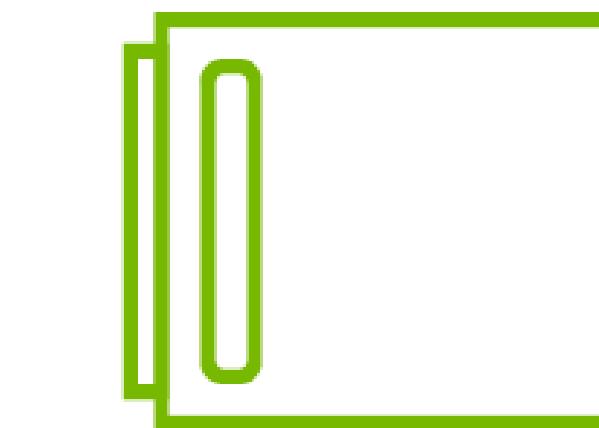


Accelerated Computing — 2 GPU servers

Application
Re-Engineered for Acceleration

CUDA-X Acceleration Libraries

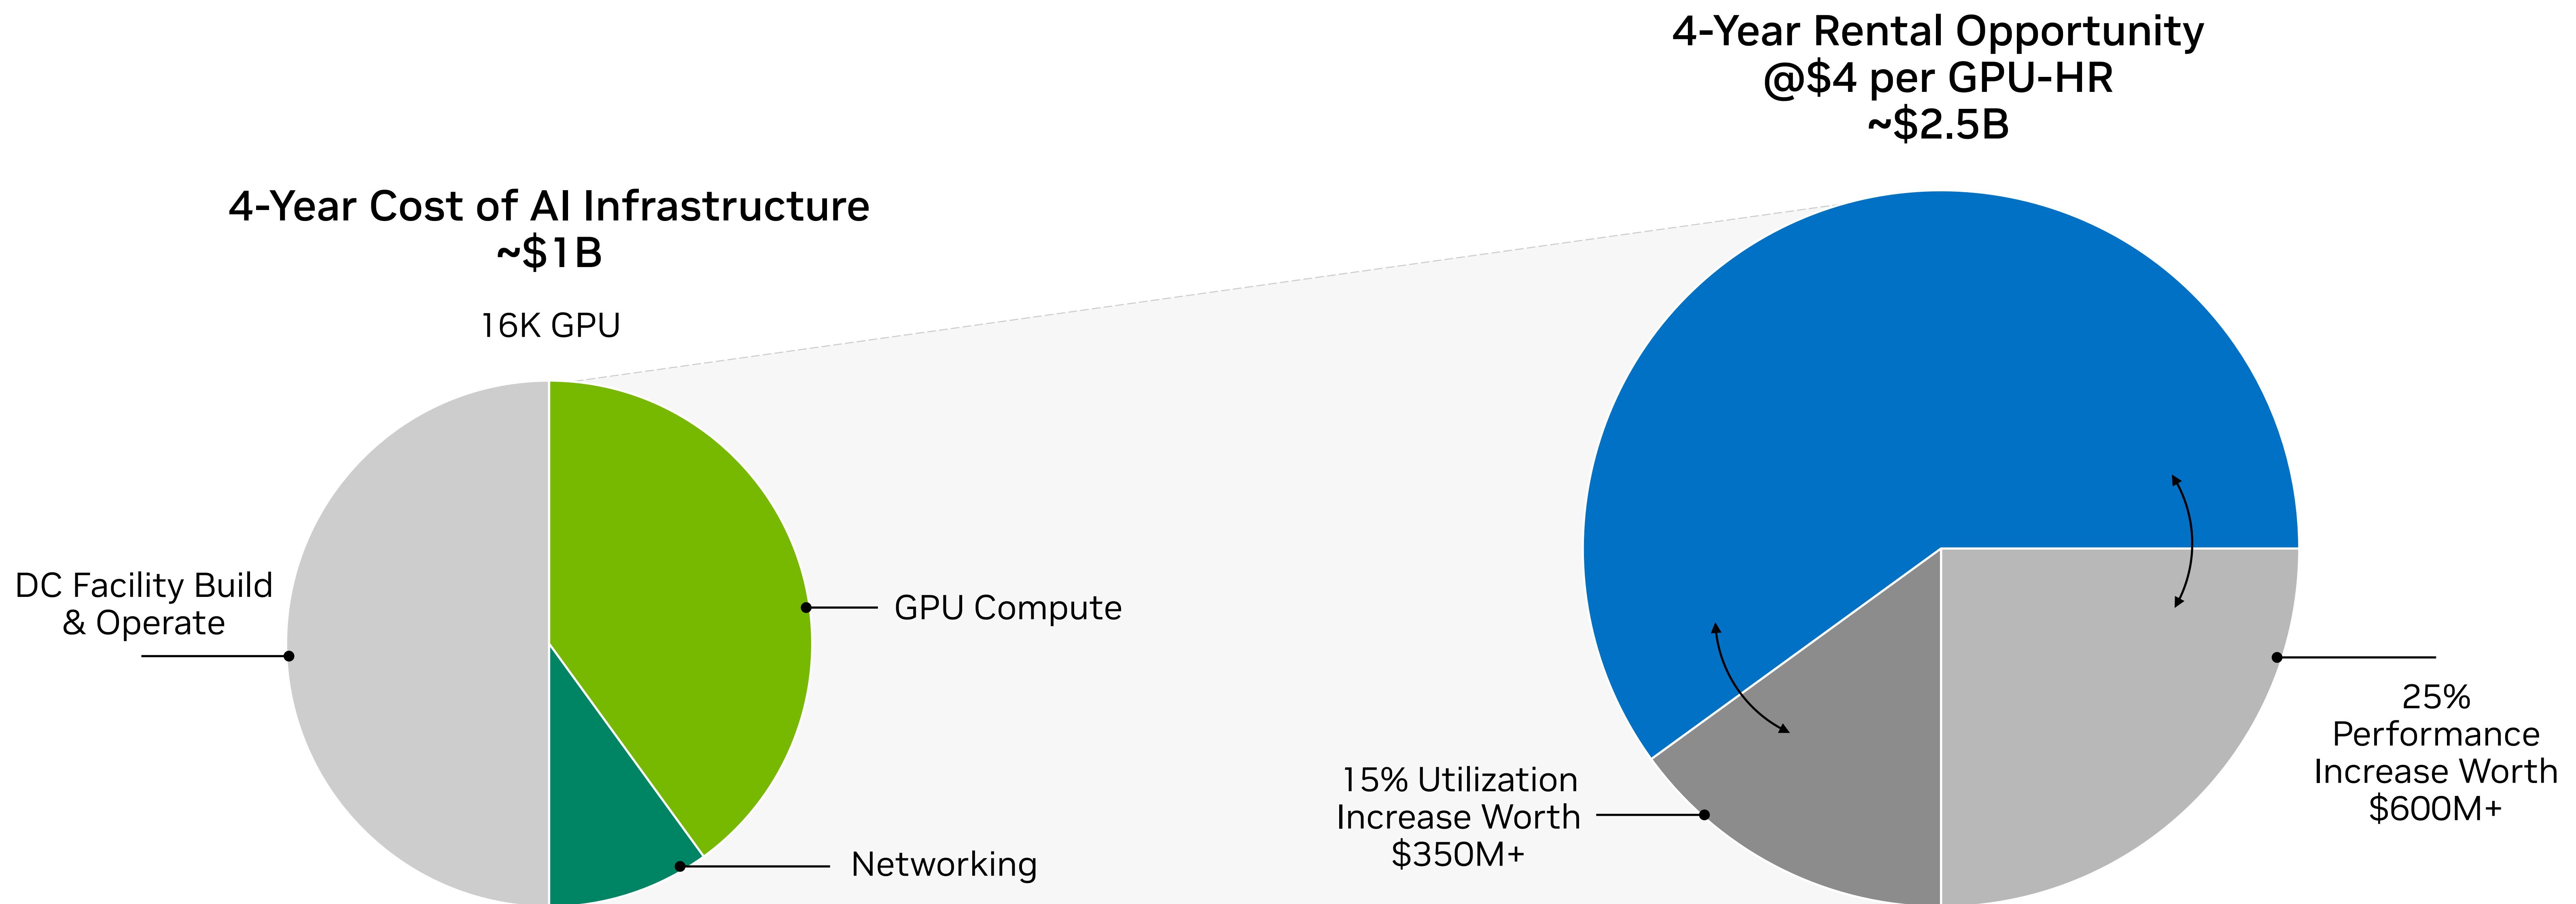
Magnum IO



**25X lower cost
84X better energy-efficiency**

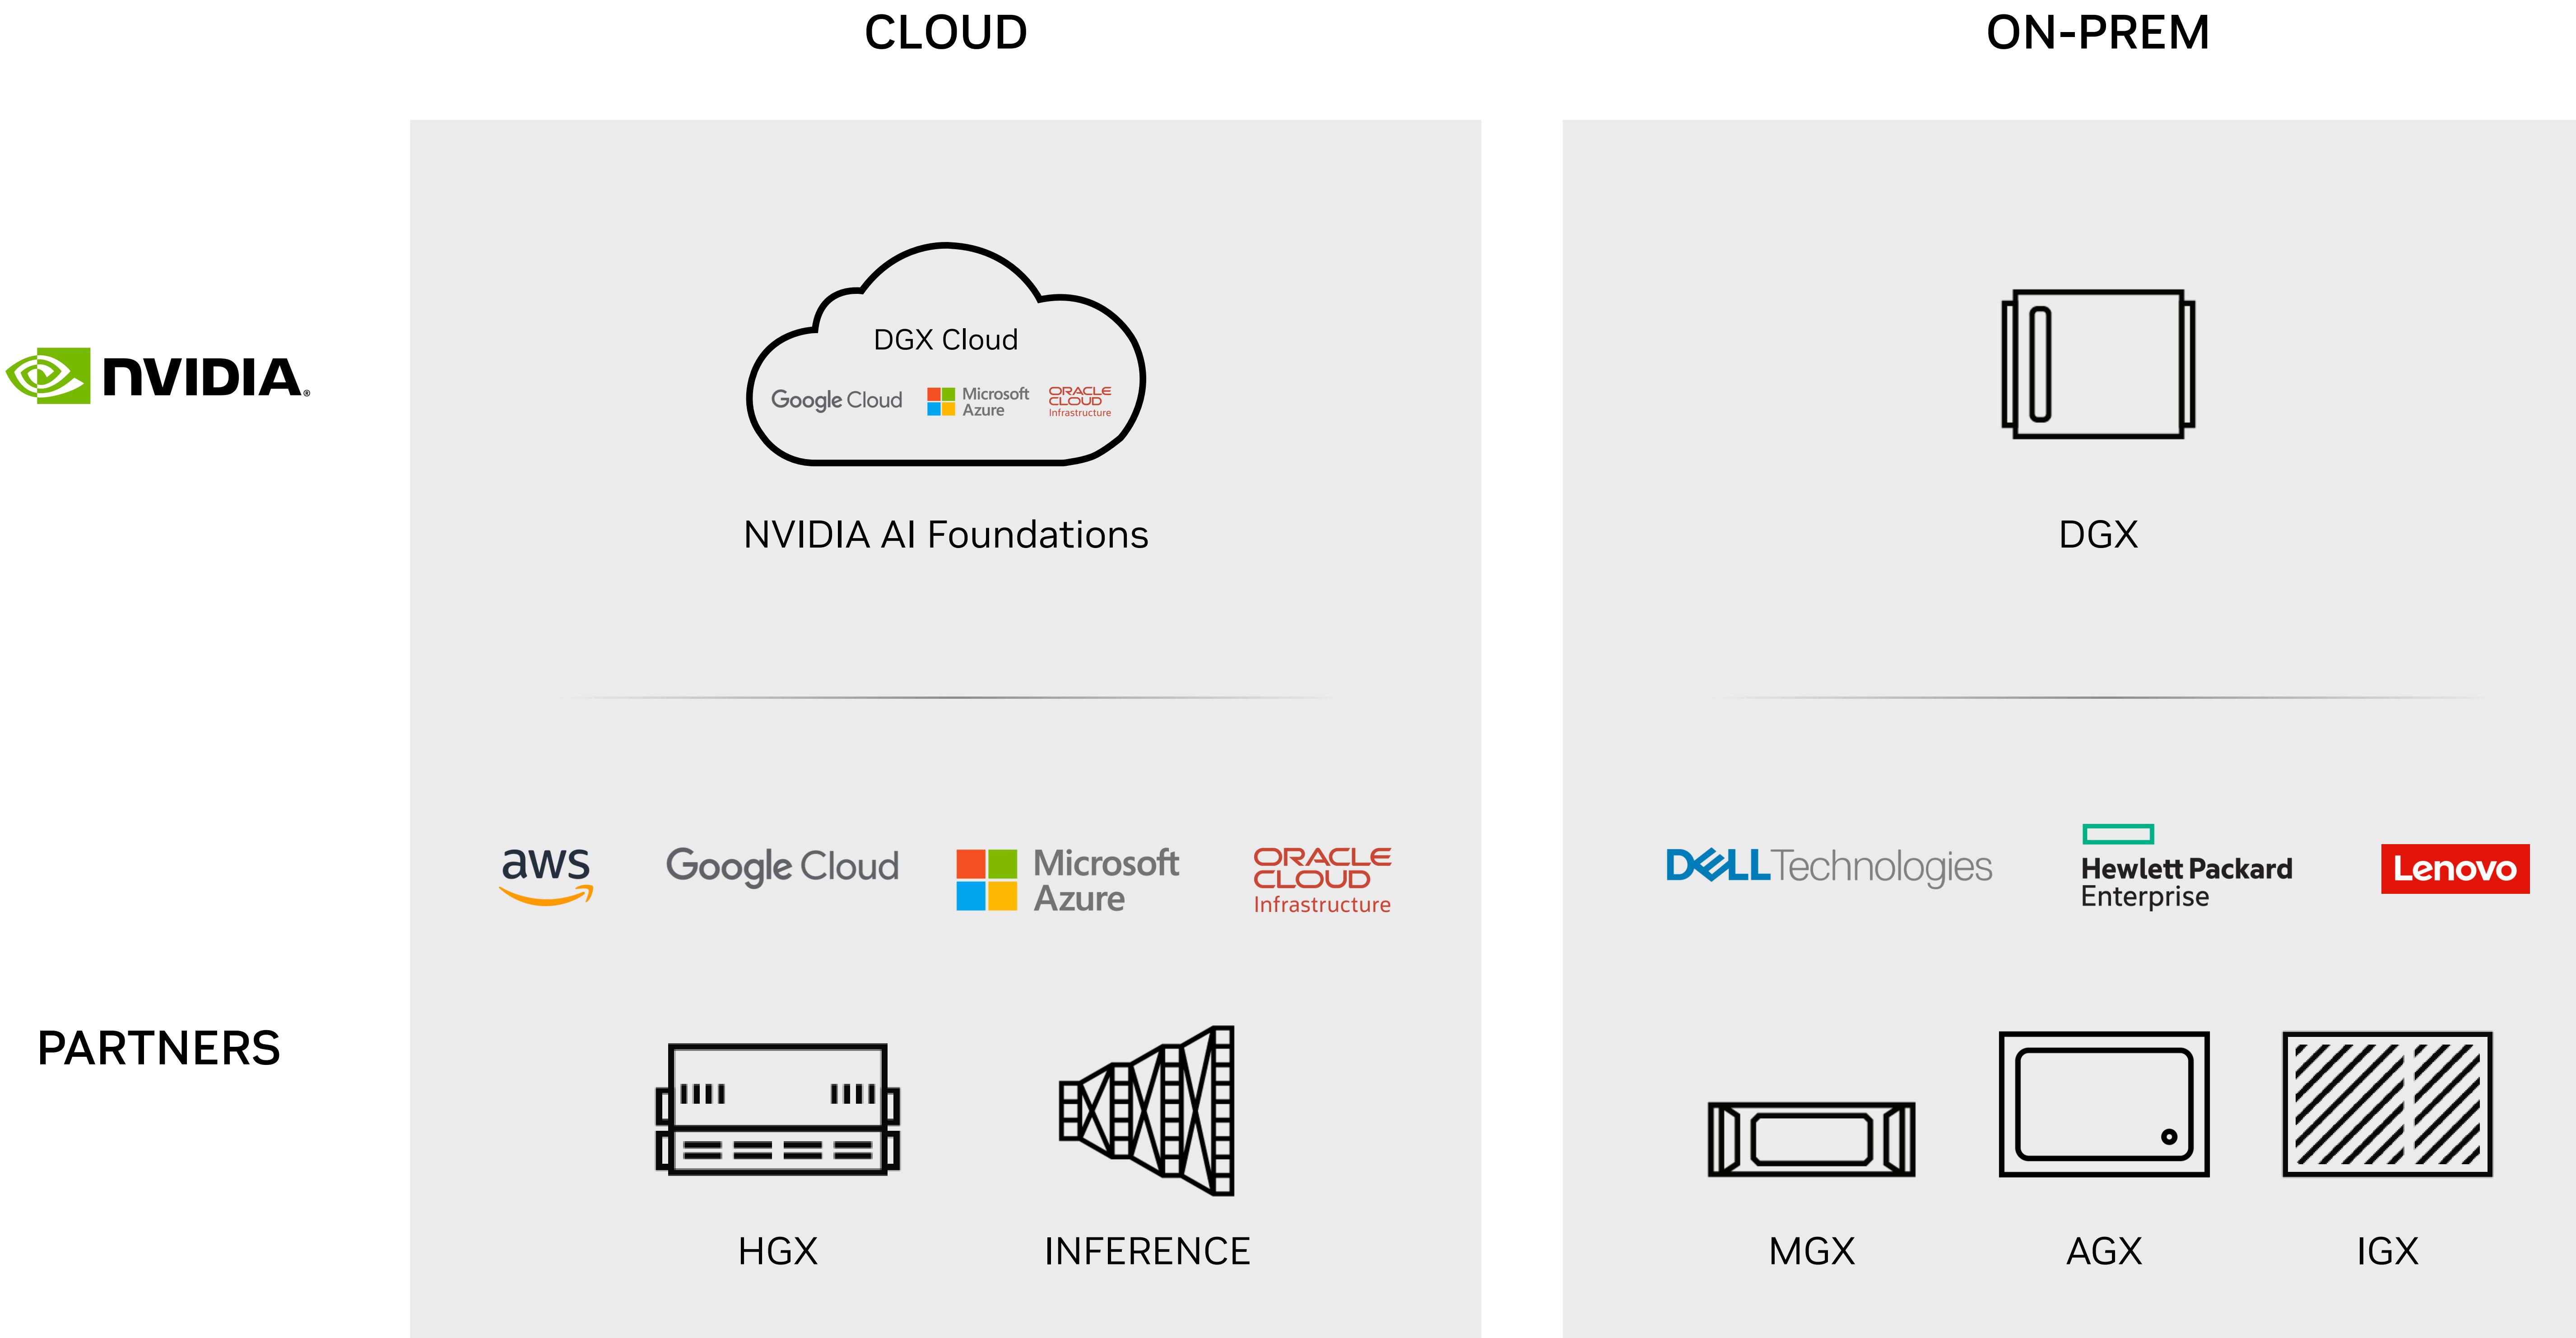
LLM Workload: Bert-Large Training and Inference | CPU Server: Dual-EYPC 7763 | GPU Server: Dual-EPYC 7763 + 8X H100 PCIe GPUs

The High ROI of High Compute Performance



NVIDIA Go-to-Market Across Cloud and On-Premises

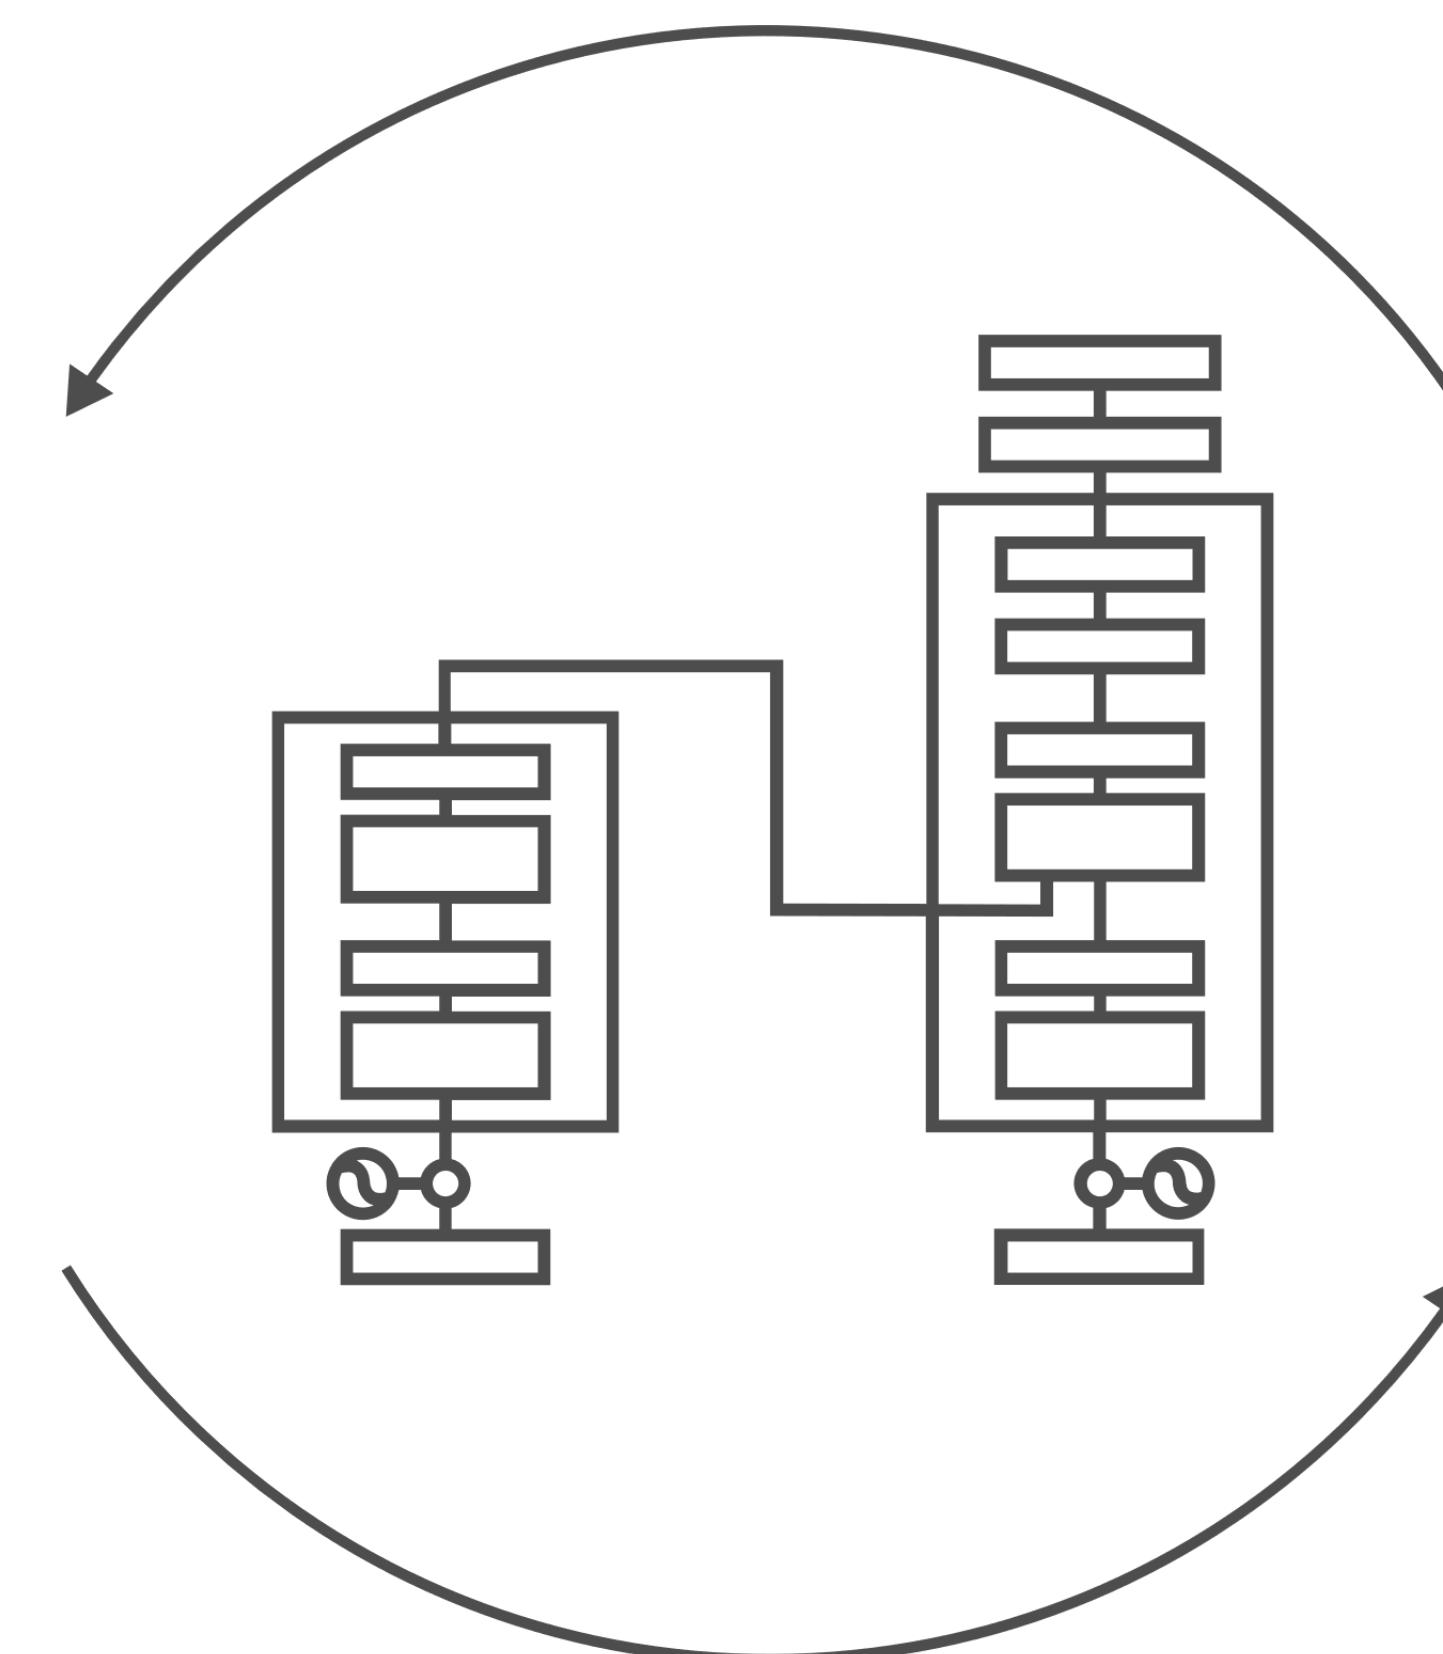
Reaching customers everywhere



Training & Inference — One Architecture

Cloud | On-Prem | Edge

TRAINING



INFERENCE

IN THE DATA CENTER

NVIDIA L40
Image Generation

NVIDIA L4
AI Video

NVIDIA H100 | L40S
Universal GPUs

NVIDIA Grace Hopper
RecSys, Gen AI

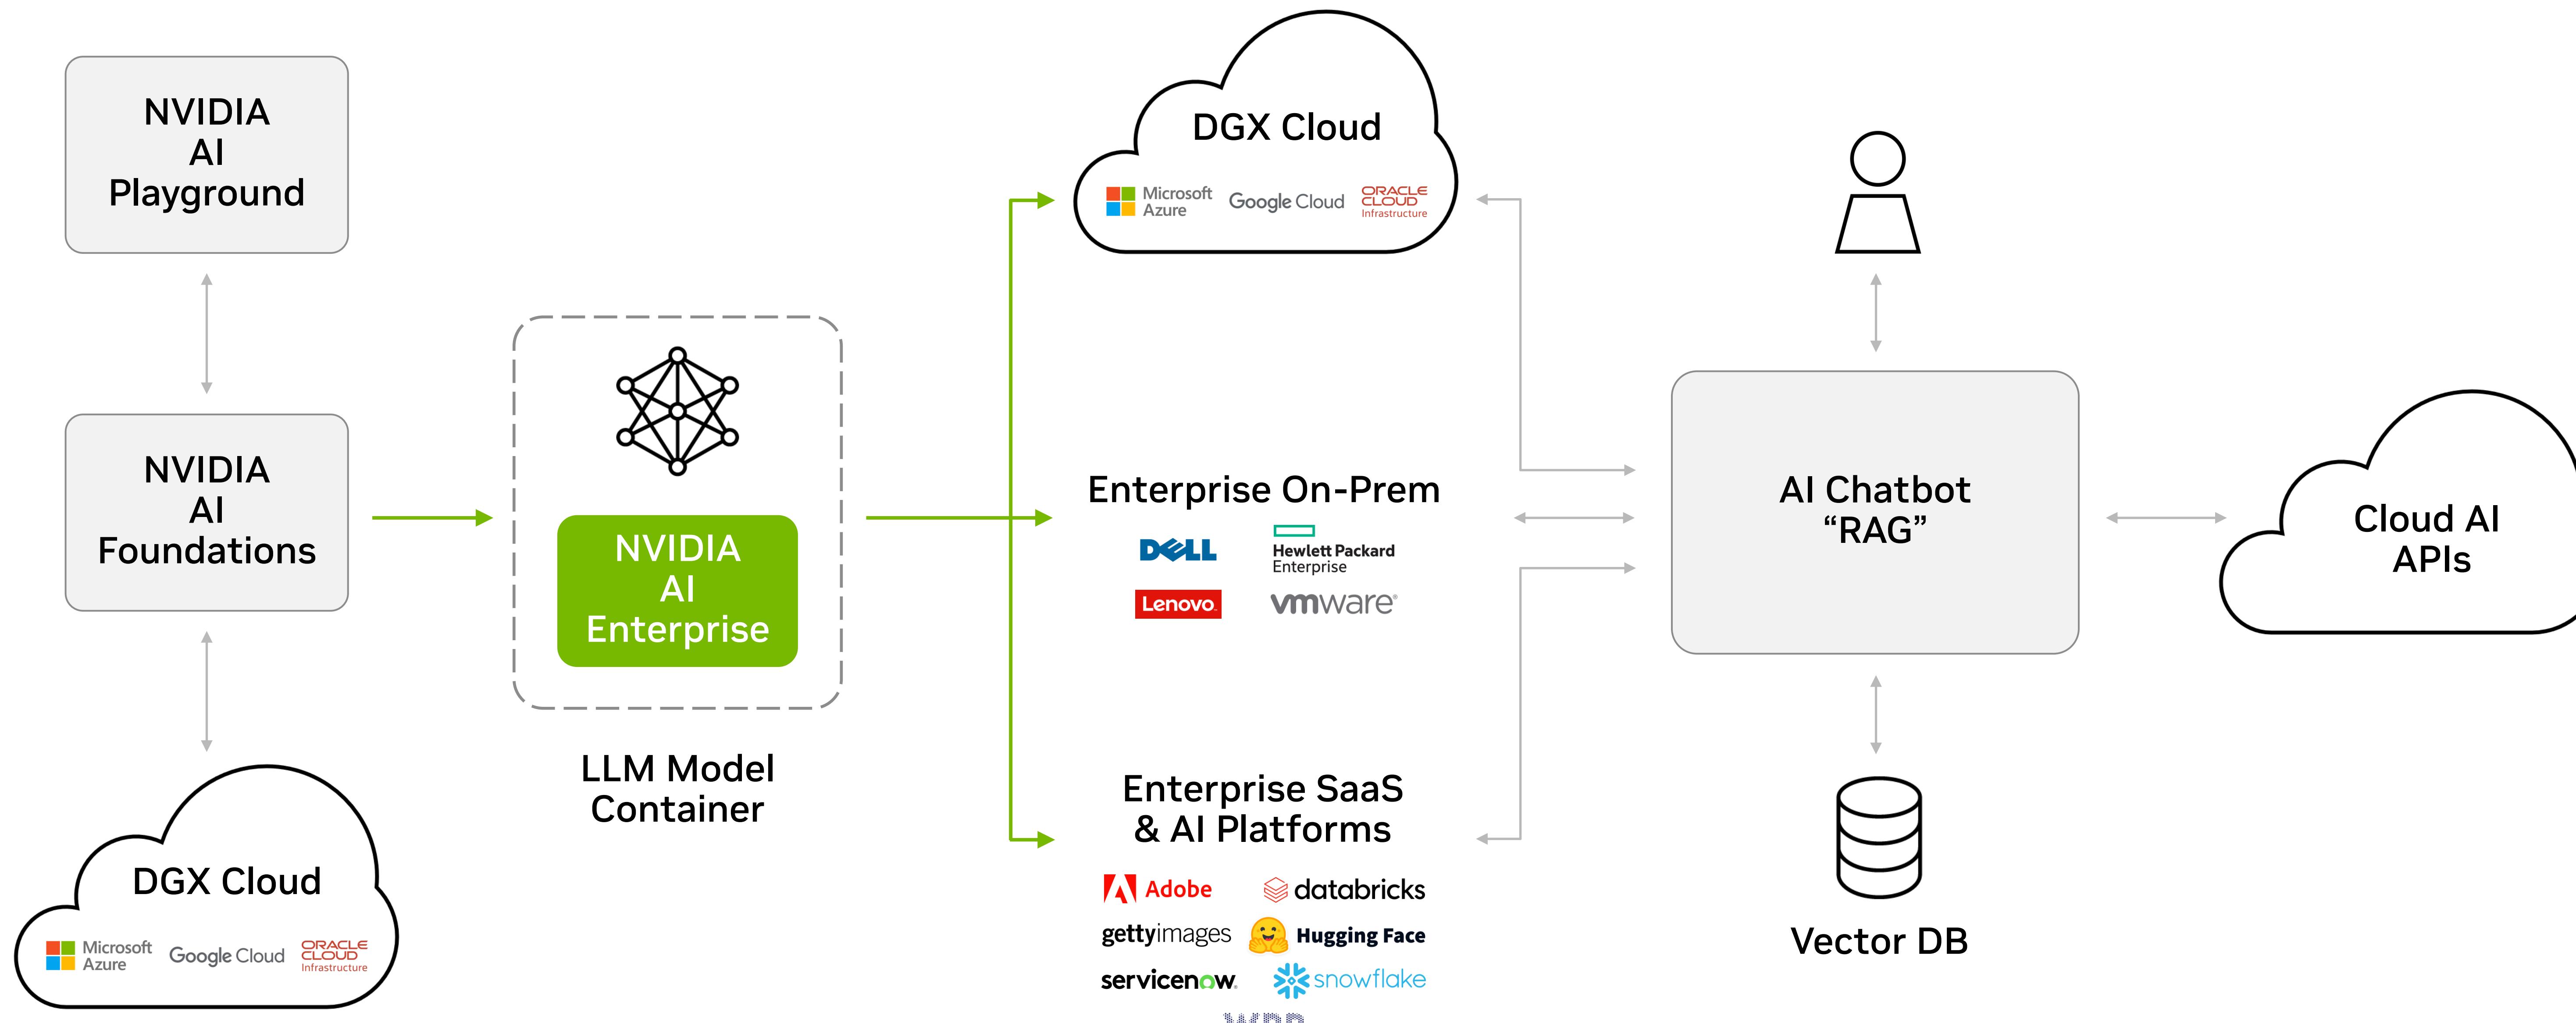
AT THE EDGE

IGX
Industrial-Grade System
for Healthcare, Logistics,
Manufacturing

AGX
Functionally-Safe System
for Autonomous Vehicles

NVIDIA AI Enterprise

Enterprise is the Next Big Generative AI Opportunity



Enterprise AI Chatbots are built as Retrieval Augmented Generation (RAG) workflows, which augment the knowledge in the LLM with vectorized Enterprise data. These Chatbots serve as apprentices, improving the productivity of every employee in every Enterprise company.

NVIDIA delivers this capability to Enterprises by packaging LLMs with NVIDIA AI Enterprise, the runtime for hosting the LLMs, into containers that can be deployed anywhere – on any cloud, on premises, or within Enterprise SaaS applications.

NVIDIA DGX Cloud

AI-training-as-a-service platform for the era of generative AI

NVIDIA AI FOUNDATIONS

NeMo | Picasso



DGX CLOUD



Microsoft
Azure



NVIDIA DGX Cloud is a cloud service that allows enterprises immediate access to the infrastructure and software needed to train advanced models for generative AI and other groundbreaking applications.

DGX Cloud provides dedicated clusters of NVIDIA DGX AI supercomputing, paired with NVIDIA AI software.

Enterprise customers can also use the NVIDIA AI Foundations model making service, which includes NVIDIA NeMo for training custom LLMs and NVIDIA Picasso for custom generative AI models for visual design.

The service is equipped with models, tools, and accelerated computing for training, customizing, optimizing, and deploying AI.

NVIDIA AI Enterprise

The operating system for enterprise AI

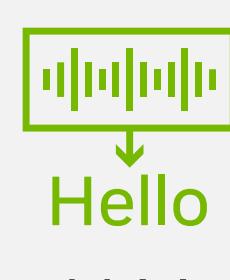
NVIDIA AI Enterprise

NVIDIA AI Enterprise is software for deploying and running AI with enterprise-grade security, API stability, manageability and support.

Cloud-native and available in every major cloud marketplace — AWS, Microsoft Azure, Google Cloud Platform and Oracle Cloud.

Certified to run on servers and workstations from all major OEMs.

AI Use Cases and Workflows



LLM



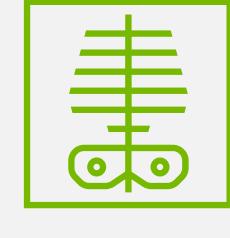
Speech AI



Recommenders



Cybersecurity



Medical Imaging



Video Analytics

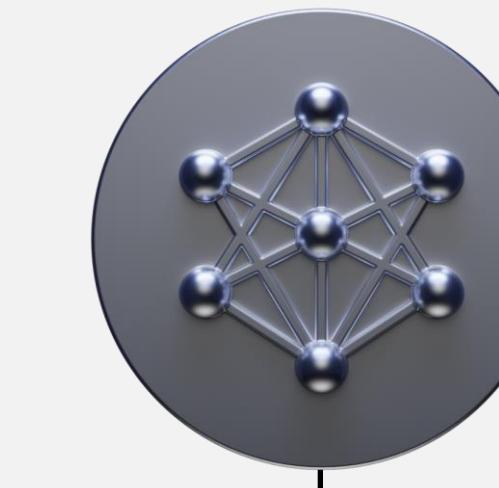


Route Optimization

...

More

Run Anywhere



NVIDIA AI Enterprise



Azure | GCP | OCI | AWS



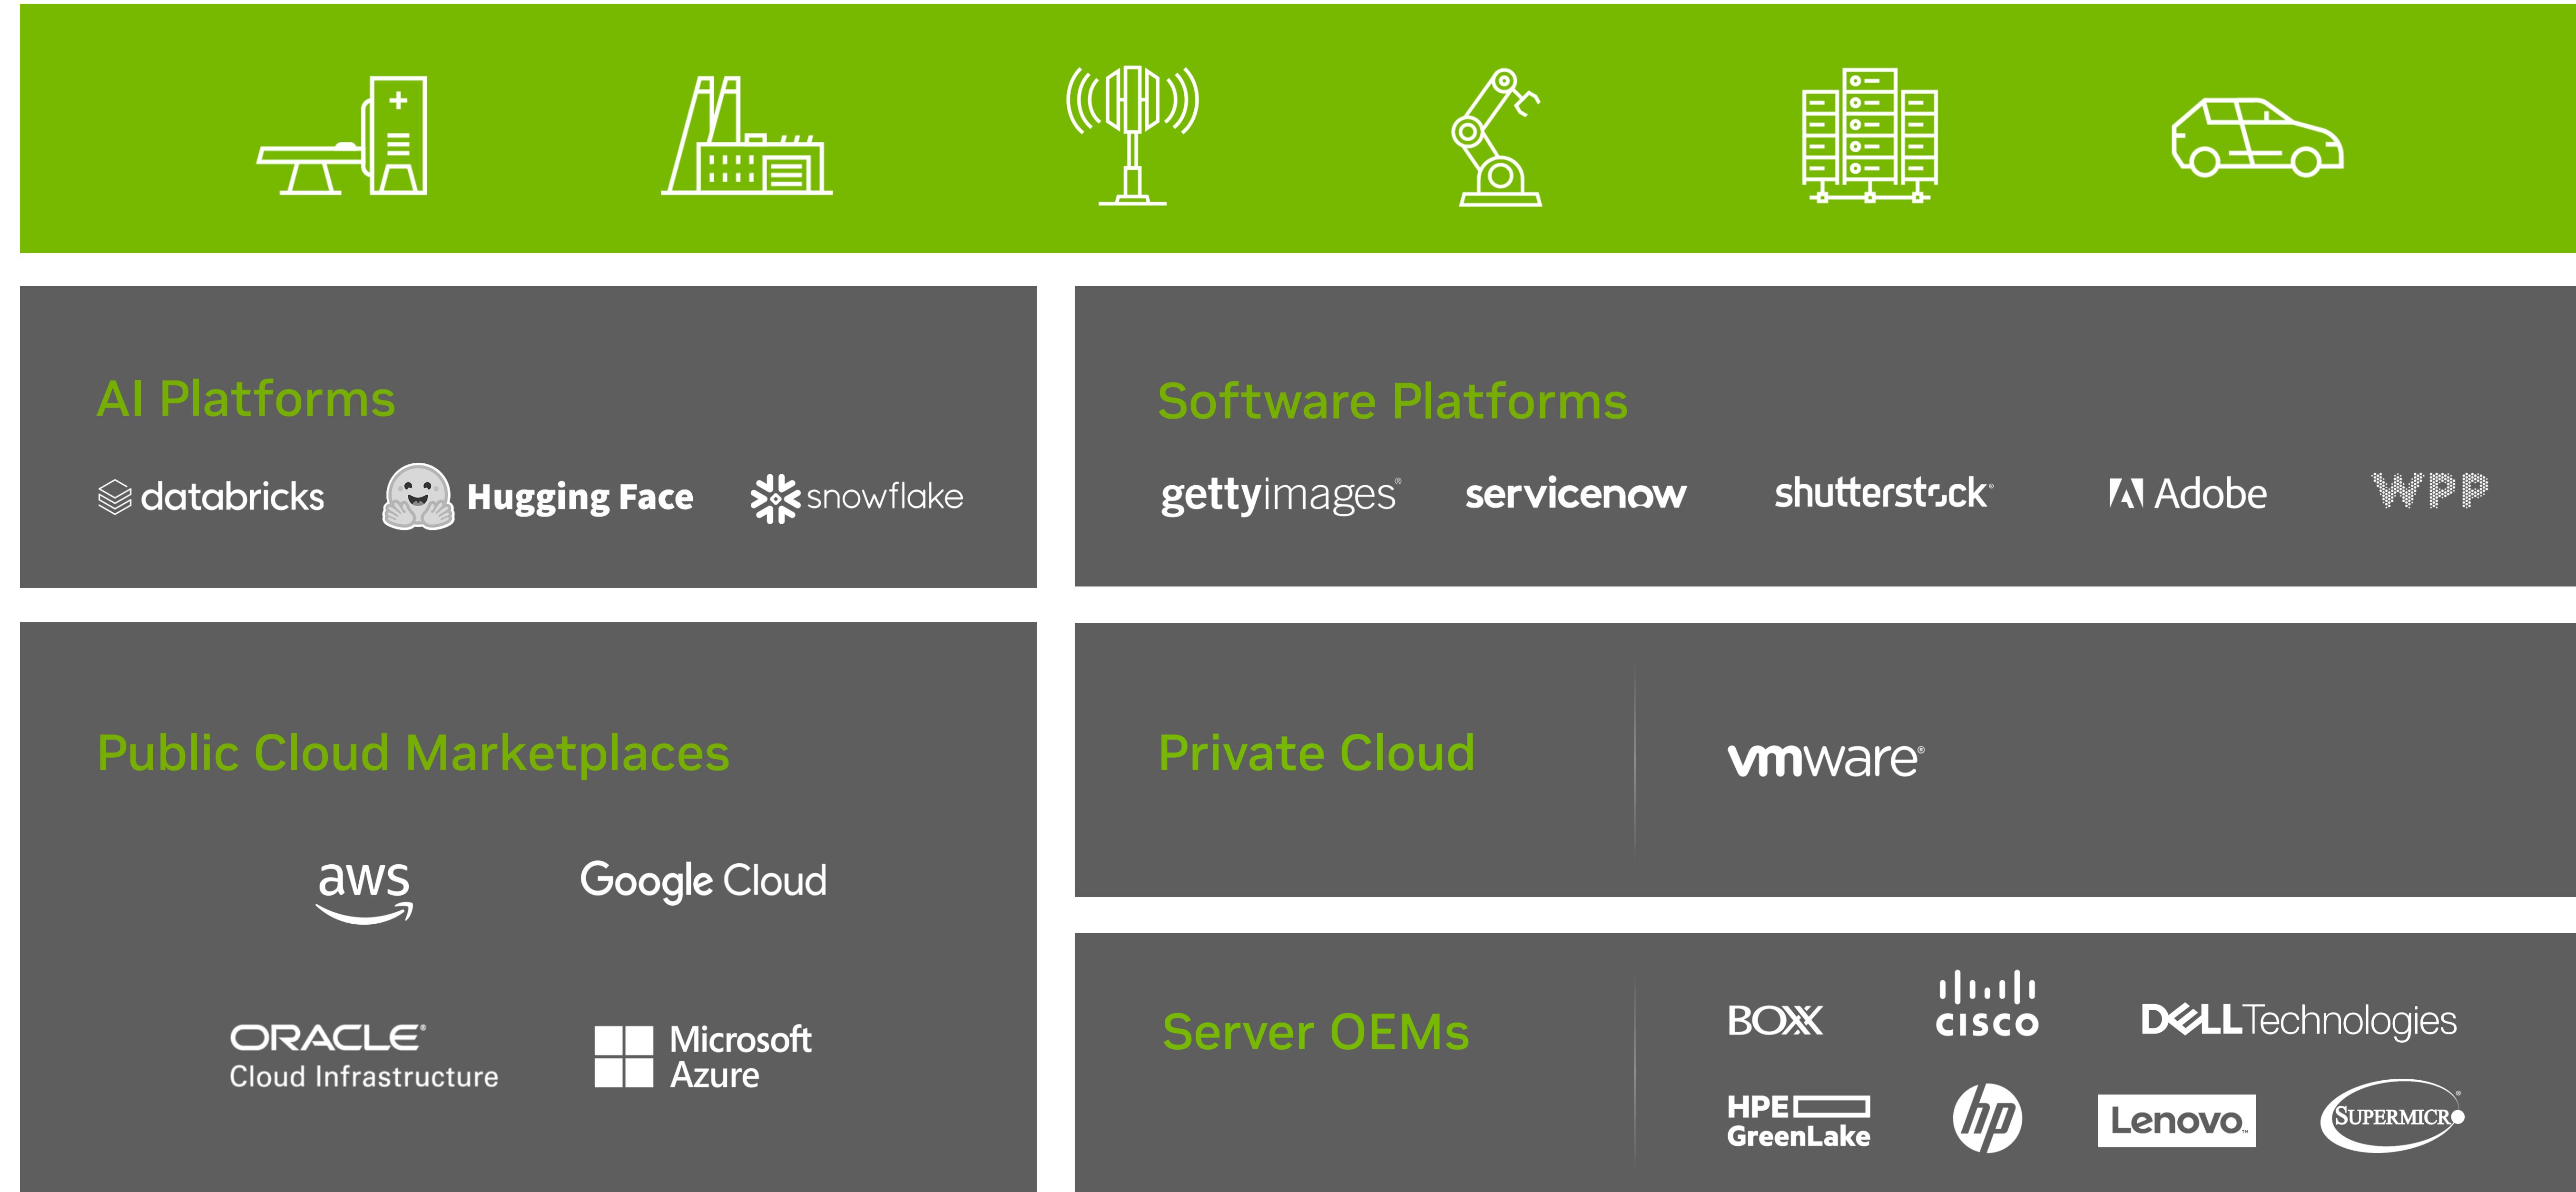
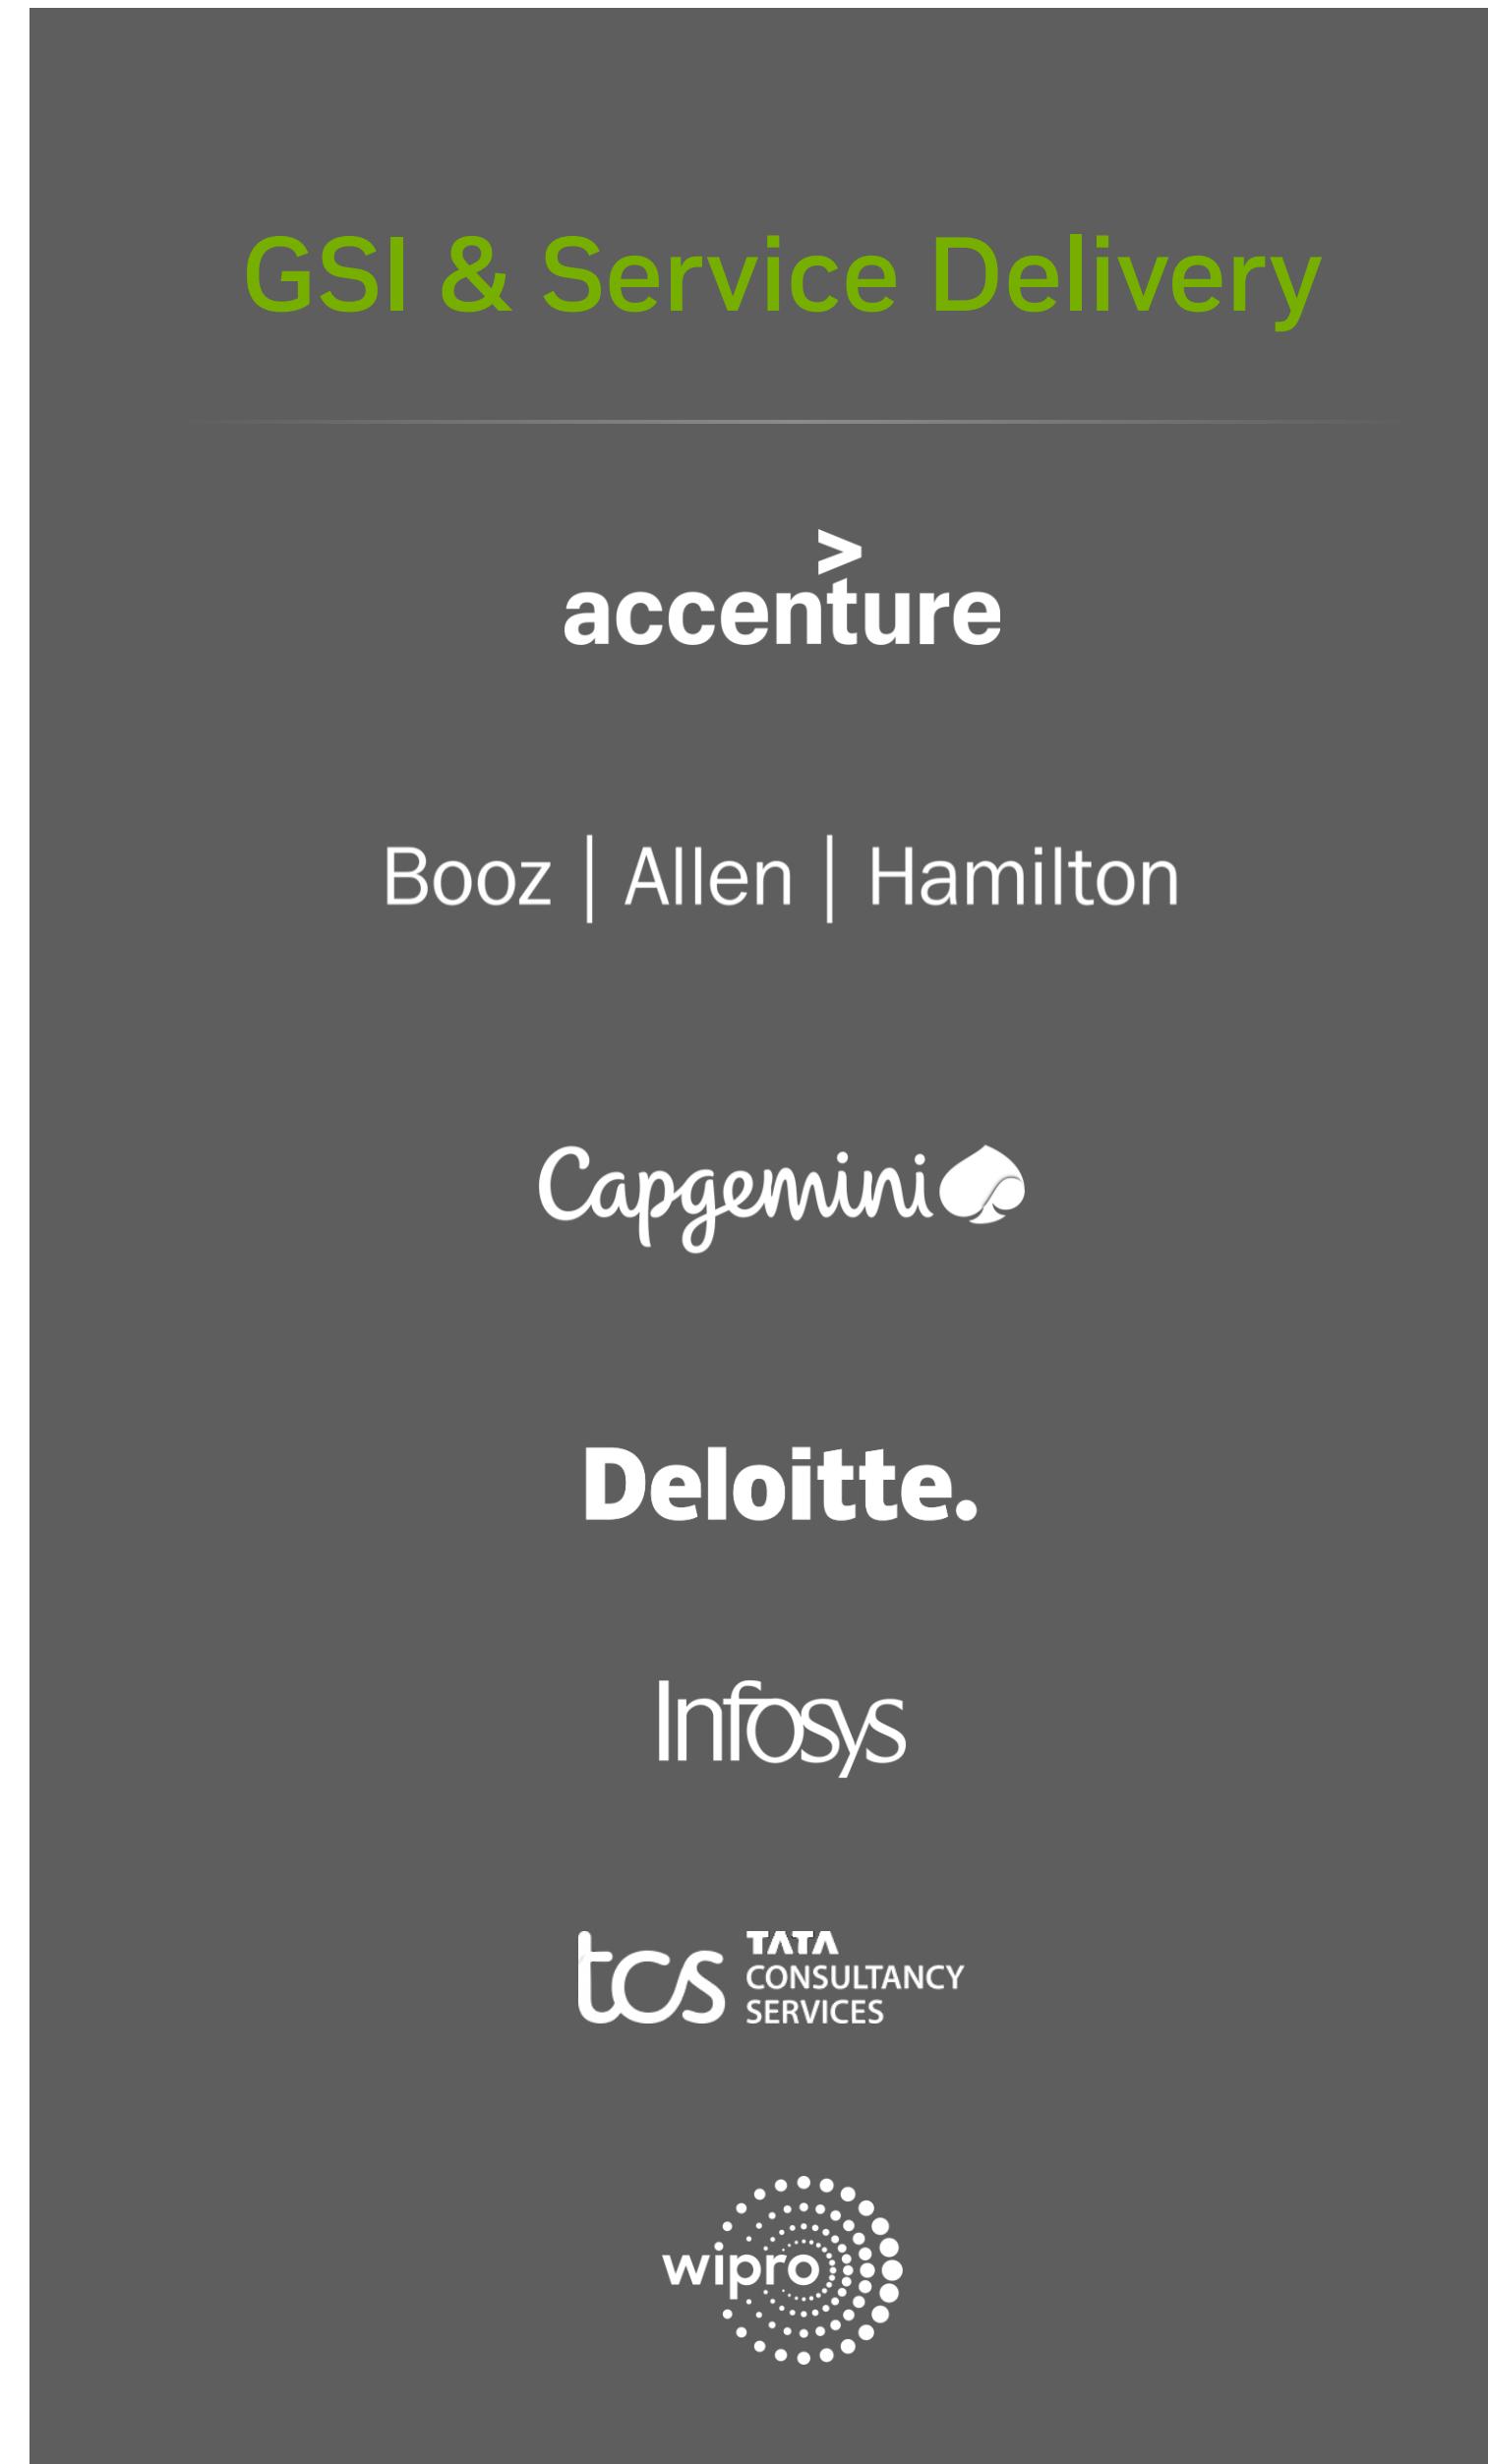
NVIDIA Certified Server
Dell | HPE | Lenovo

Consumption pricing
per GPU-hour

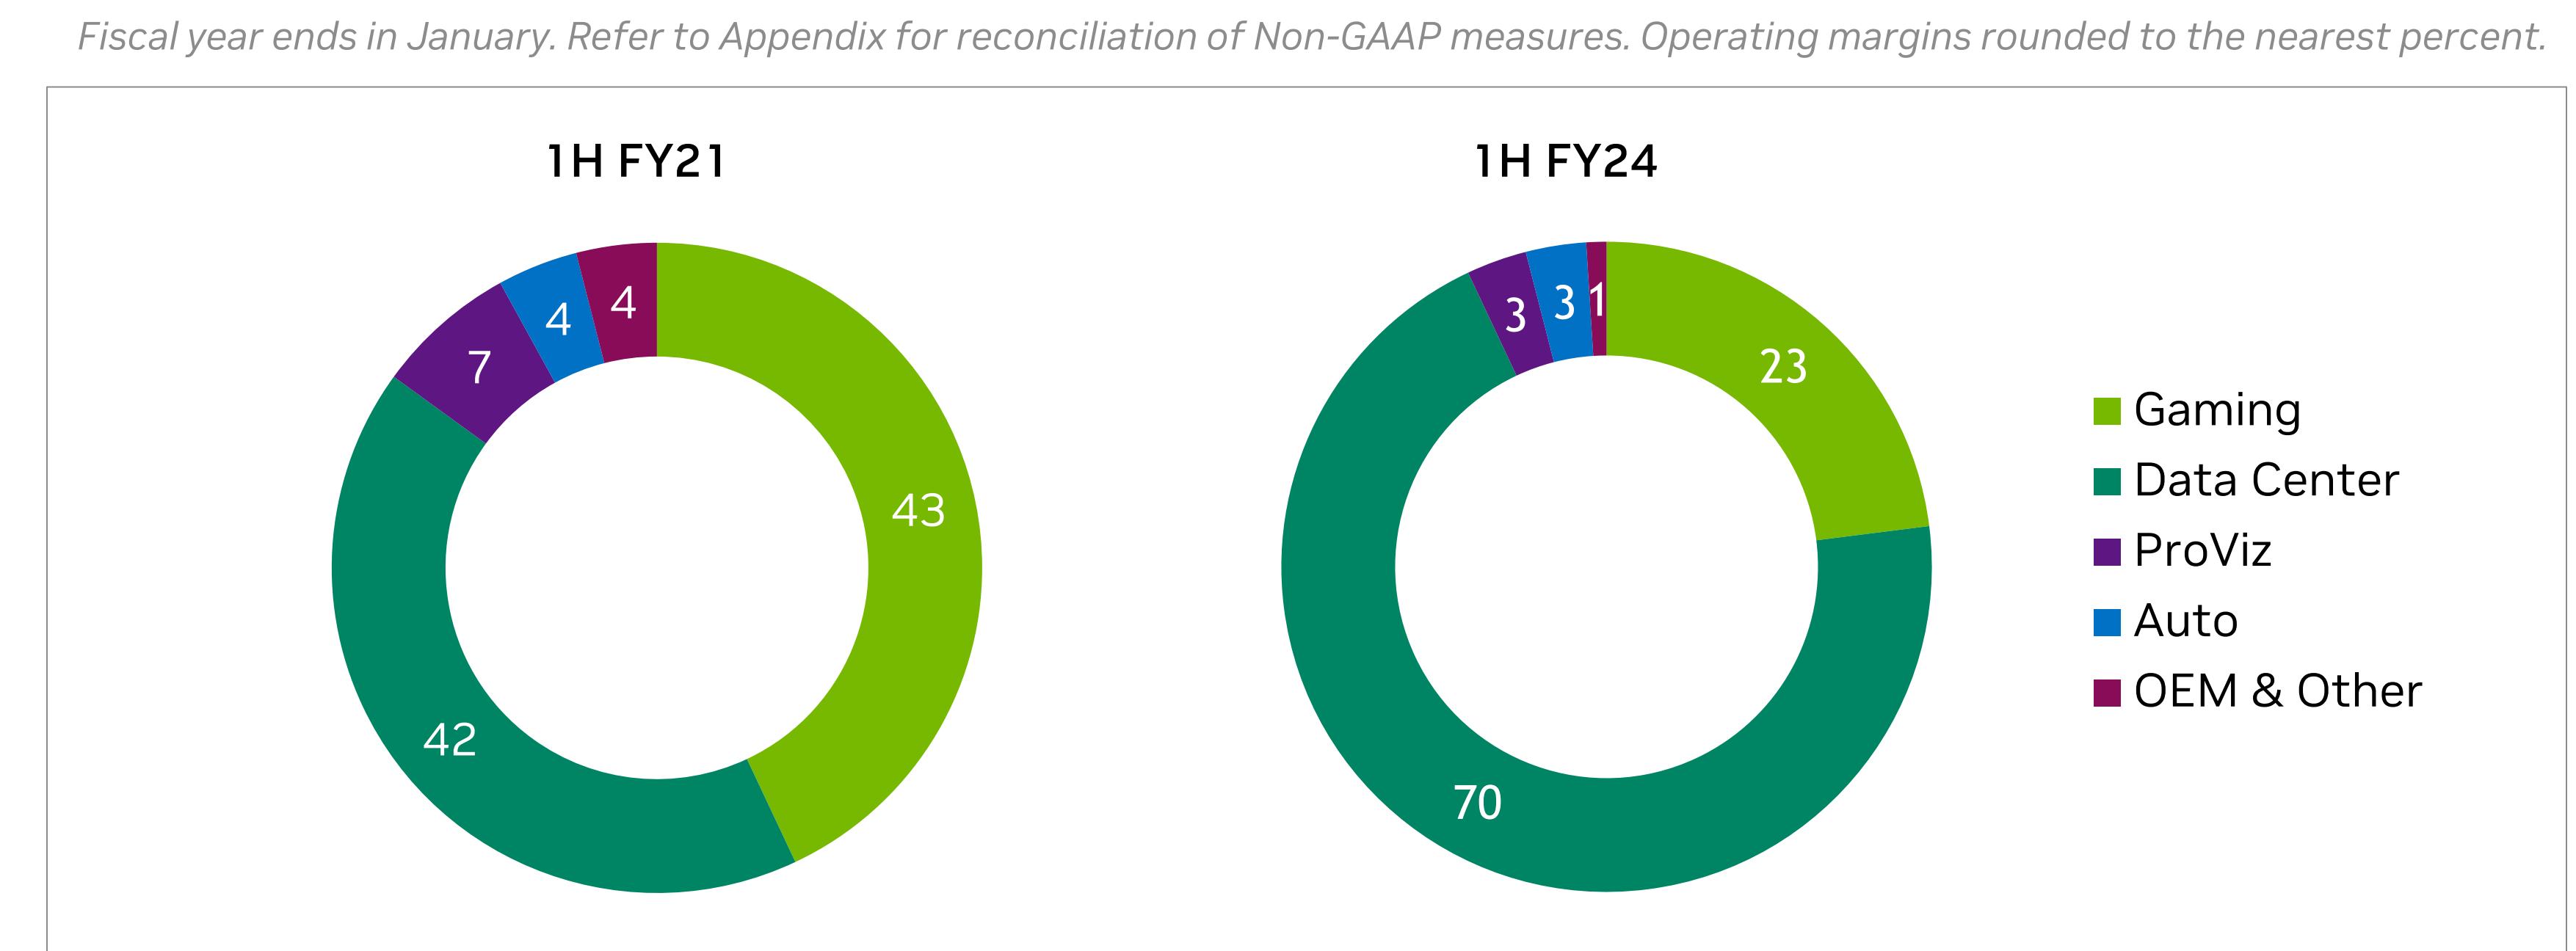
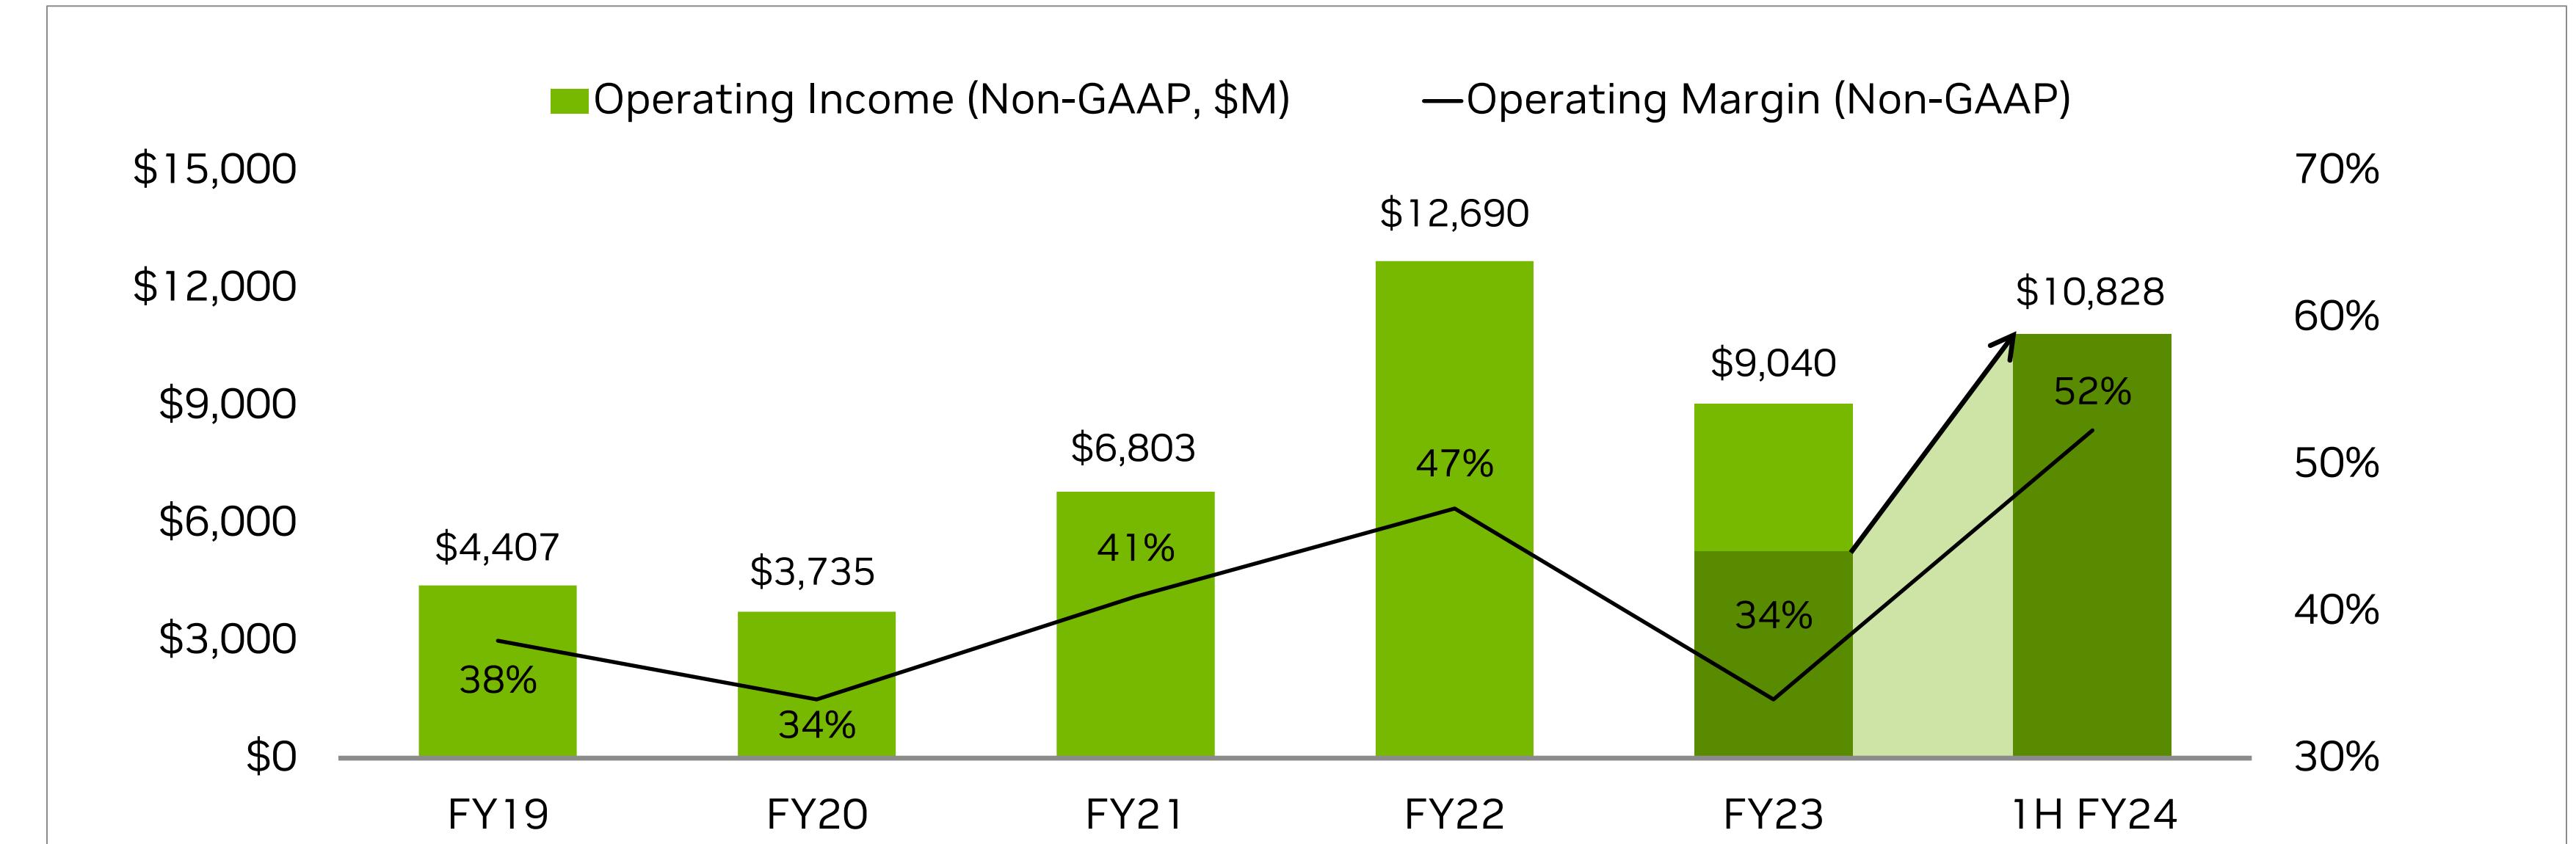
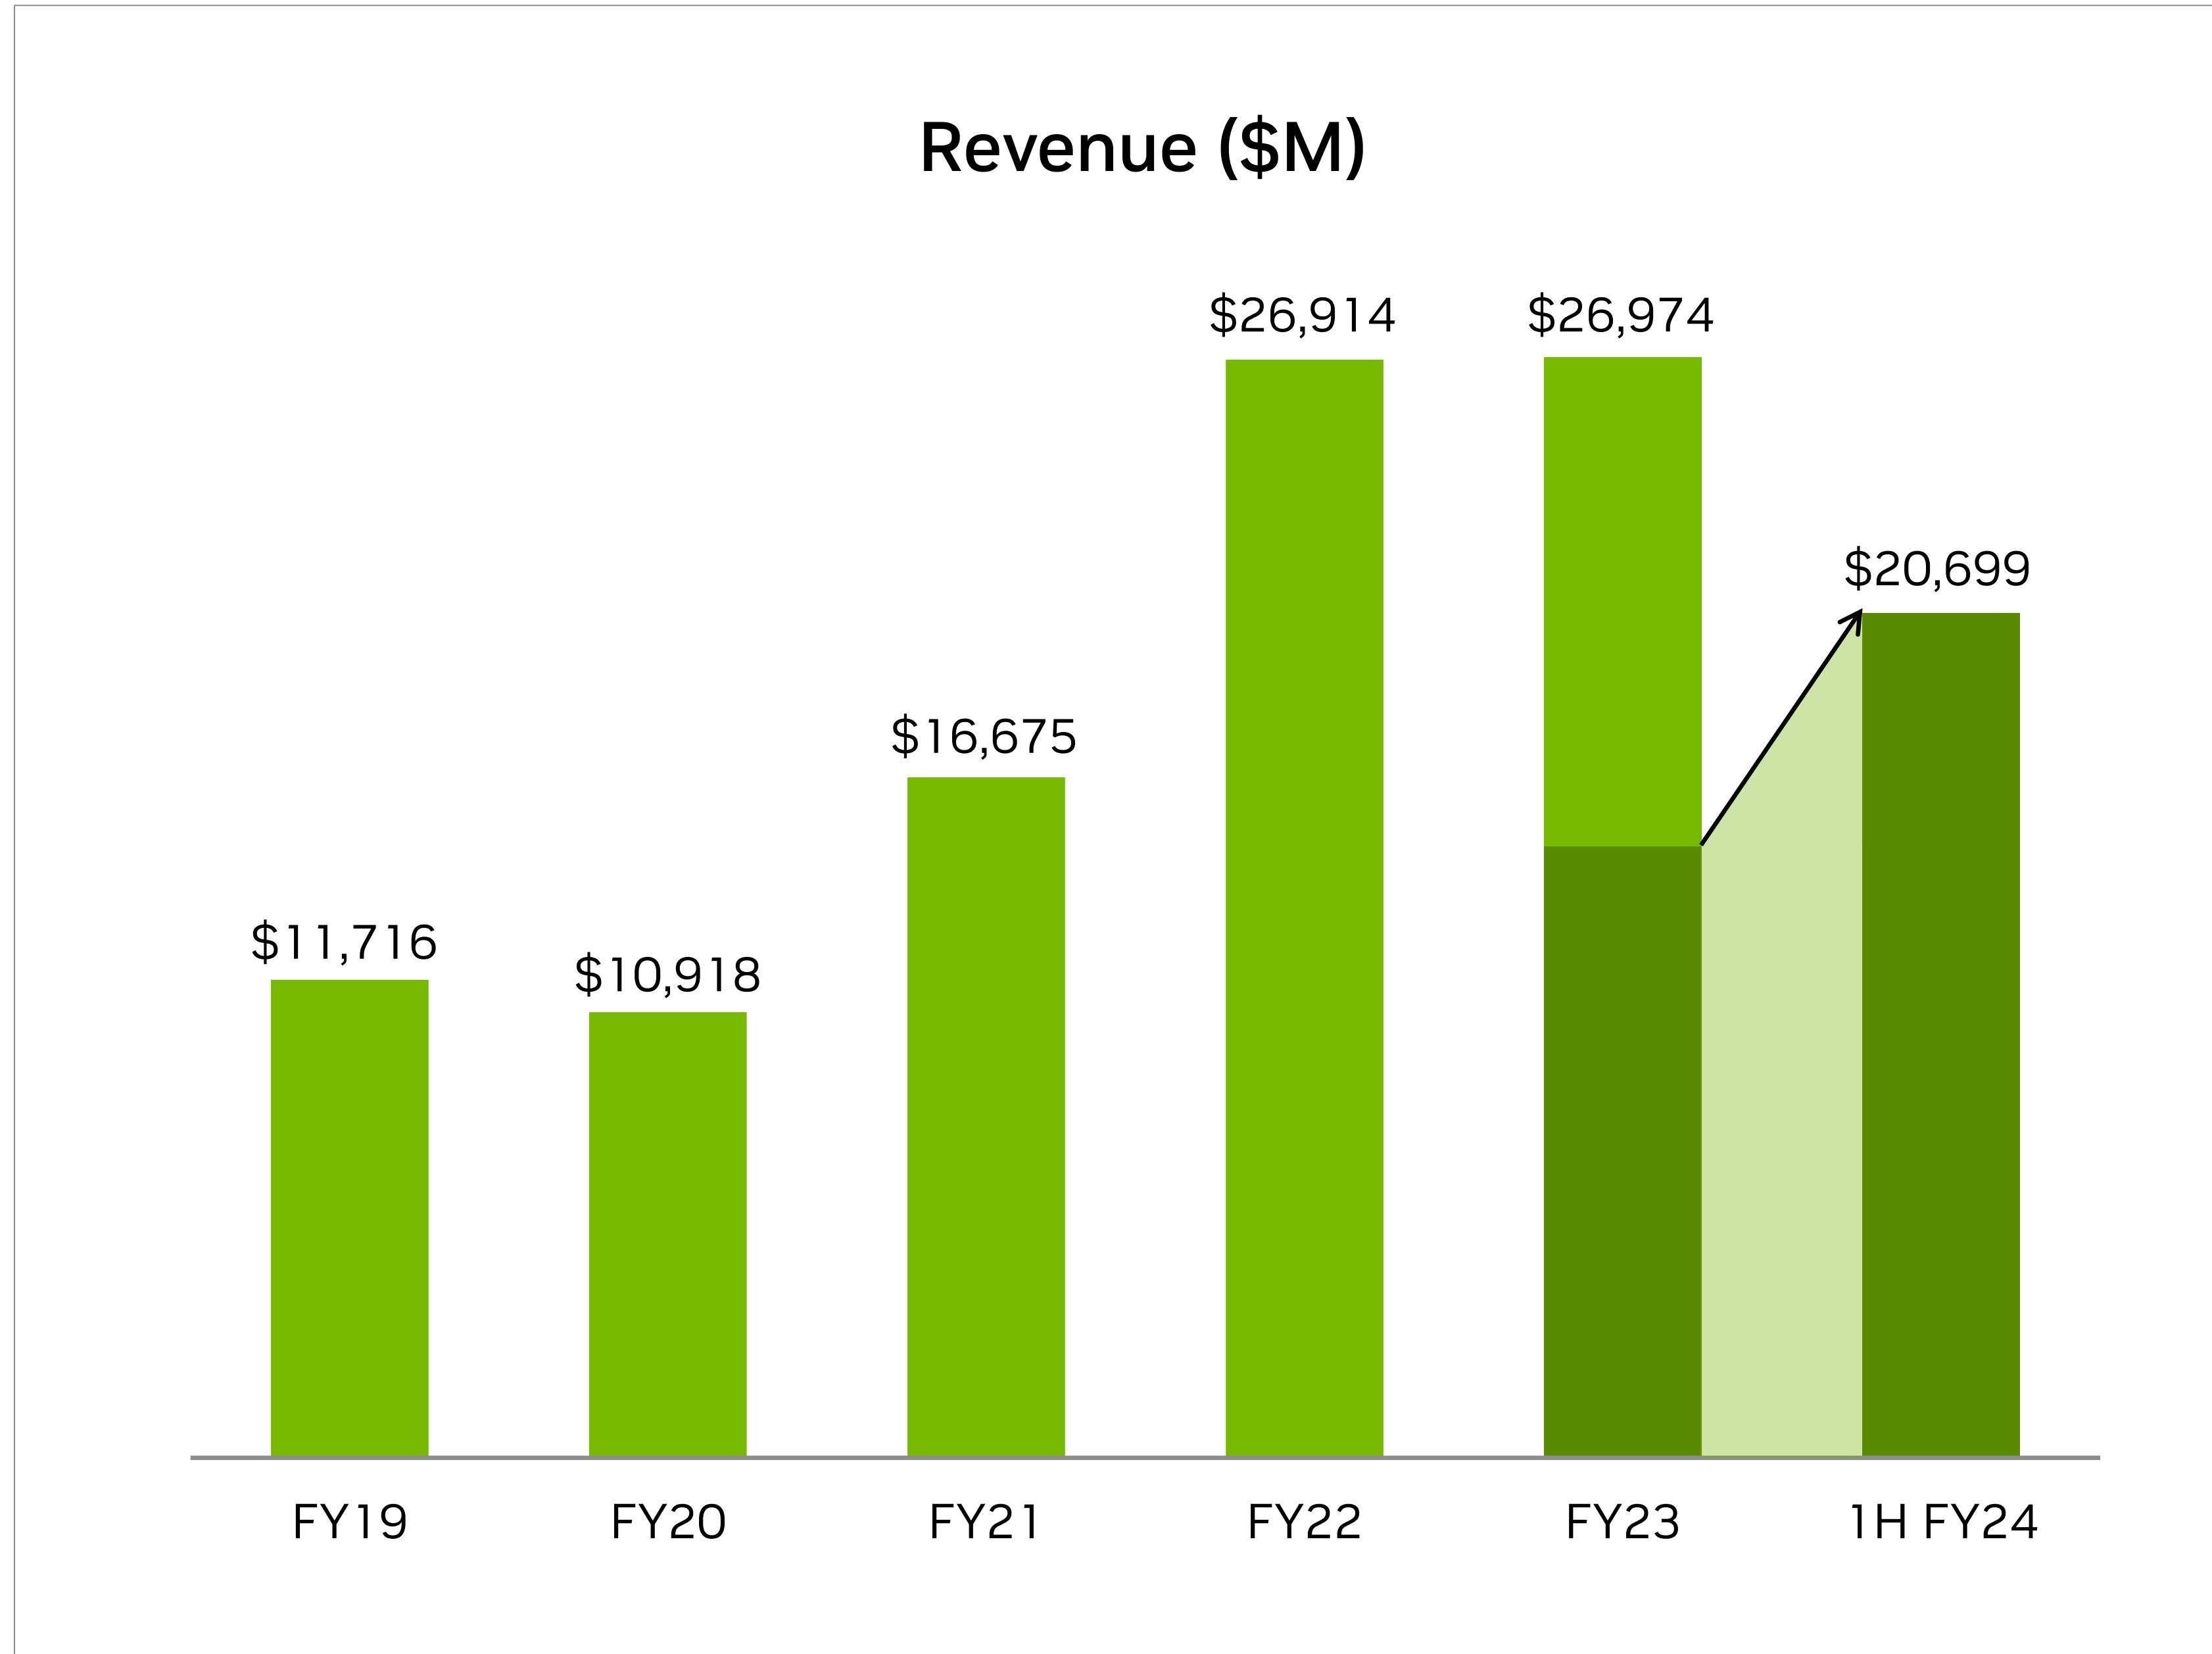
Subscription pricing
per GPU/year
(included with H100 PCIe/DGX)

NVIDIA AI Enterprise

Broad and deep ecosystem and distribution to reach every enterprise



Driving Strong & Profitable Growth



FY23 financial metrics reflect a \$2.2B charge for inventory and related reserves primarily related to Data Center and Gaming.

NVIDIA Gross Margins Reflect Value of Acceleration

Accelerated computing requires full-stack and data center-scale innovation across silicon, systems, algorithms and applications.

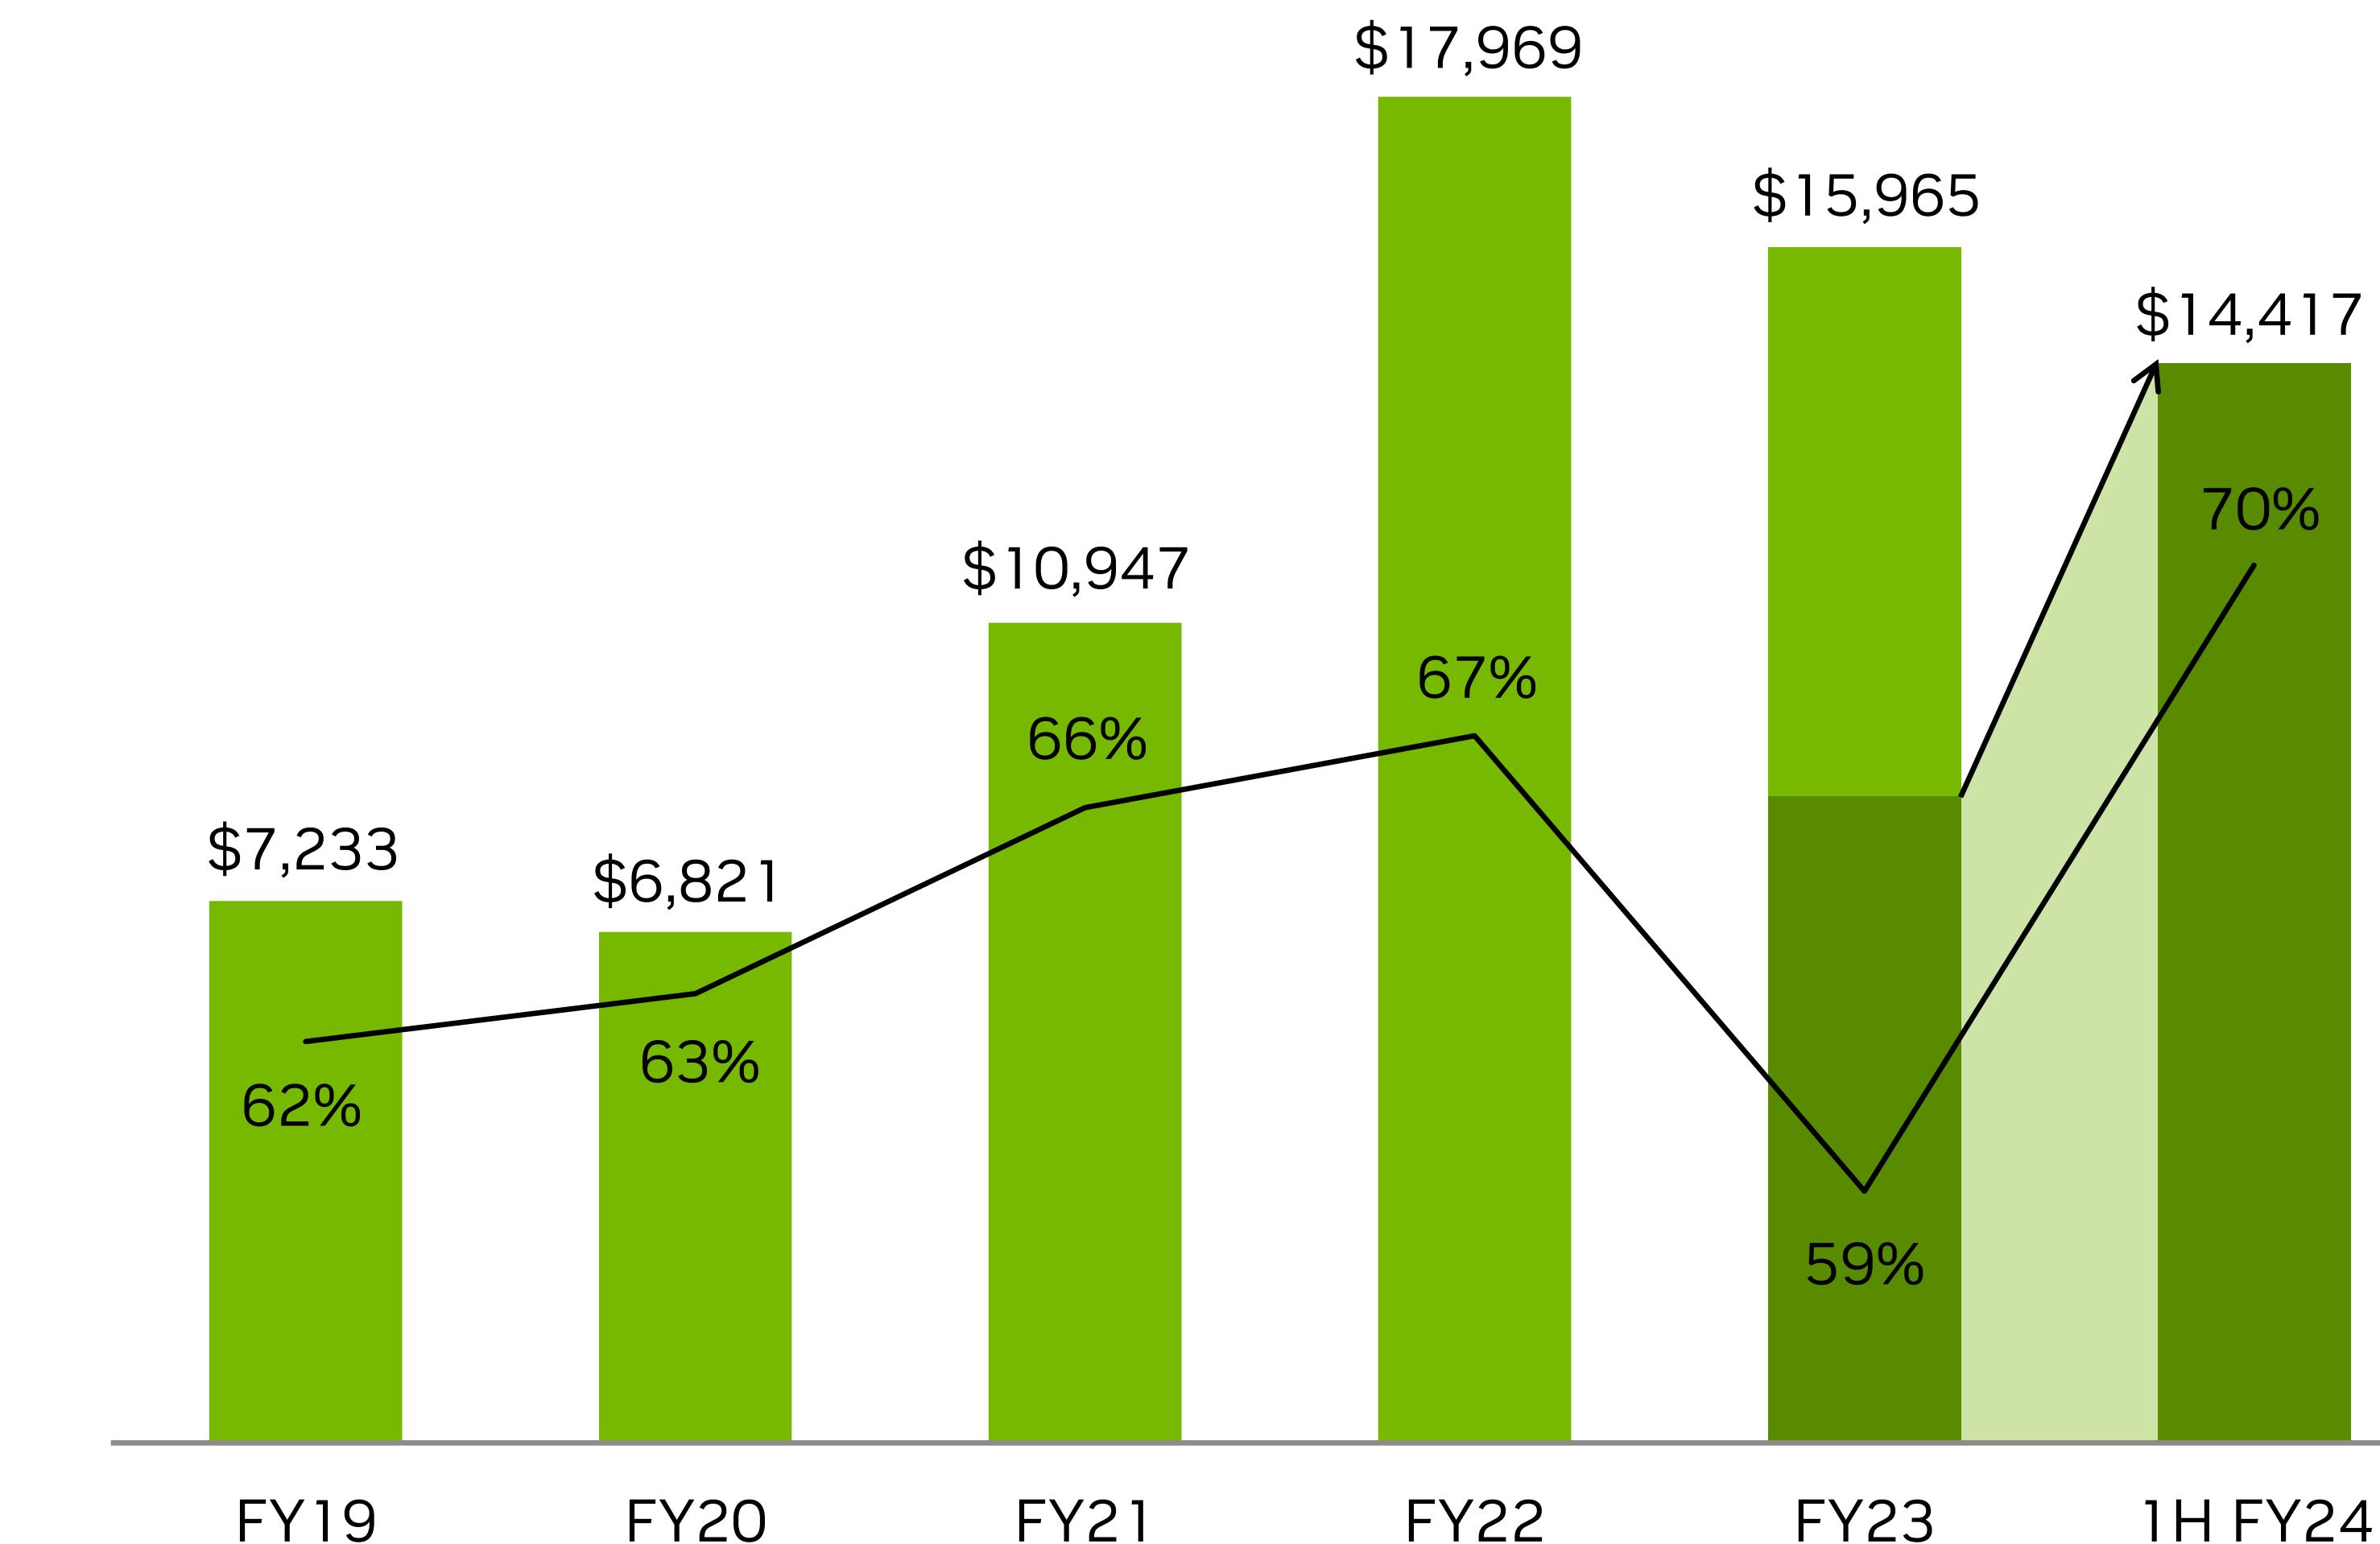
Significant expertise and effort are required, but application speed-ups can be incredible, resulting in dramatic cost and time-to-solution savings.

For example, 2 NVIDIA HGX nodes with 16 NVIDIA H100 GPUs that cost \$400K can replace 960 nodes of CPU servers that cost \$10M for the same LLM workload.

NVIDIA chips carry the value of the full-stack, not just the chip.

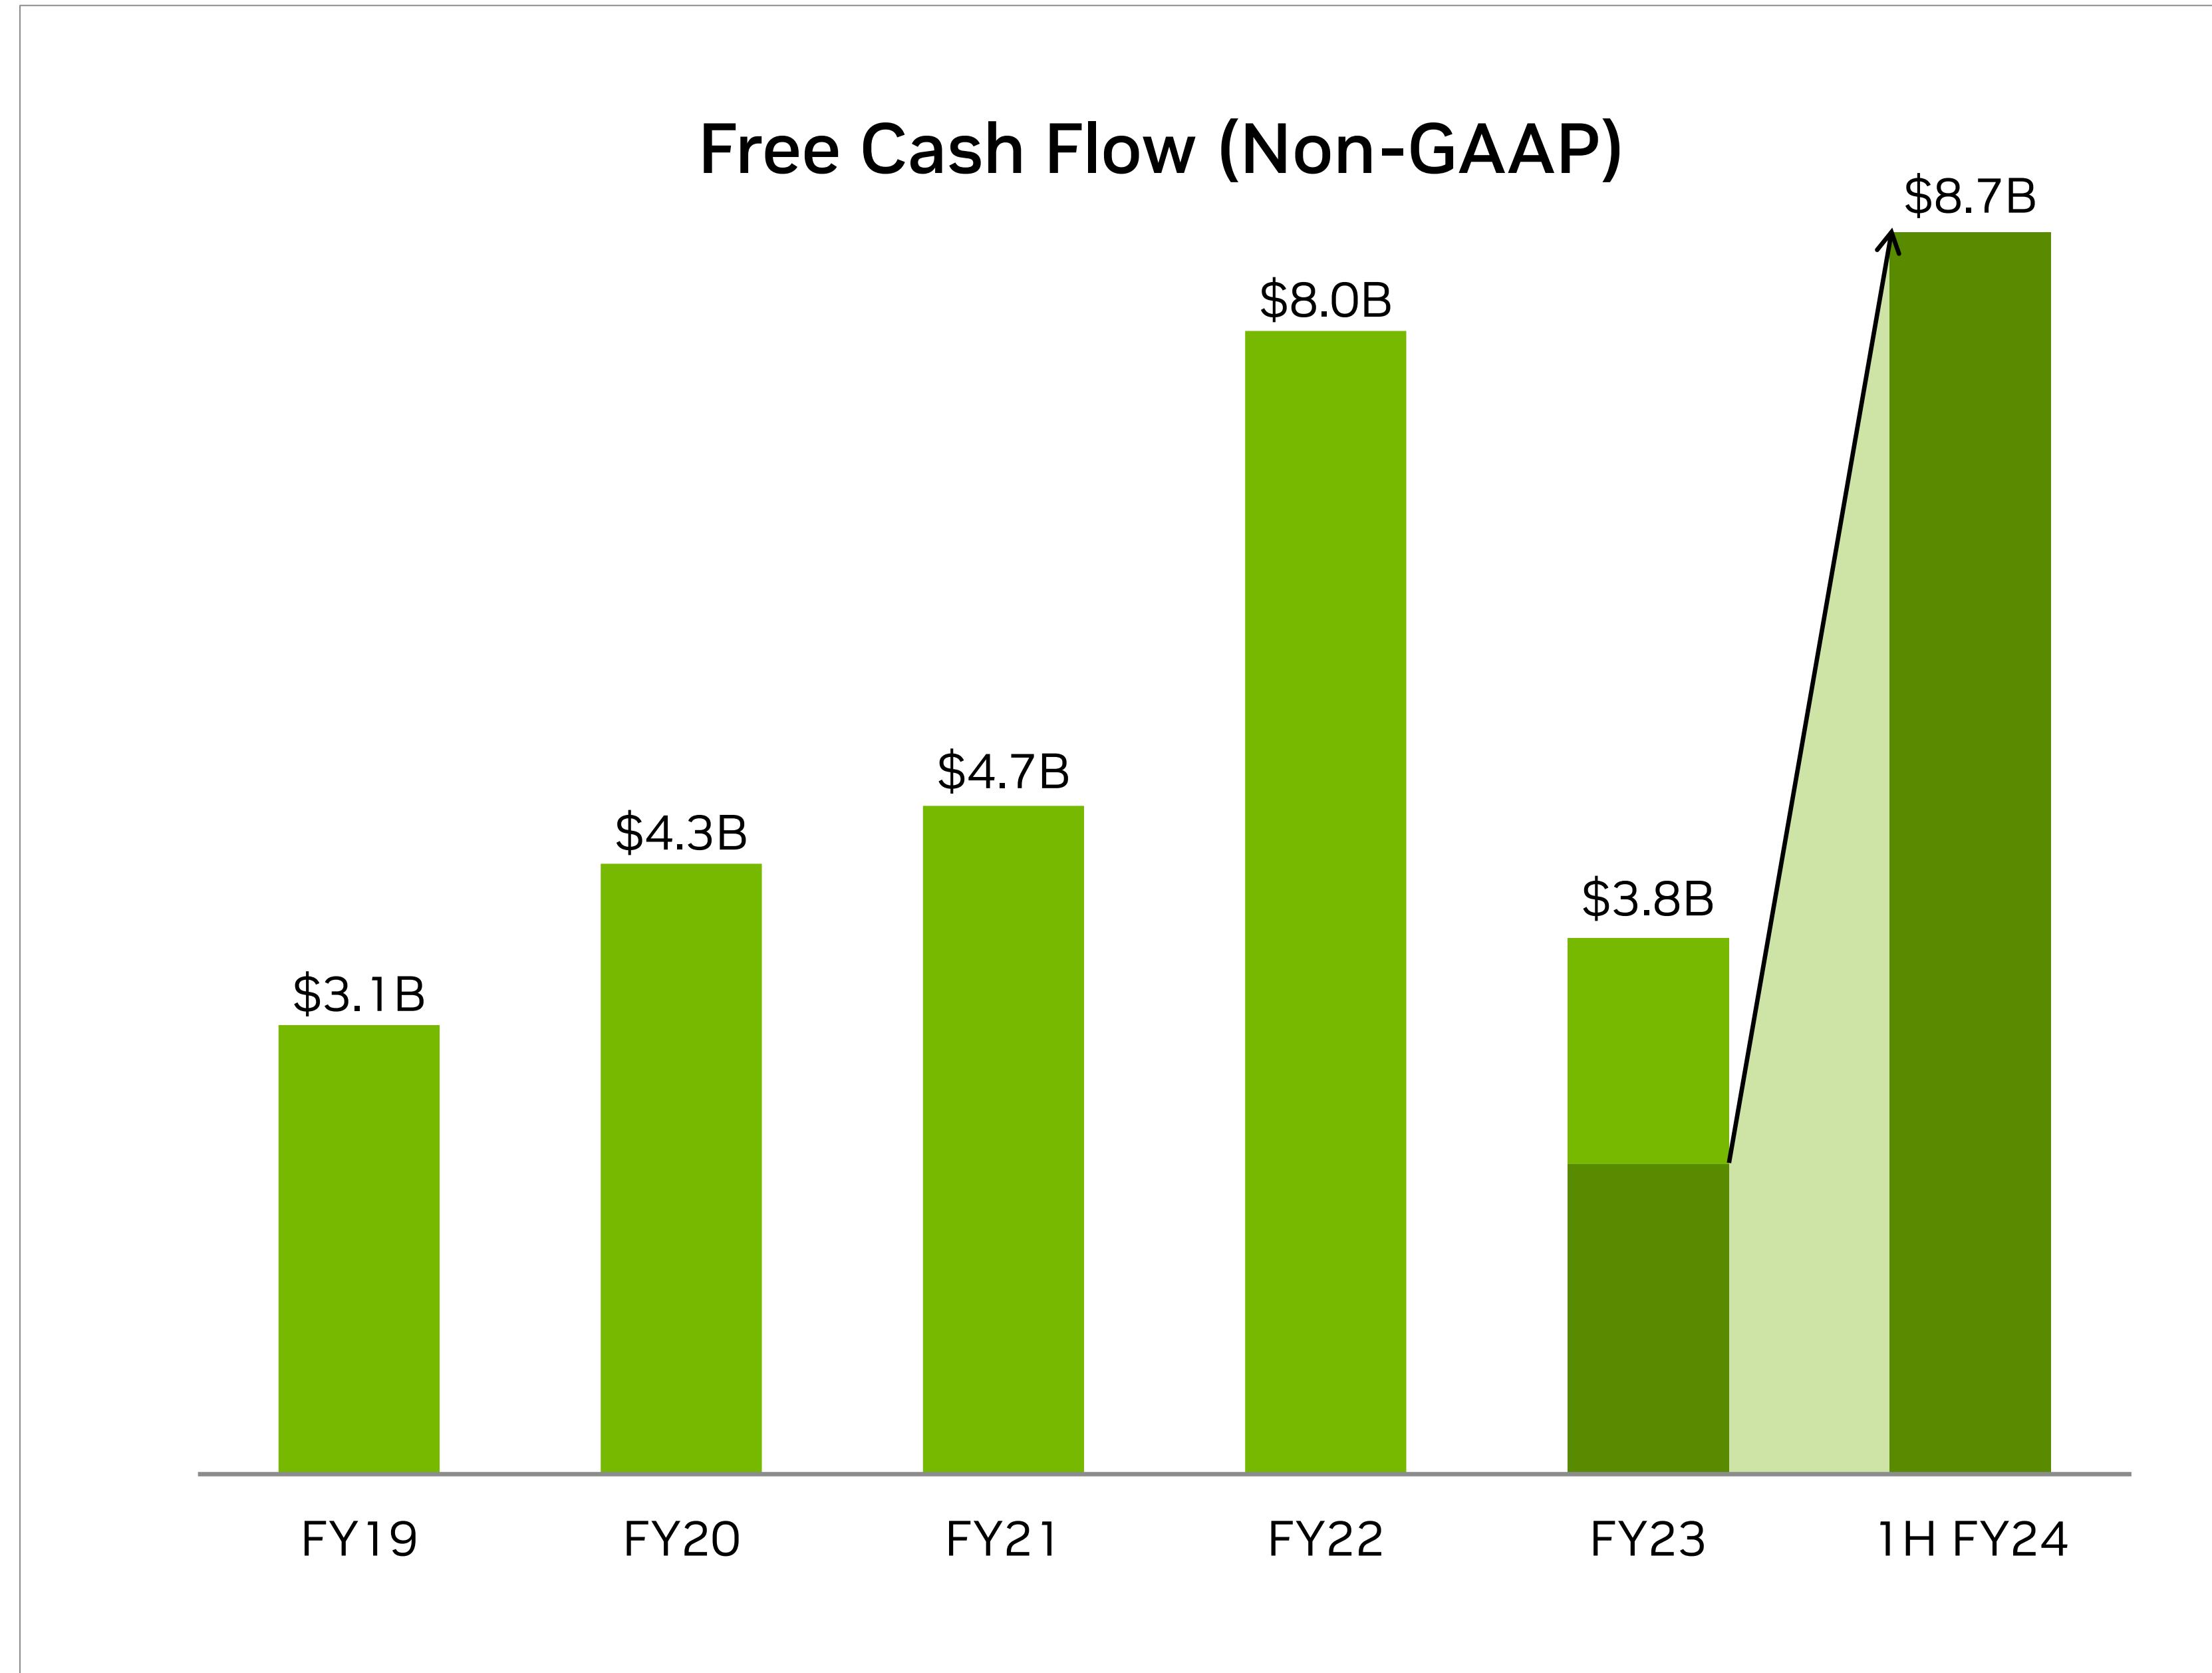
Cost comparison example based on latest available NVIDIA A100 GPU and Intel CPU inference results in the commercially available category of the MLPerf industry benchmark; includes related infrastructure costs such as networking.

■ Gross Profit (Non-GAAP, \$M) — Gross Margin (Non-GAAP)



FY23 financial metrics reflect a \$2.2B charge for inventory and related reserves primarily related to Data Center and Gaming. Fiscal year ends in January. Refer to Appendix for reconciliation of Non-GAAP measures. Gross margins are rounded to the nearest percent.

Strong Cash Flow Generation



Fiscal year ends in January. Refer to Appendix for reconciliation of Non-GAAP measures.

¹ Subject to continuing determination by our Board of Directors.

Capital Allocation

Share Repurchase

\$10B repurchased in FY23

Additional \$25B in stock repurchases authorized, adding to \$4B which remained as of end of Q2

Dividend

\$398M in FY 2023

Plan to Maintain¹

Strategic Investments

Growing Our Talent

Platform Reach & Ecosystem

Our Market Platforms at a Glance

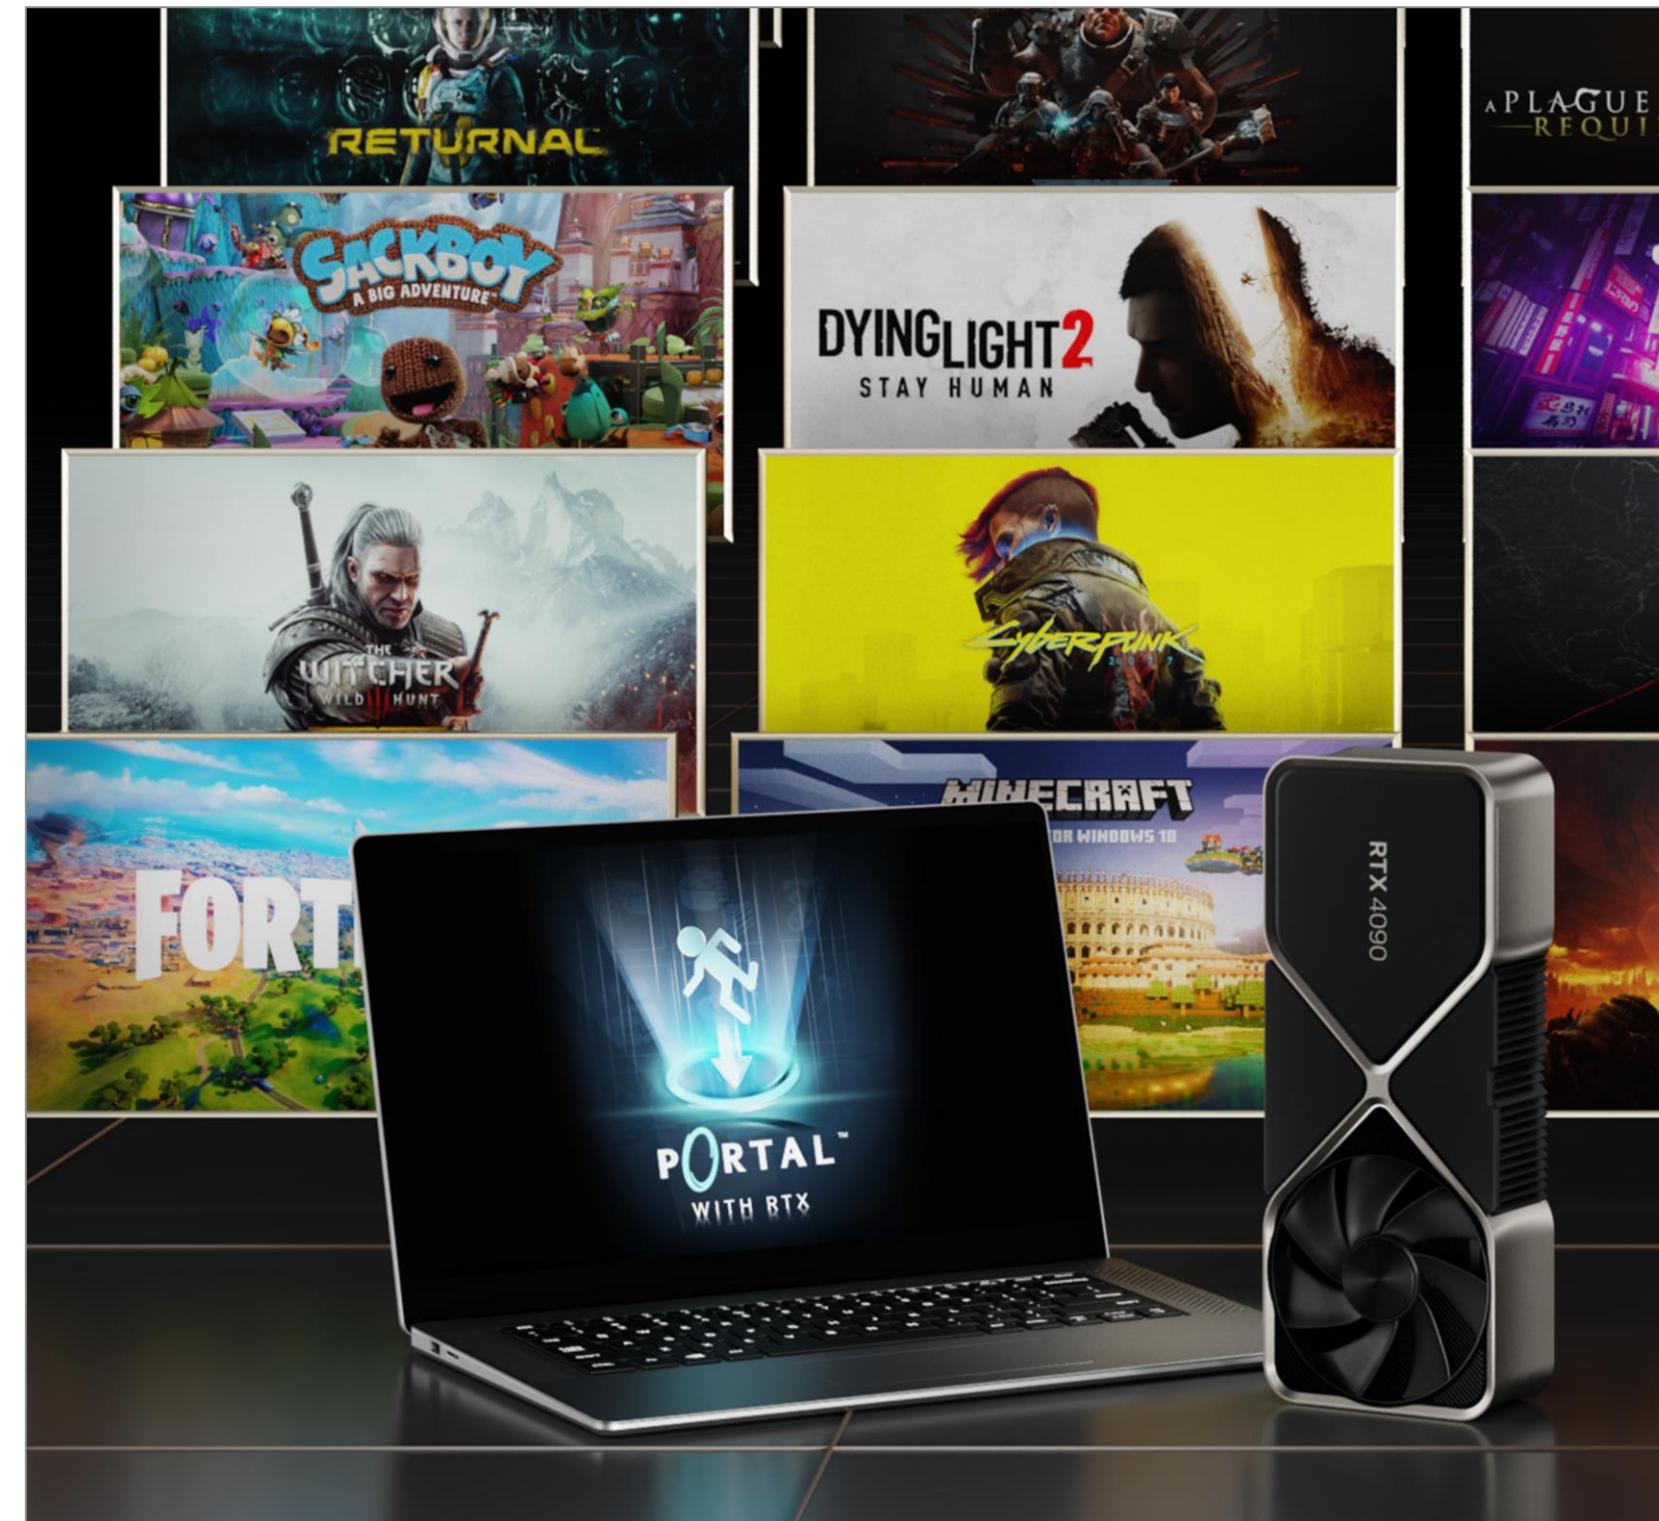


Data Center

56% of FY23 revenue

FY23 Revenue \$15.0B
5-yr CAGR 51%

DGX/HGX/MGX/IGX systems
GPU | CPU | DPU | Networking
NVIDIA AI software

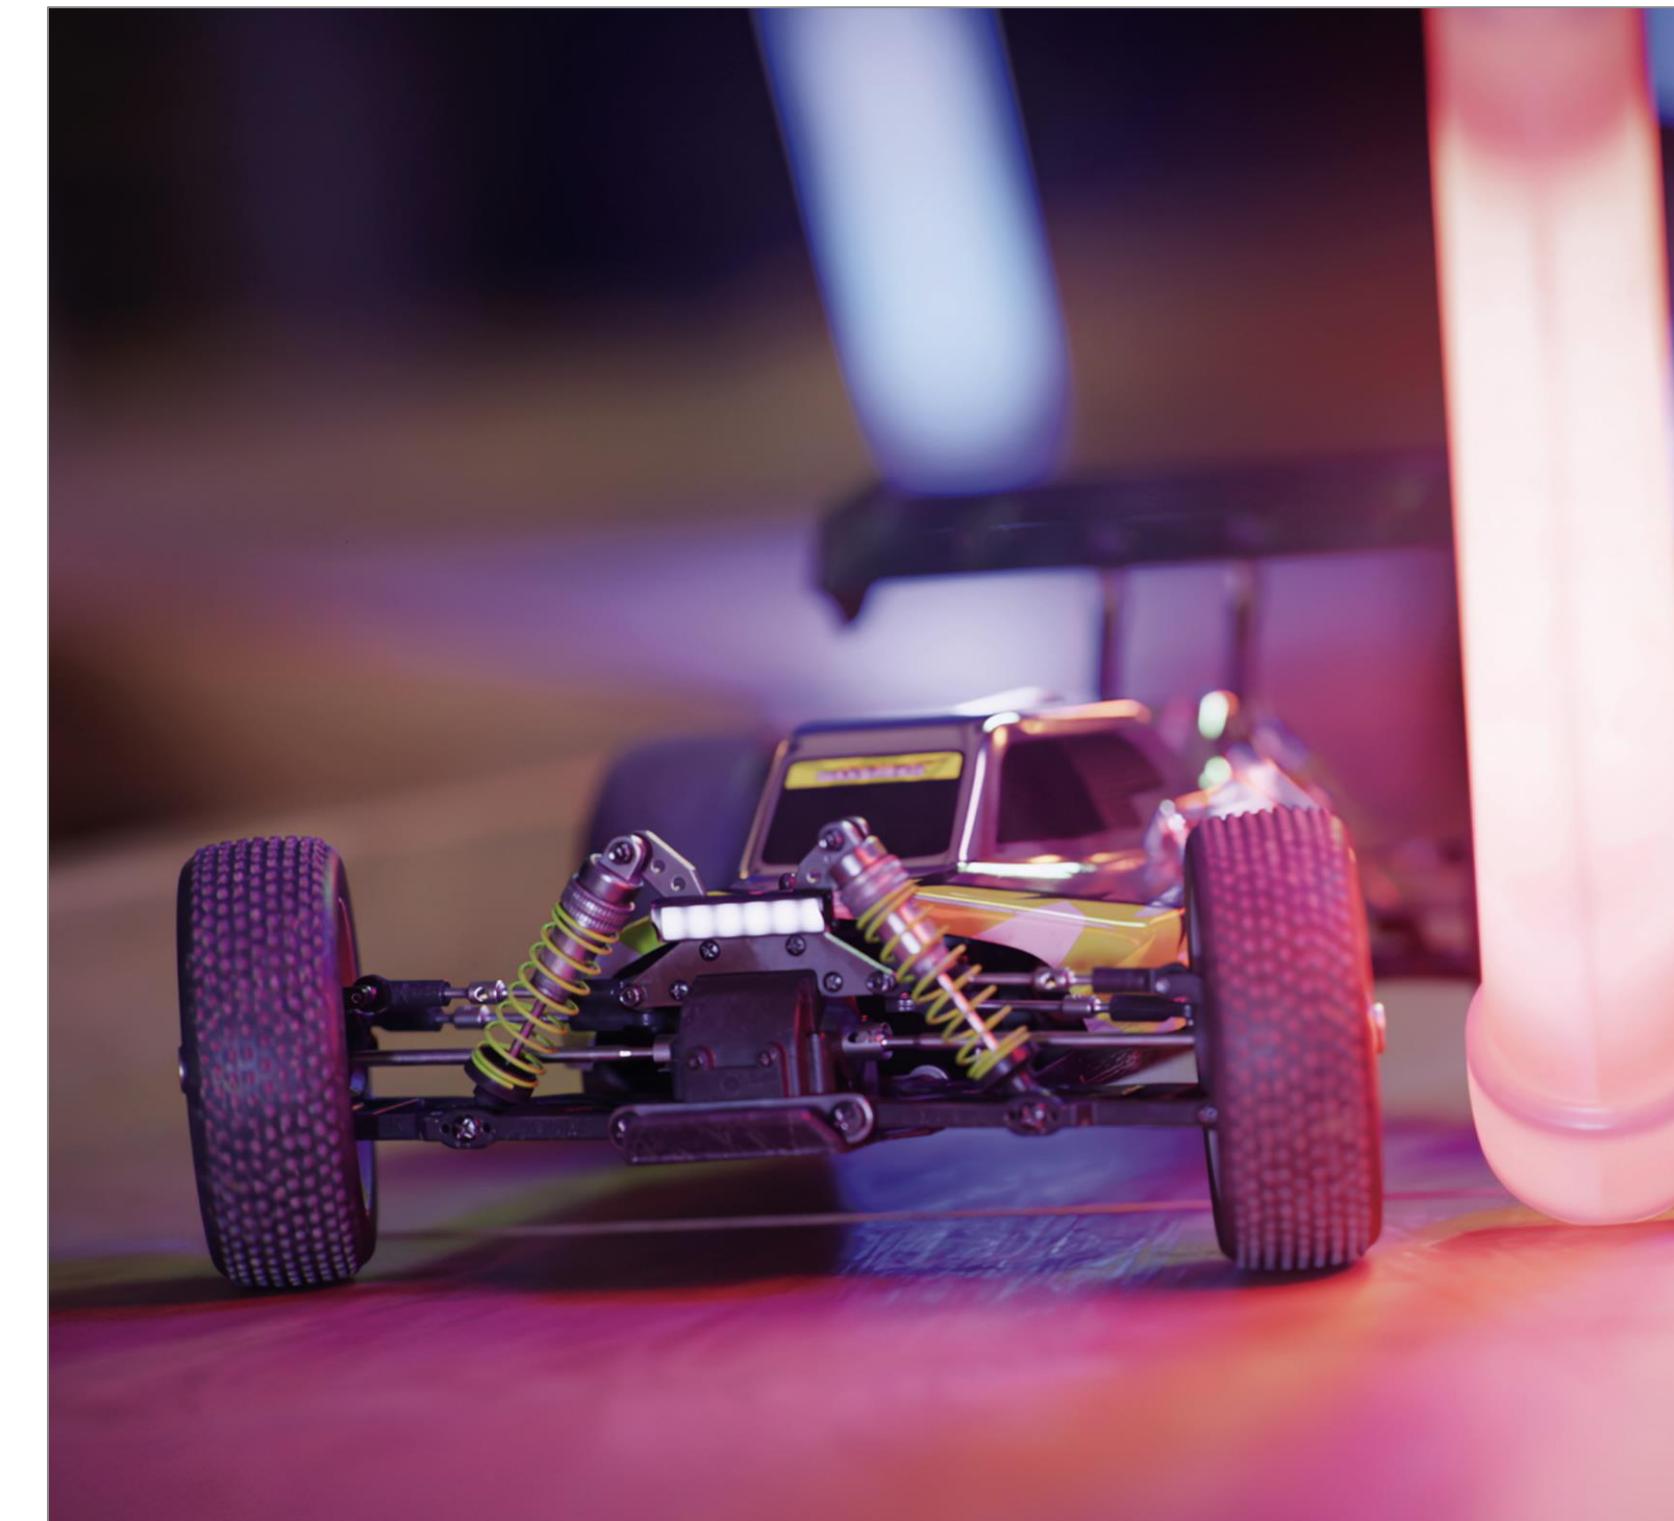


Gaming

33% of FY23 revenue

FY23 Revenue \$9.1B
5-yr CAGR 10%

GeForce GPUs for PC gaming
GeForce NOW cloud gaming



Professional Visualization

6% of FY23 revenue

FY23 Revenue \$1.5B
5-yr CAGR 11%

NVIDIA RTX GPUs
for workstations
Omniverse software



Automotive

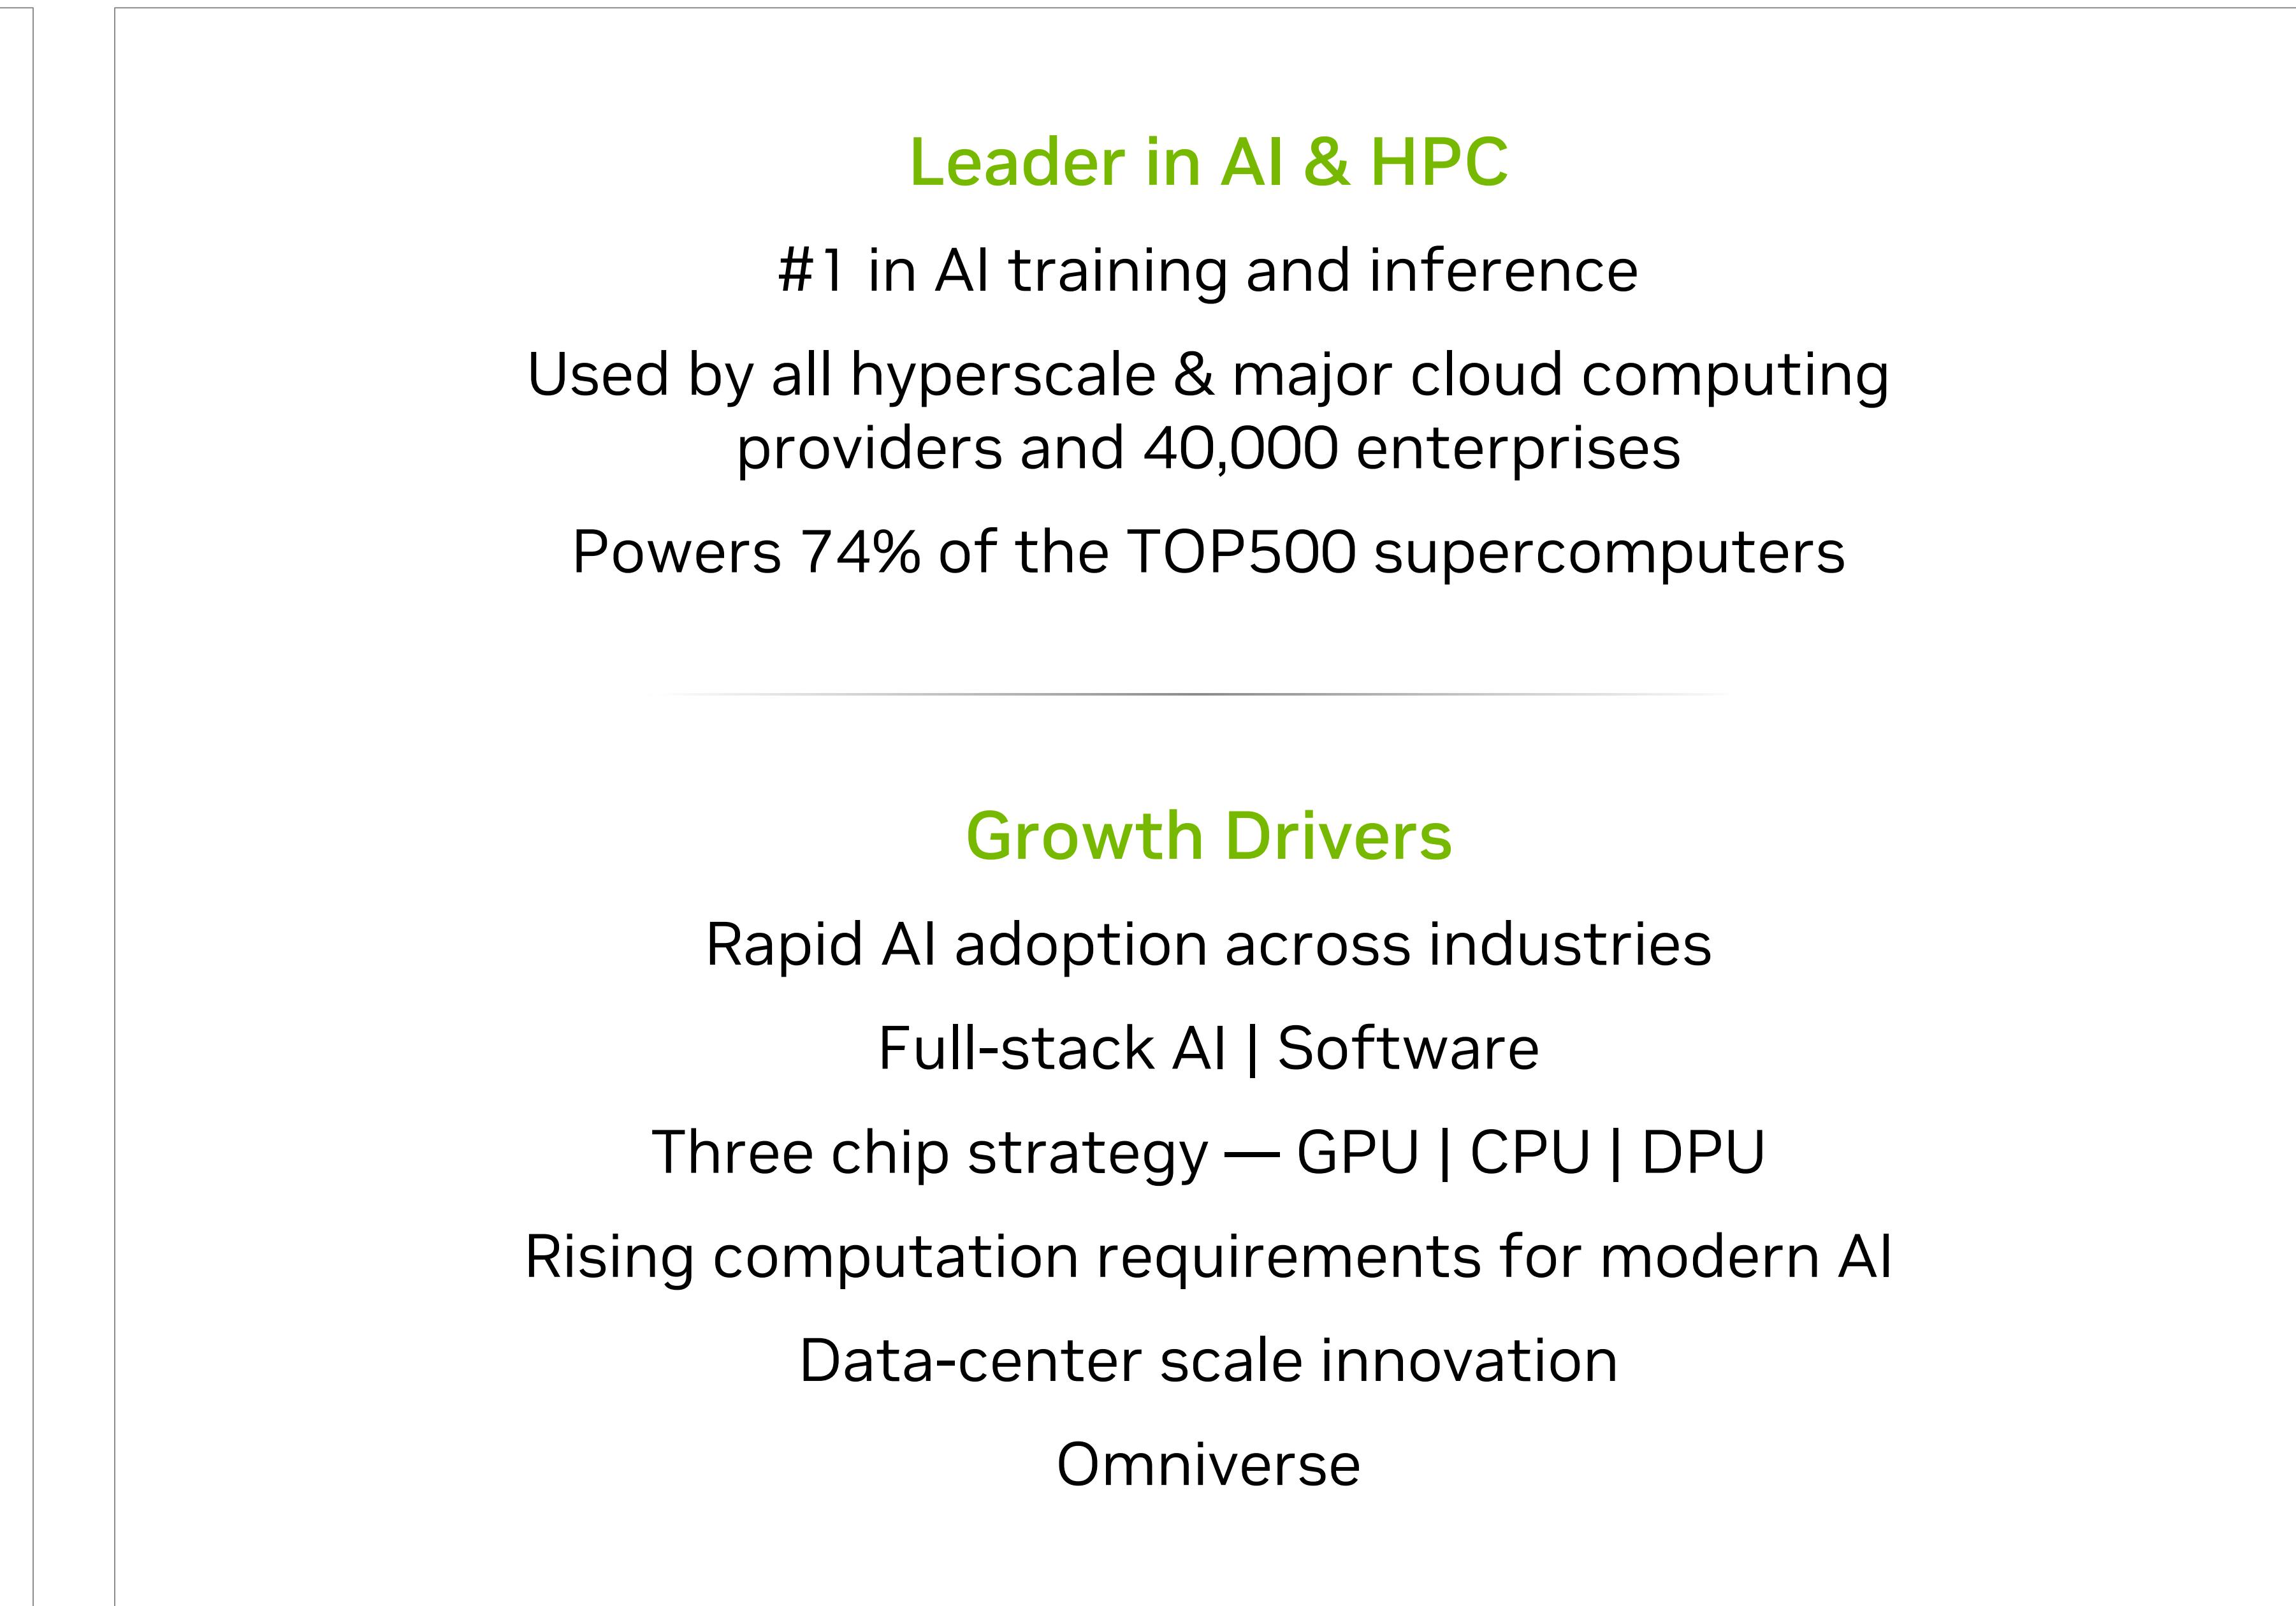
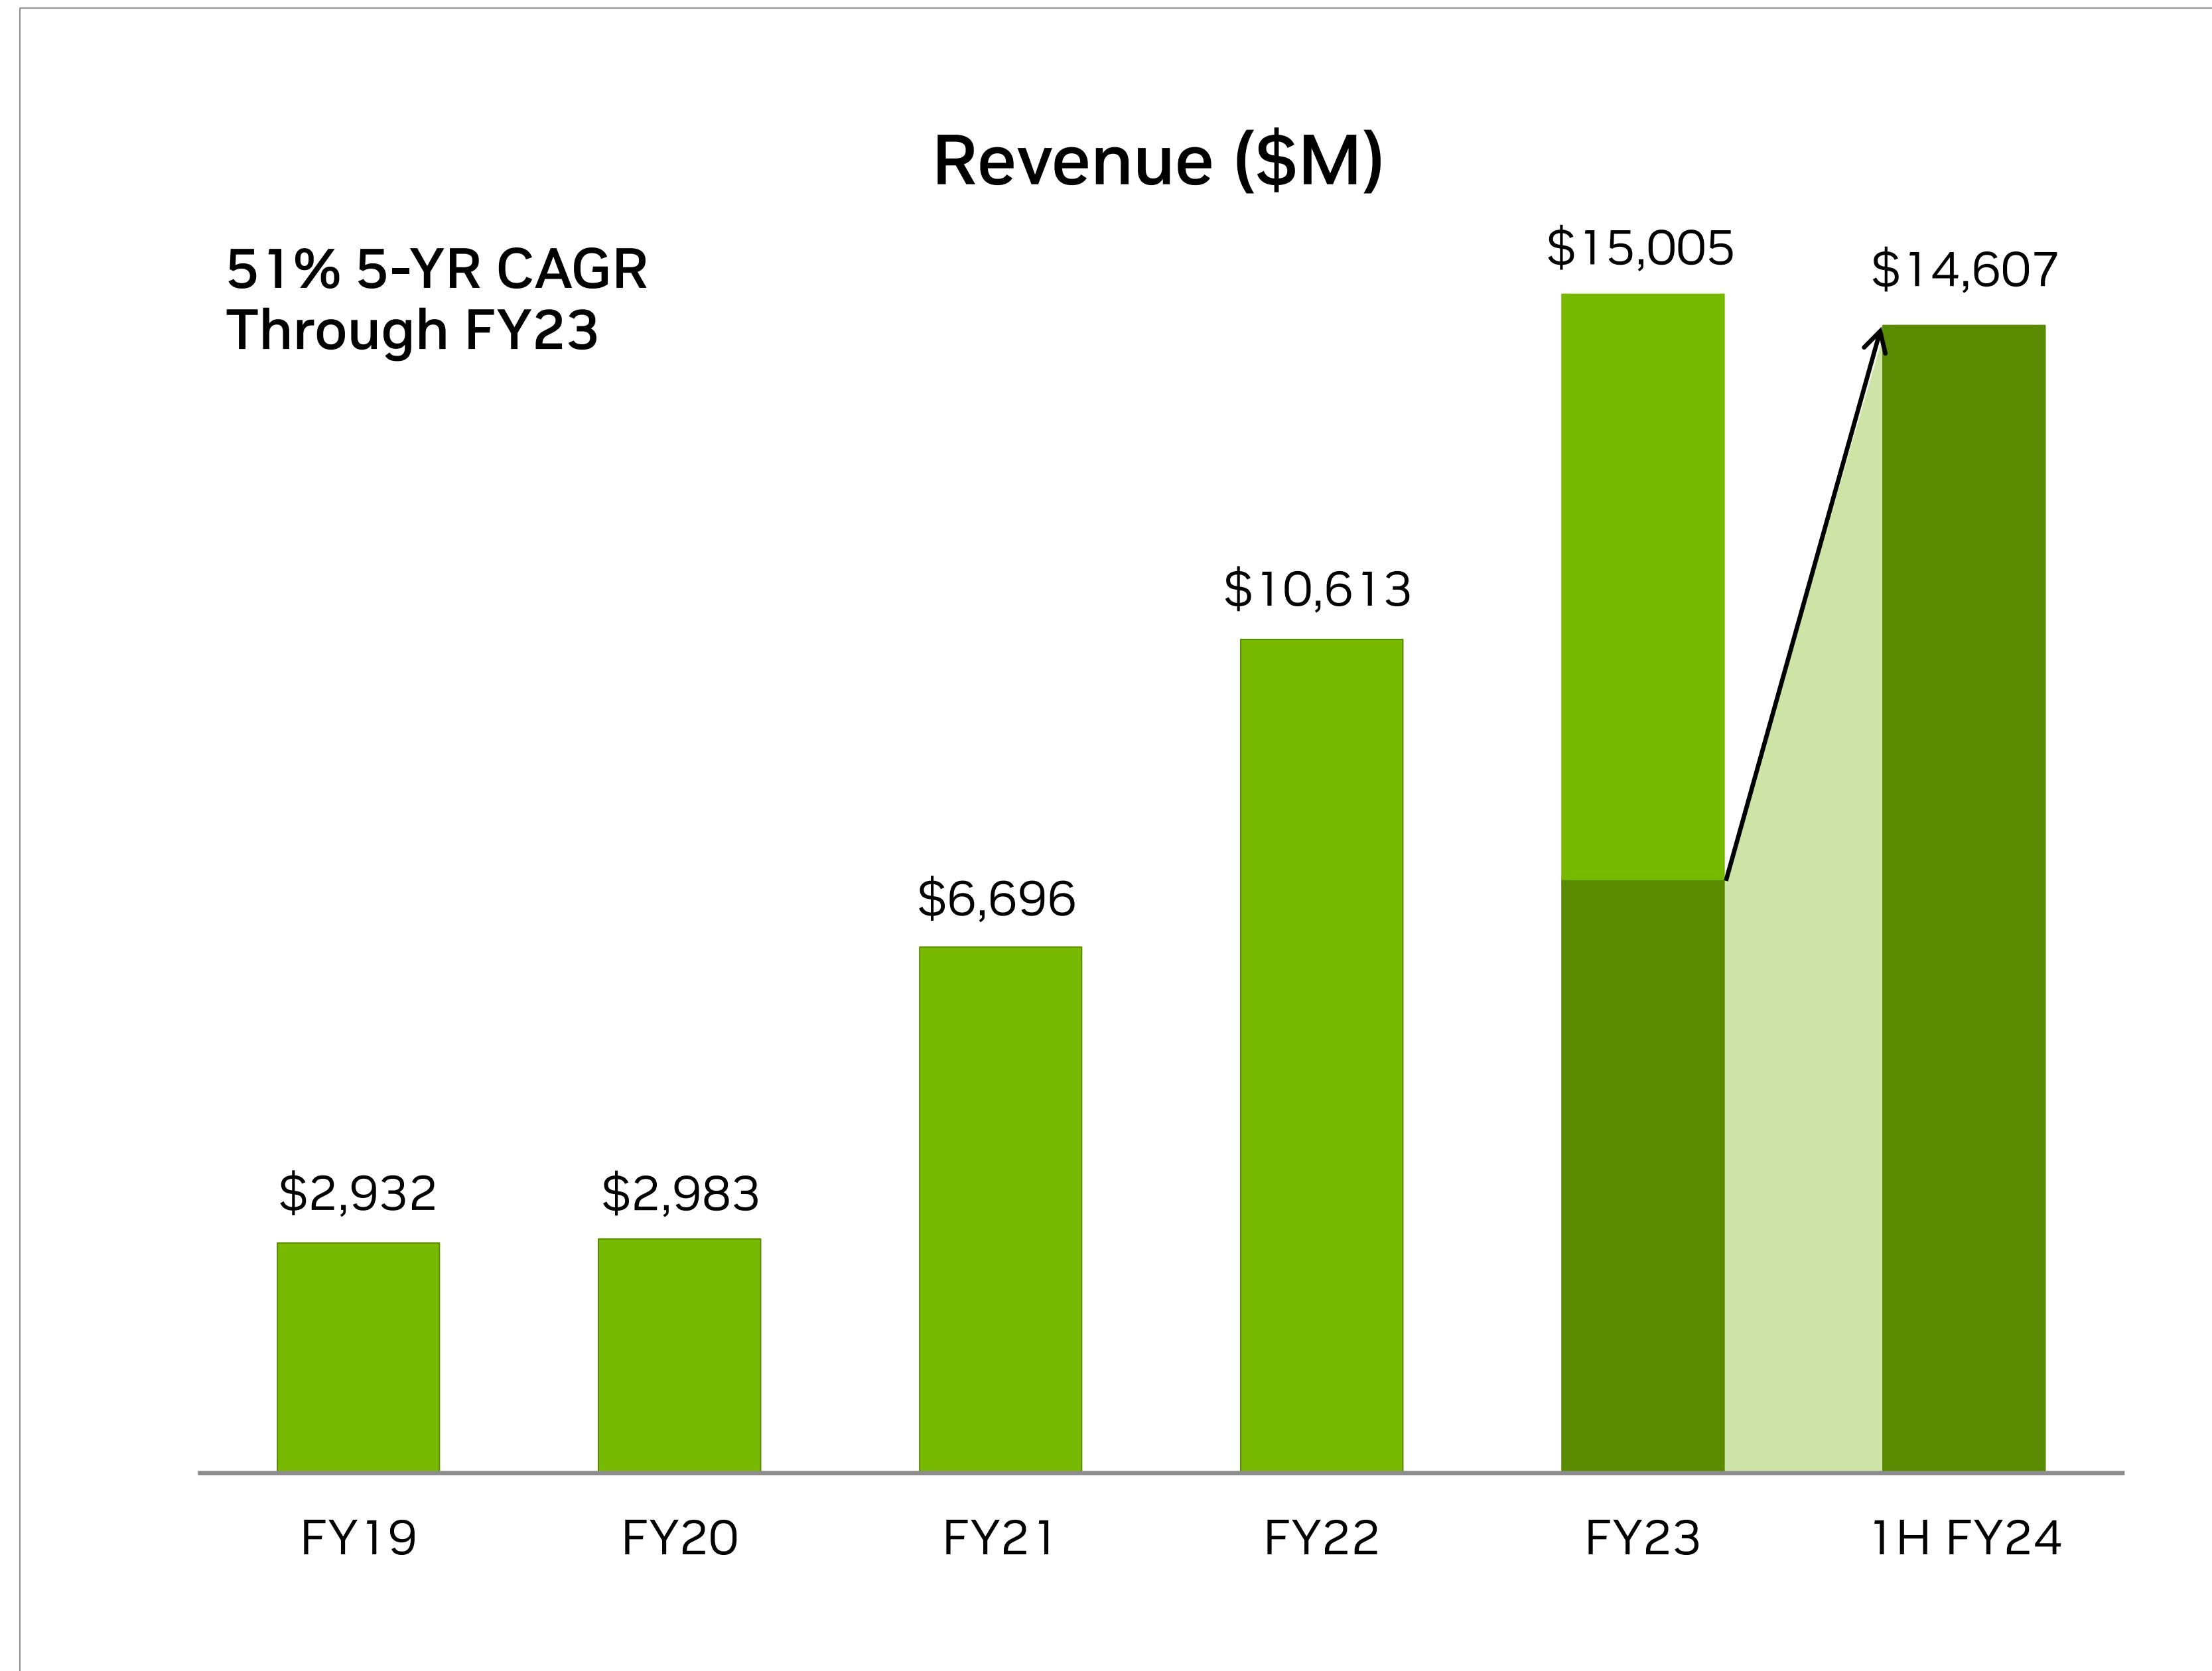
3% of FY23 revenue

FY23 Revenue \$0.9B
5-yr CAGR 10%

DRIVE Hyperion sensor architecture
with AGX compute
DRIVE AV & IX full stack software
for ADAS, AV & AI cockpit

Data Center

The leading computing platform for AI, HPC & graphics

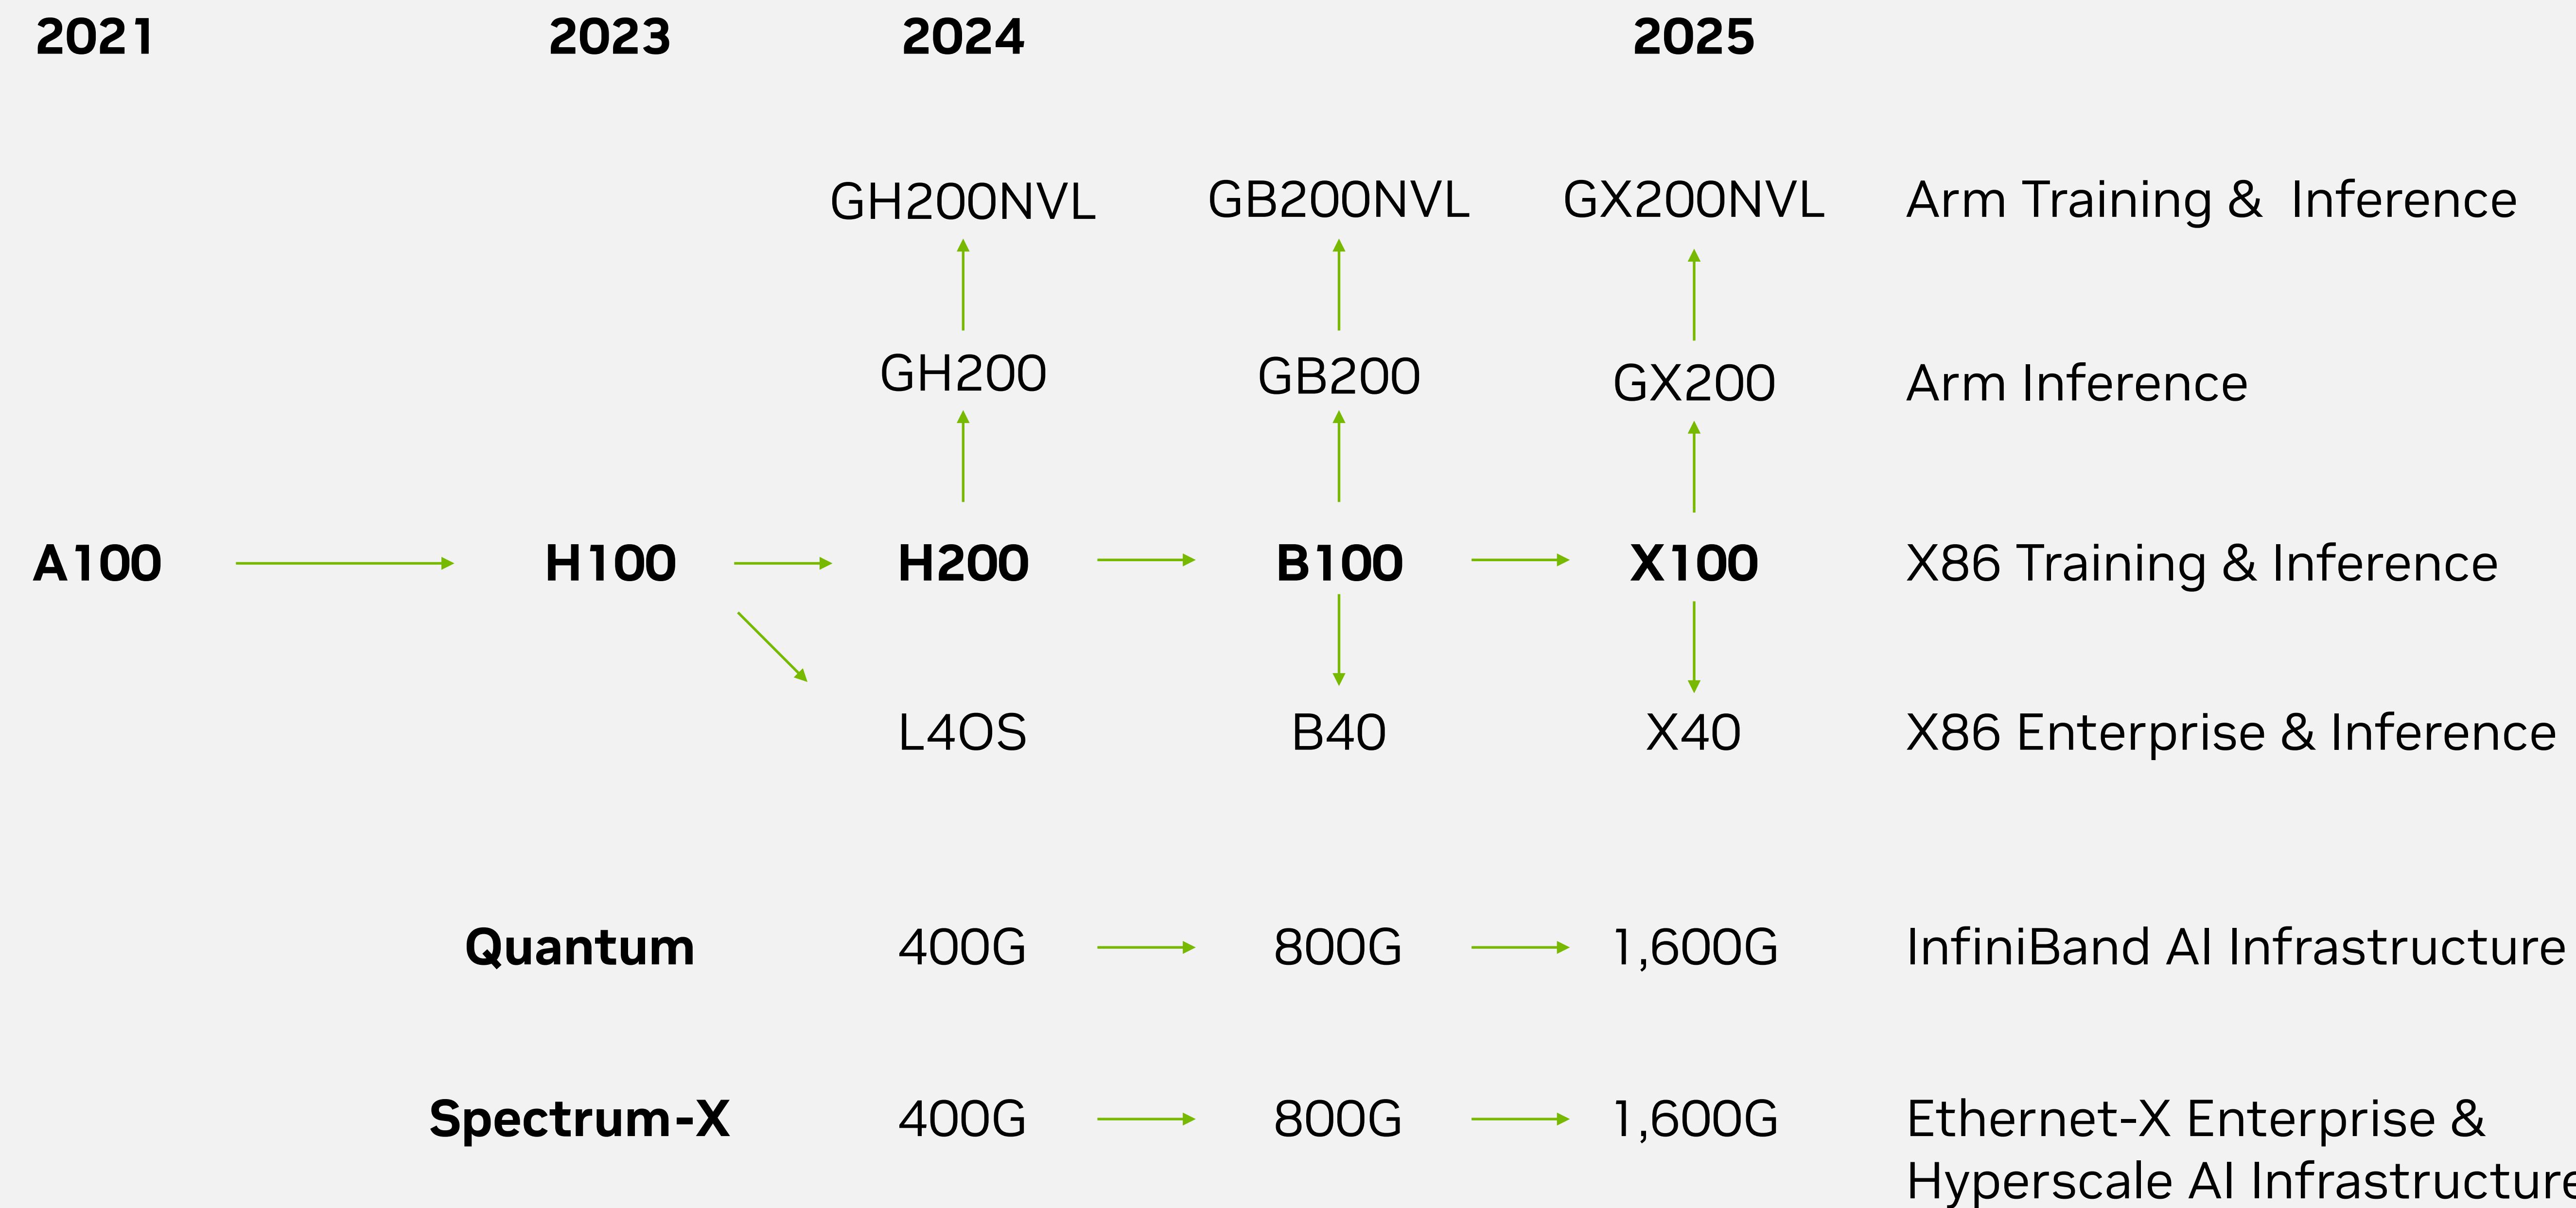


NVIDIA AI – One Architecture | Train and Deploy Everywhere

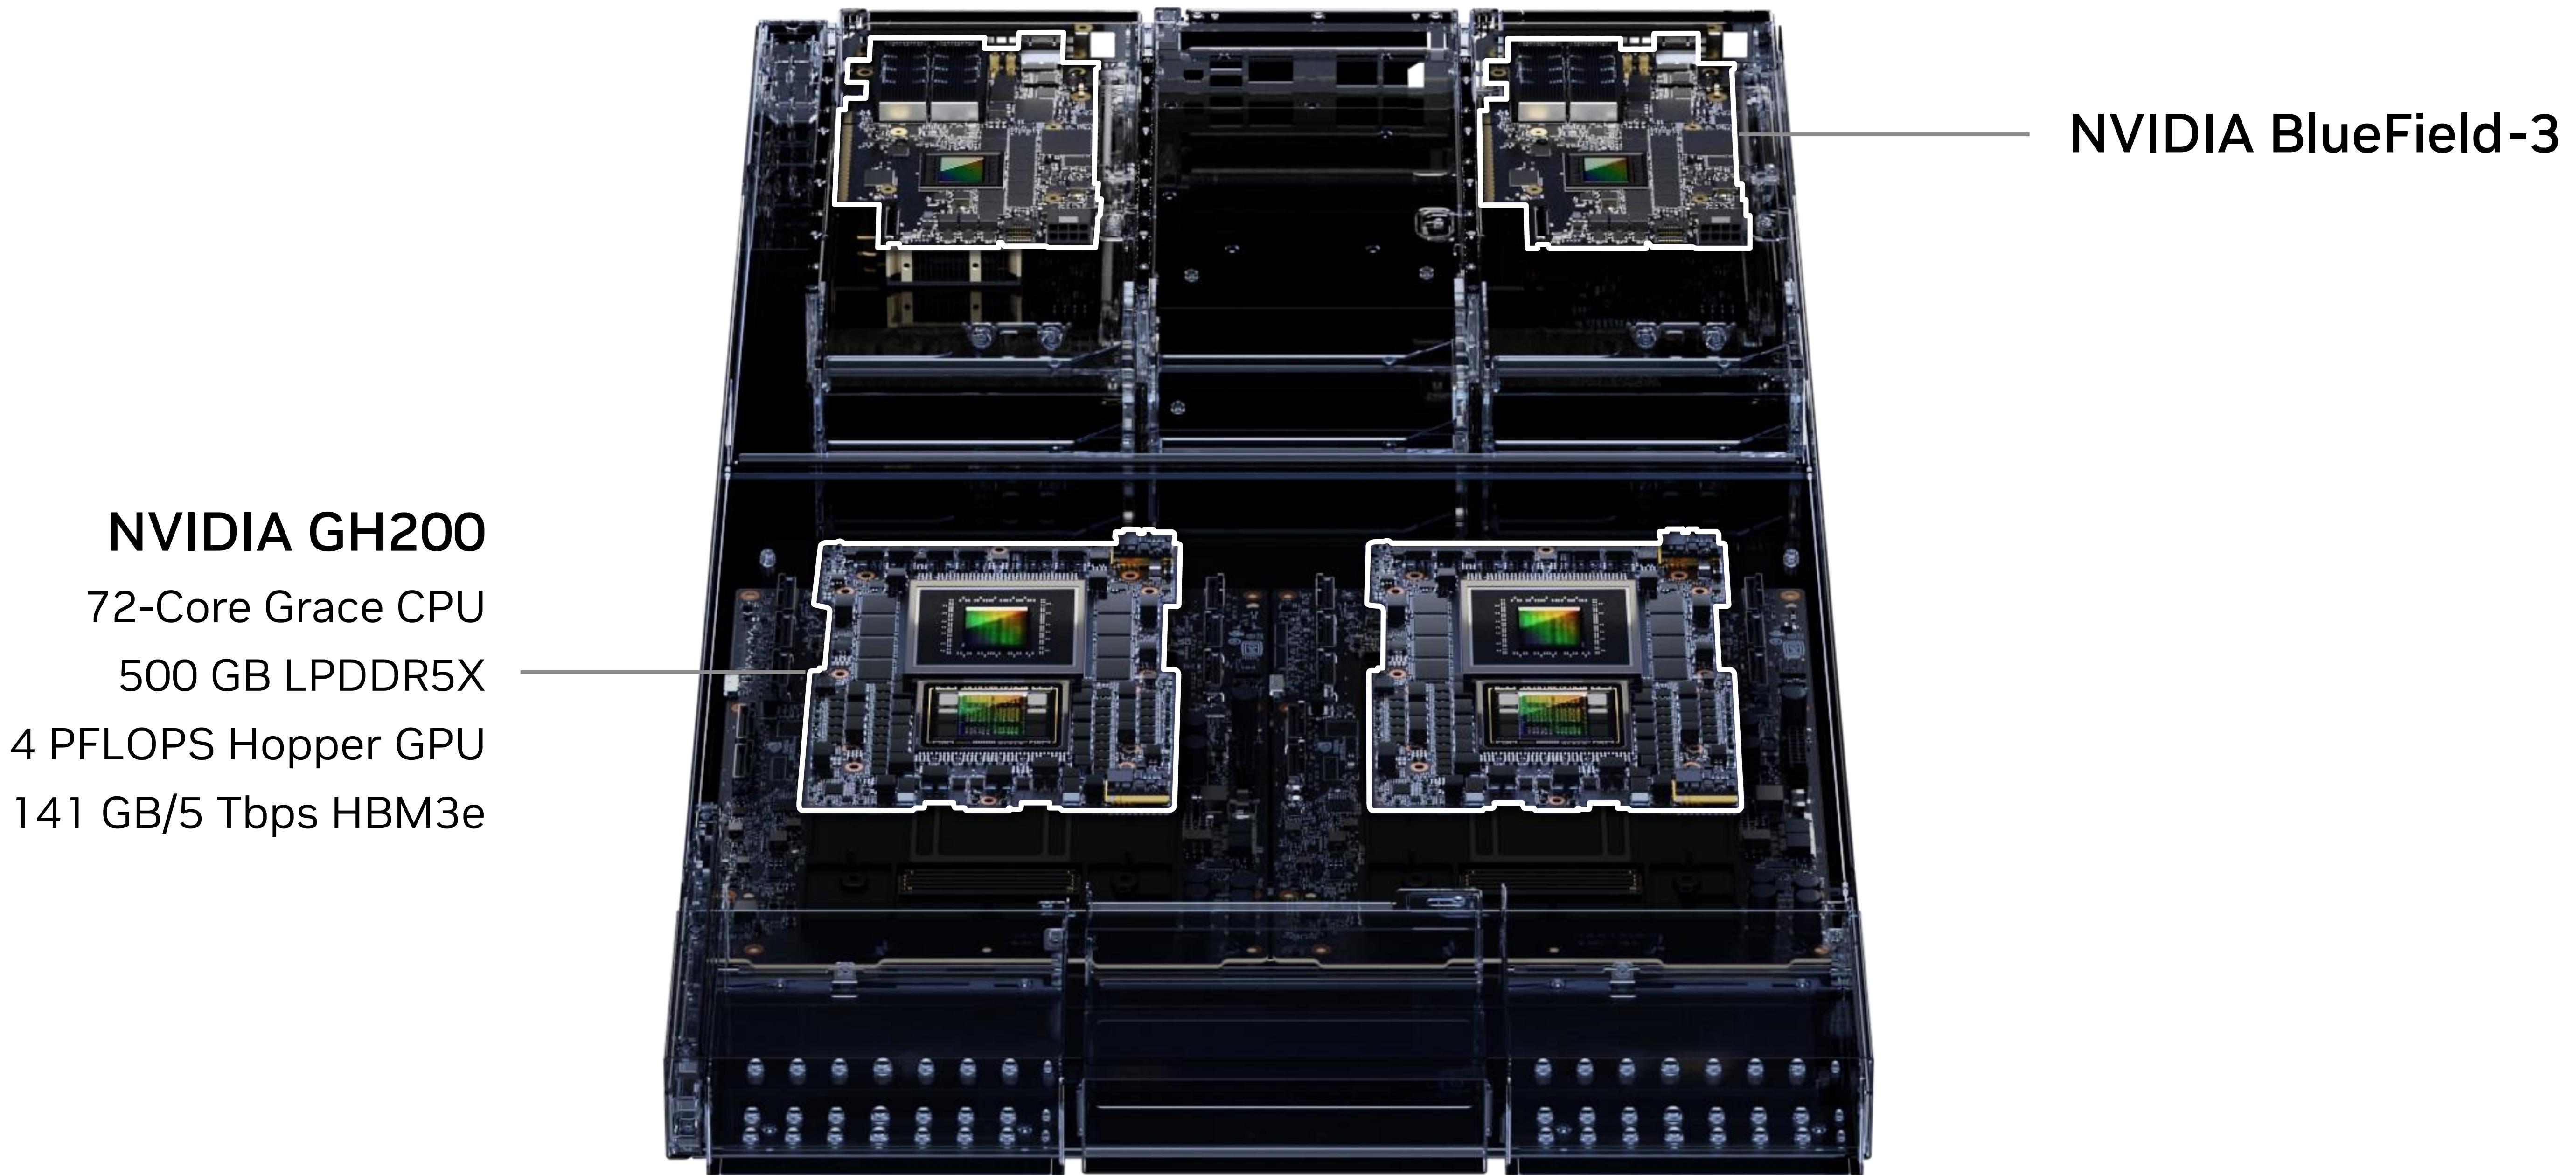
From Two-Year Rhythm

to

One-Year Rhythm | Training & Inference | x86 & Arm | Hyperscale & Enterprise

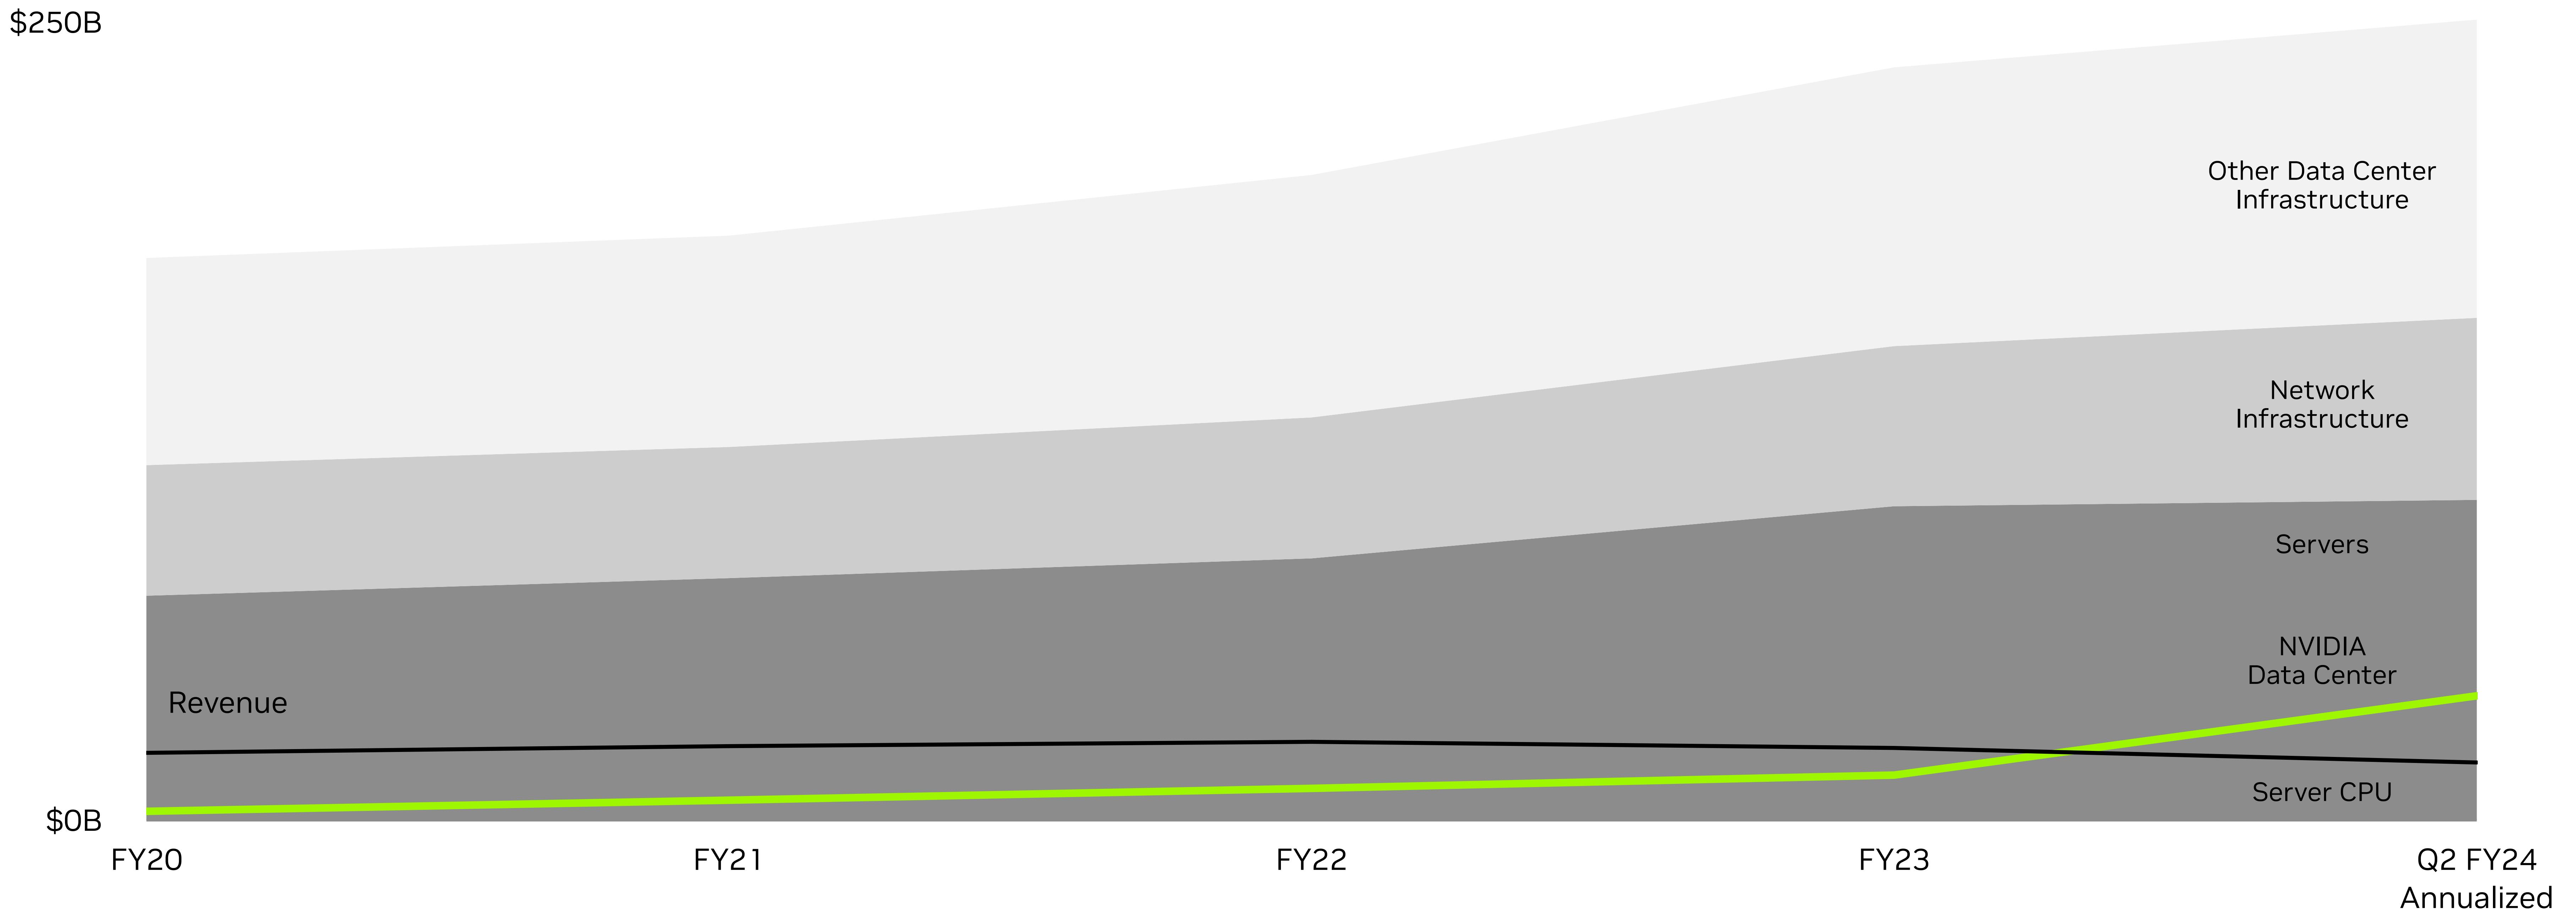


NVIDIA Grace Hopper Superchip



Addressing the Entire Data Center

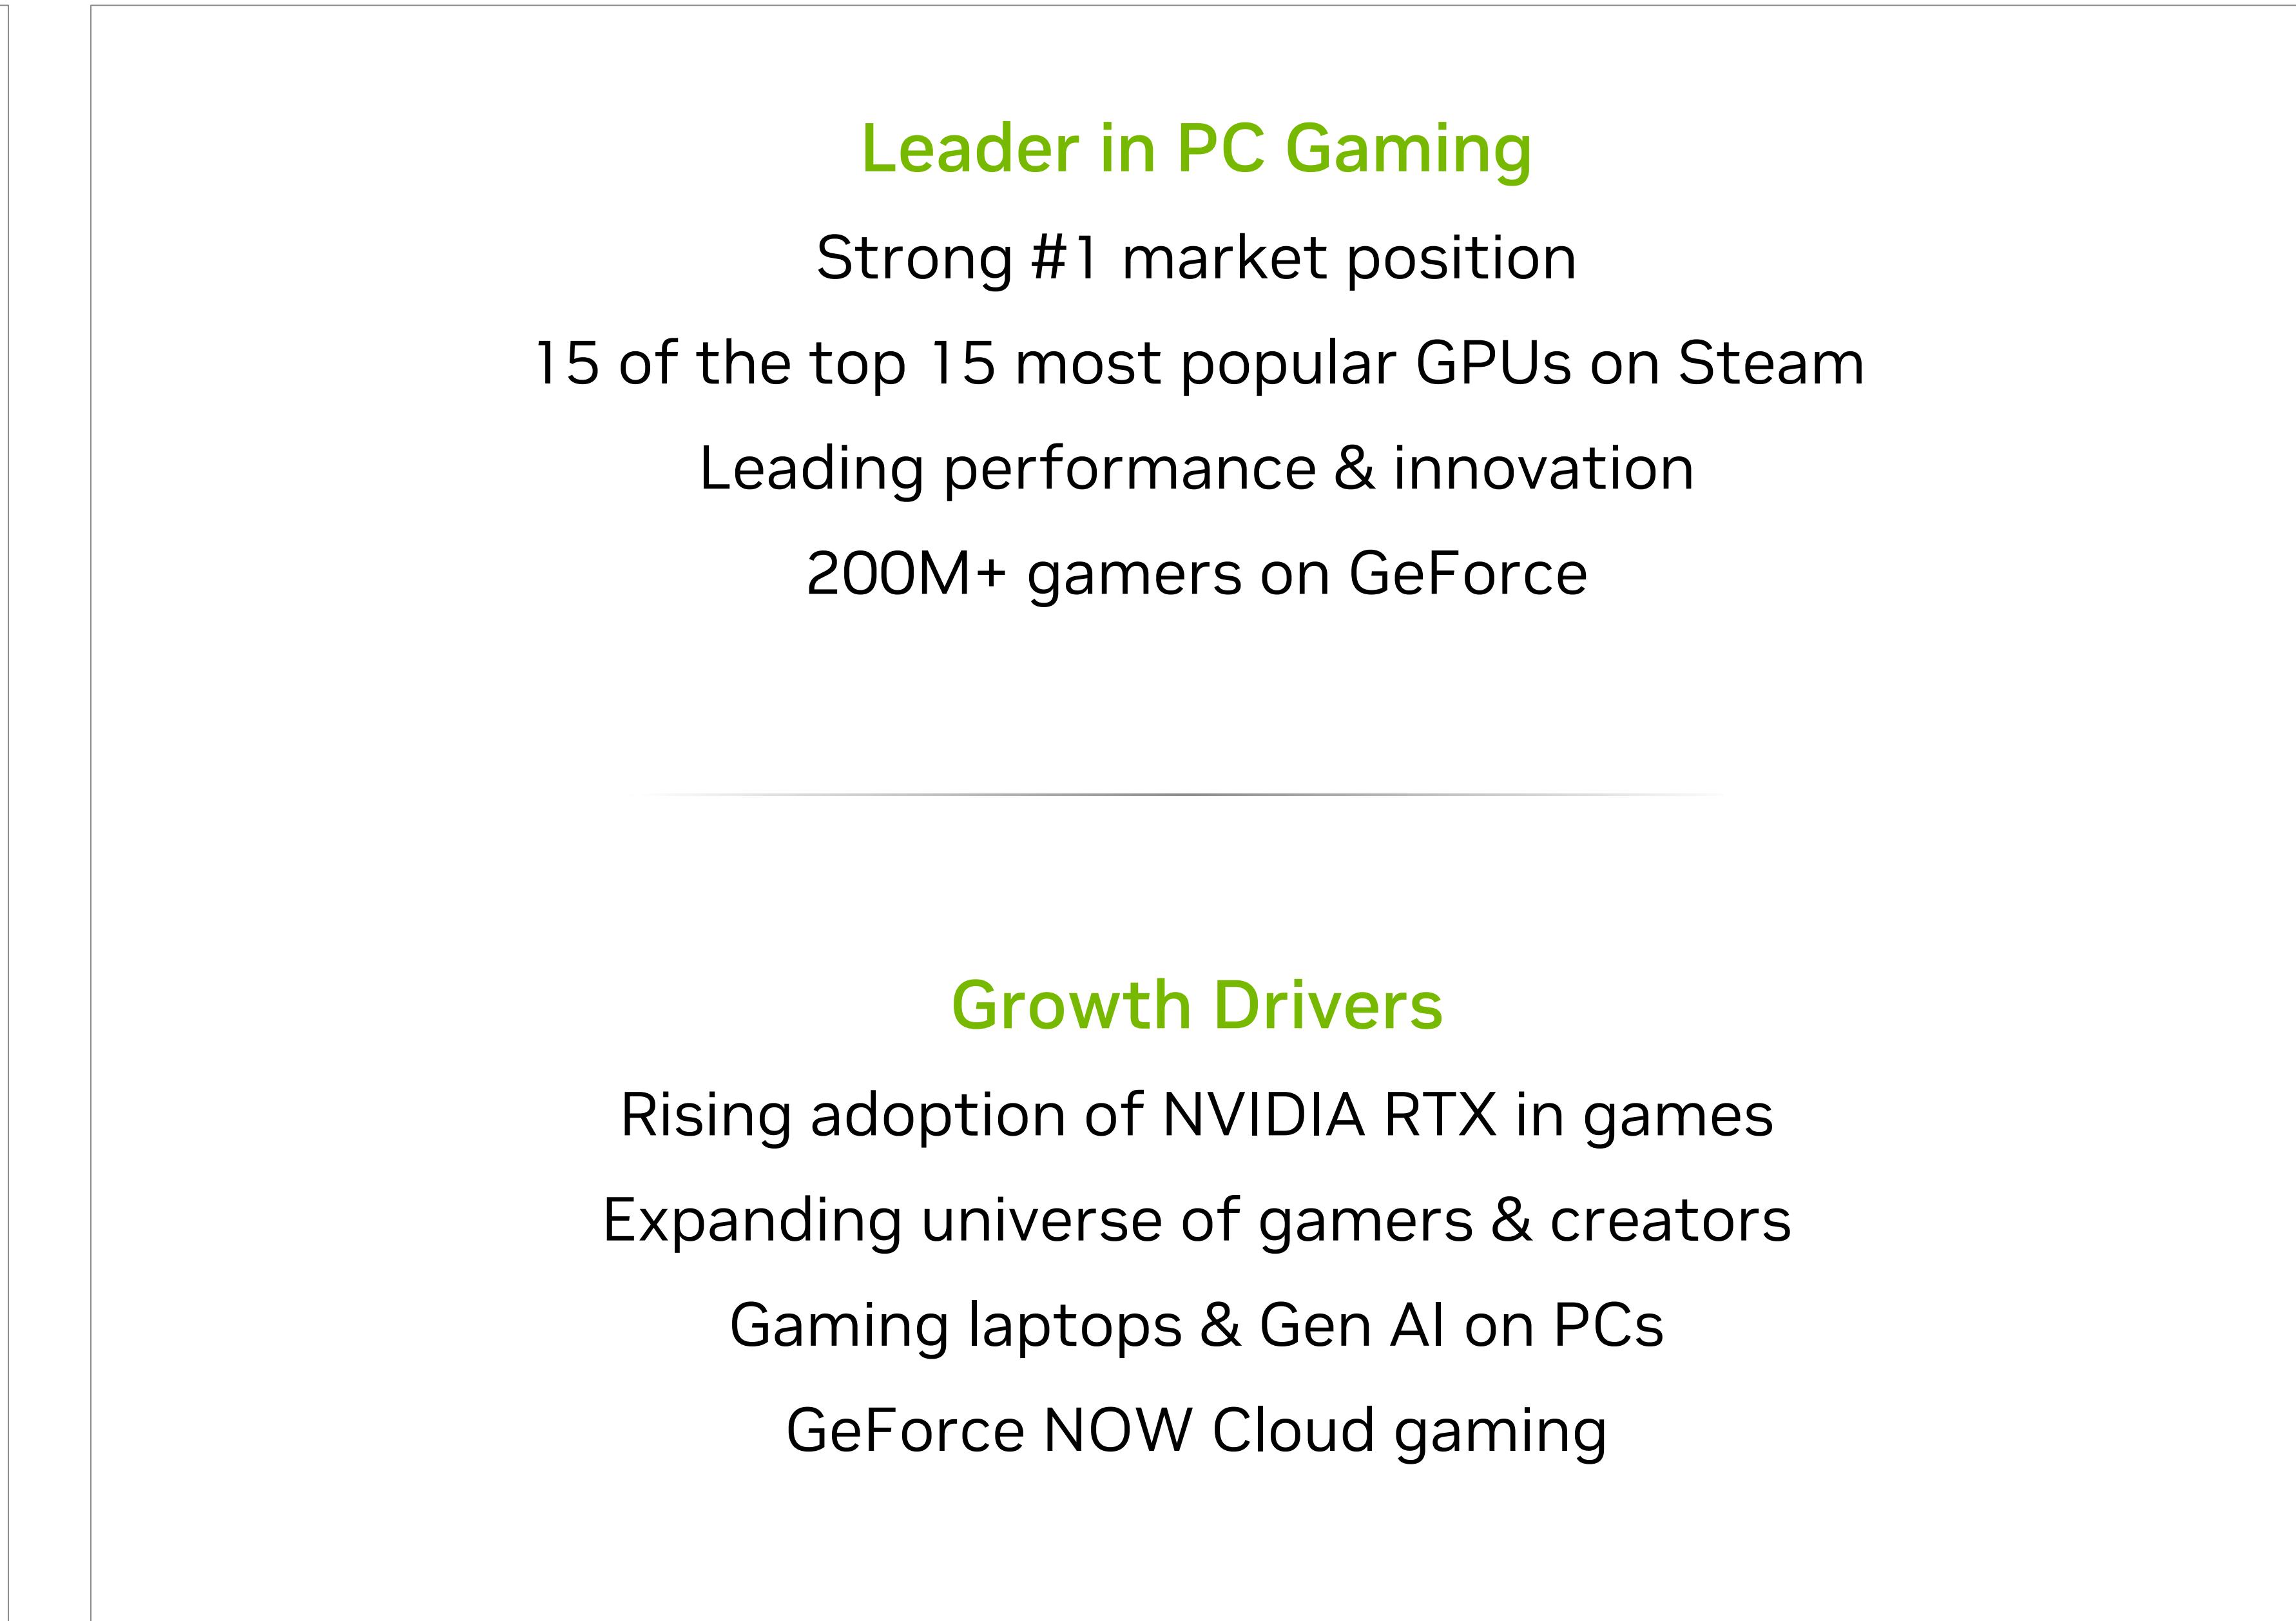
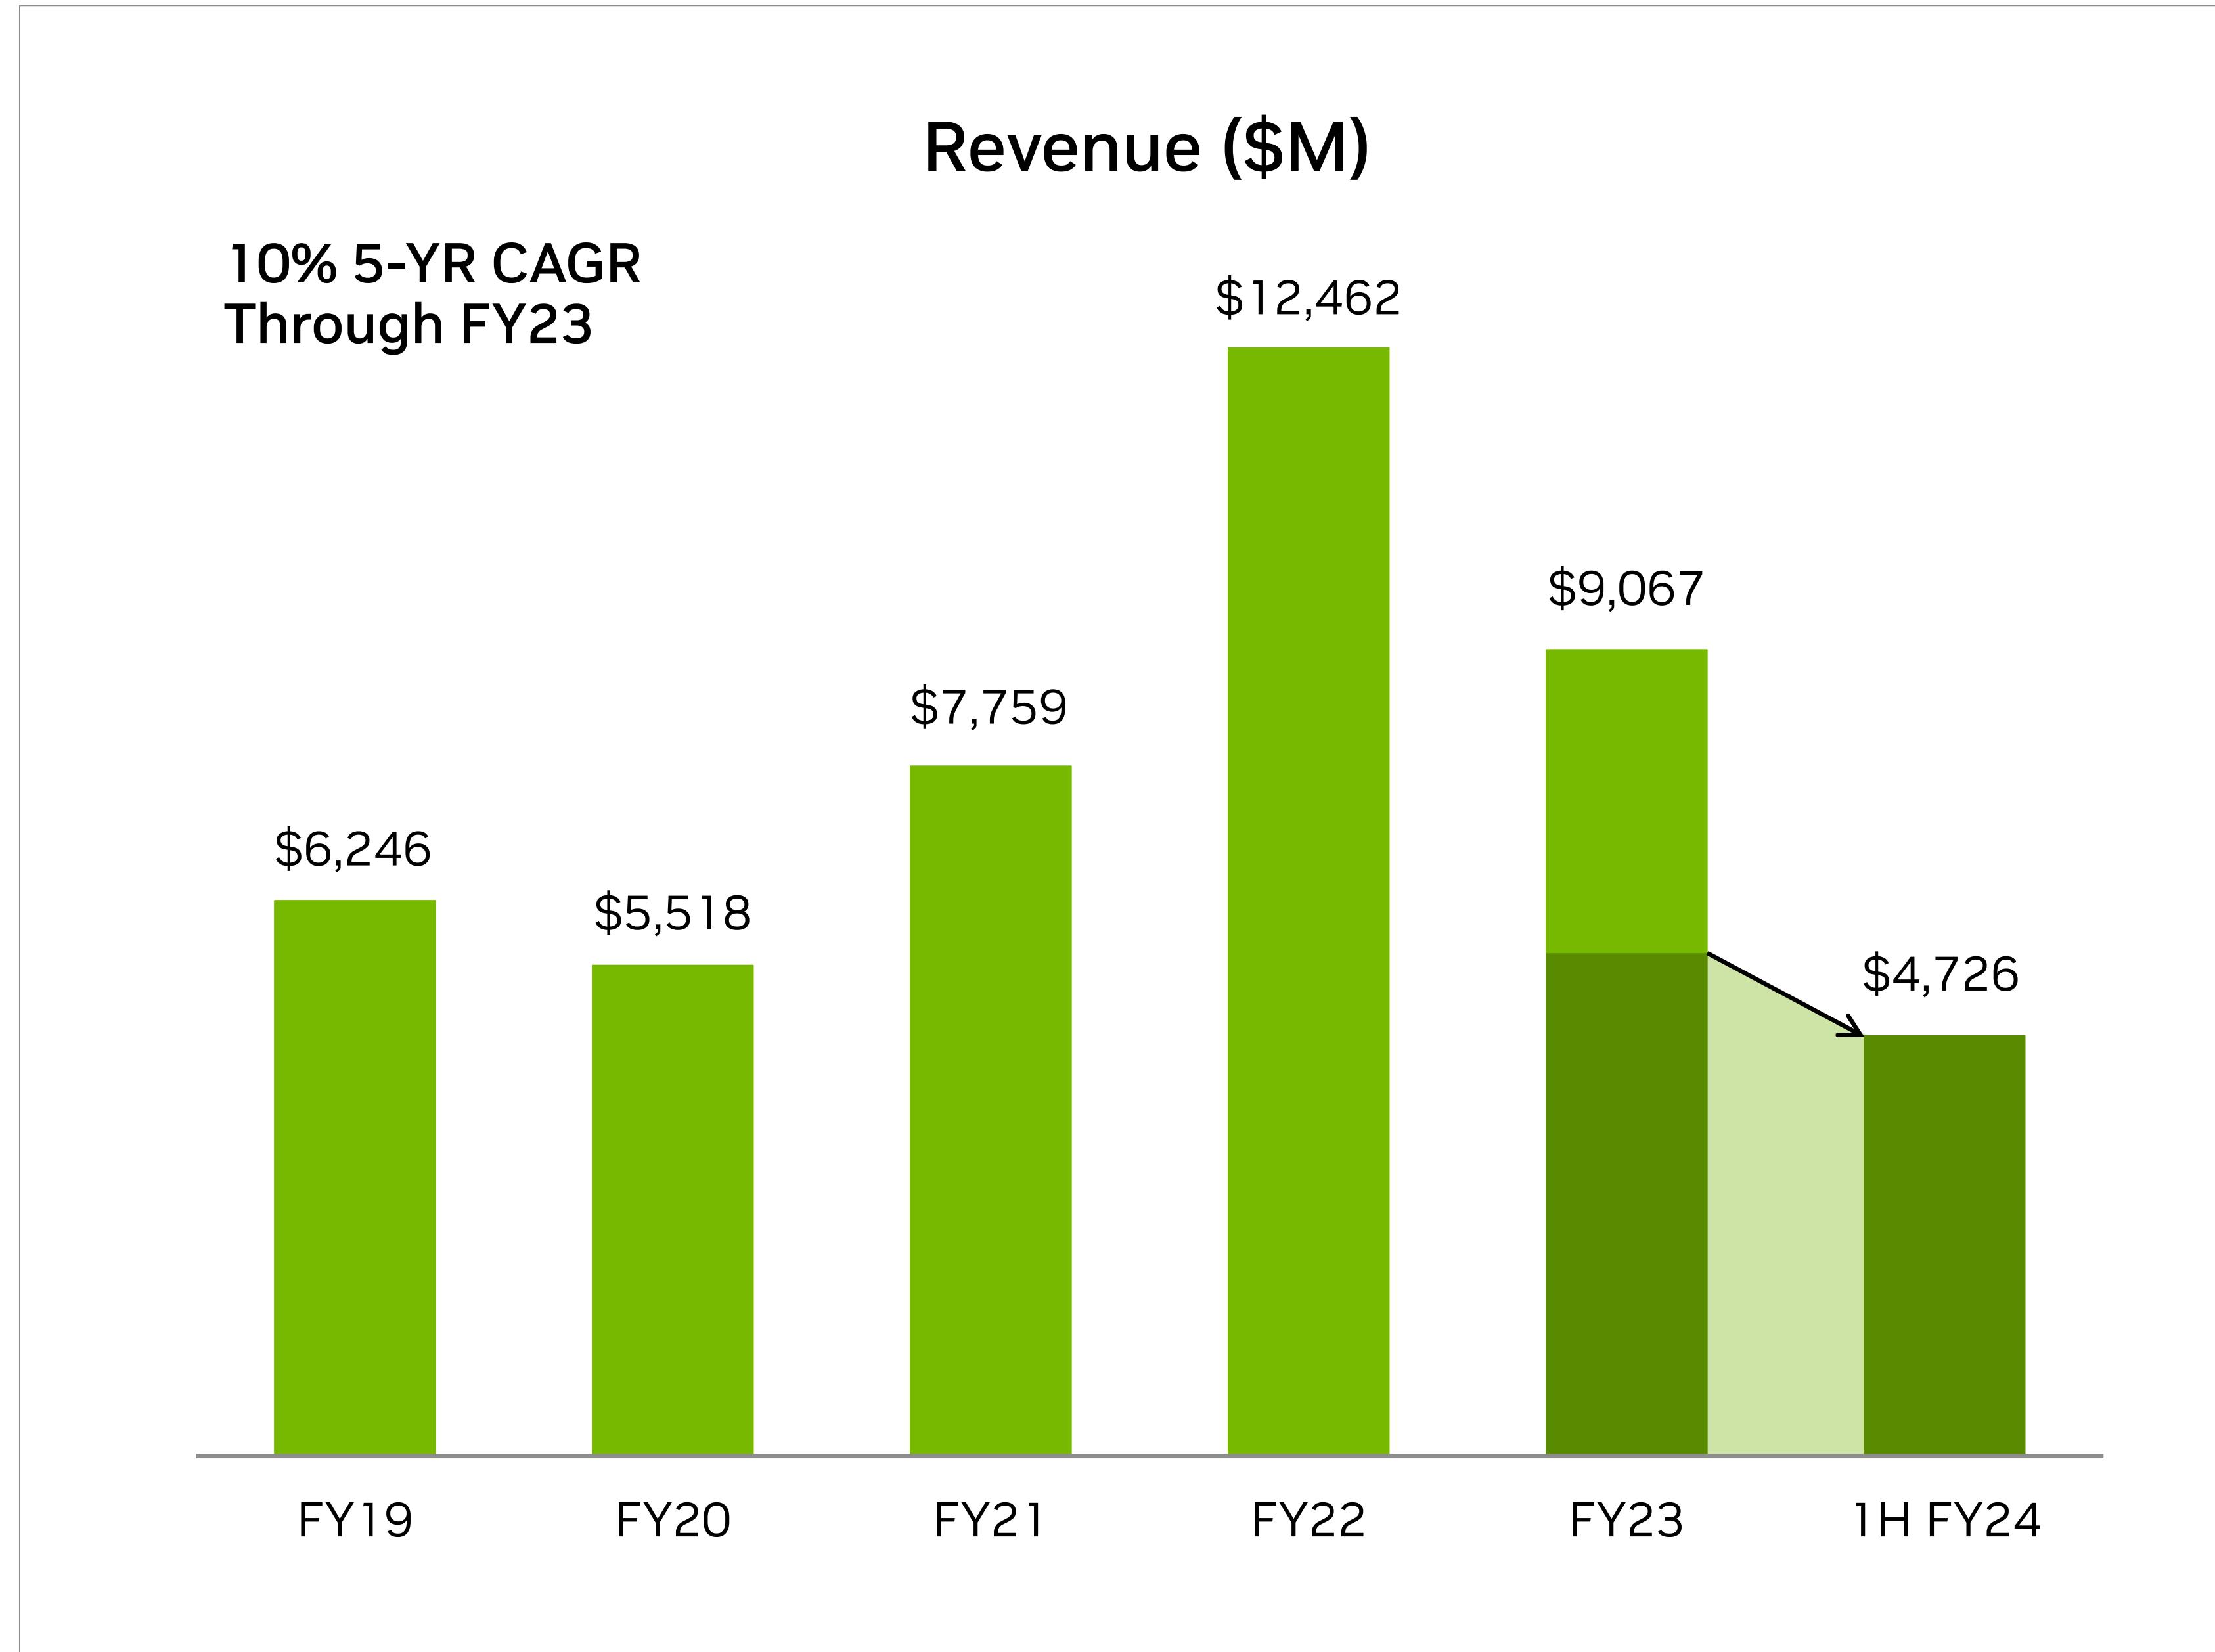
\$1T+ data center infrastructure installed base



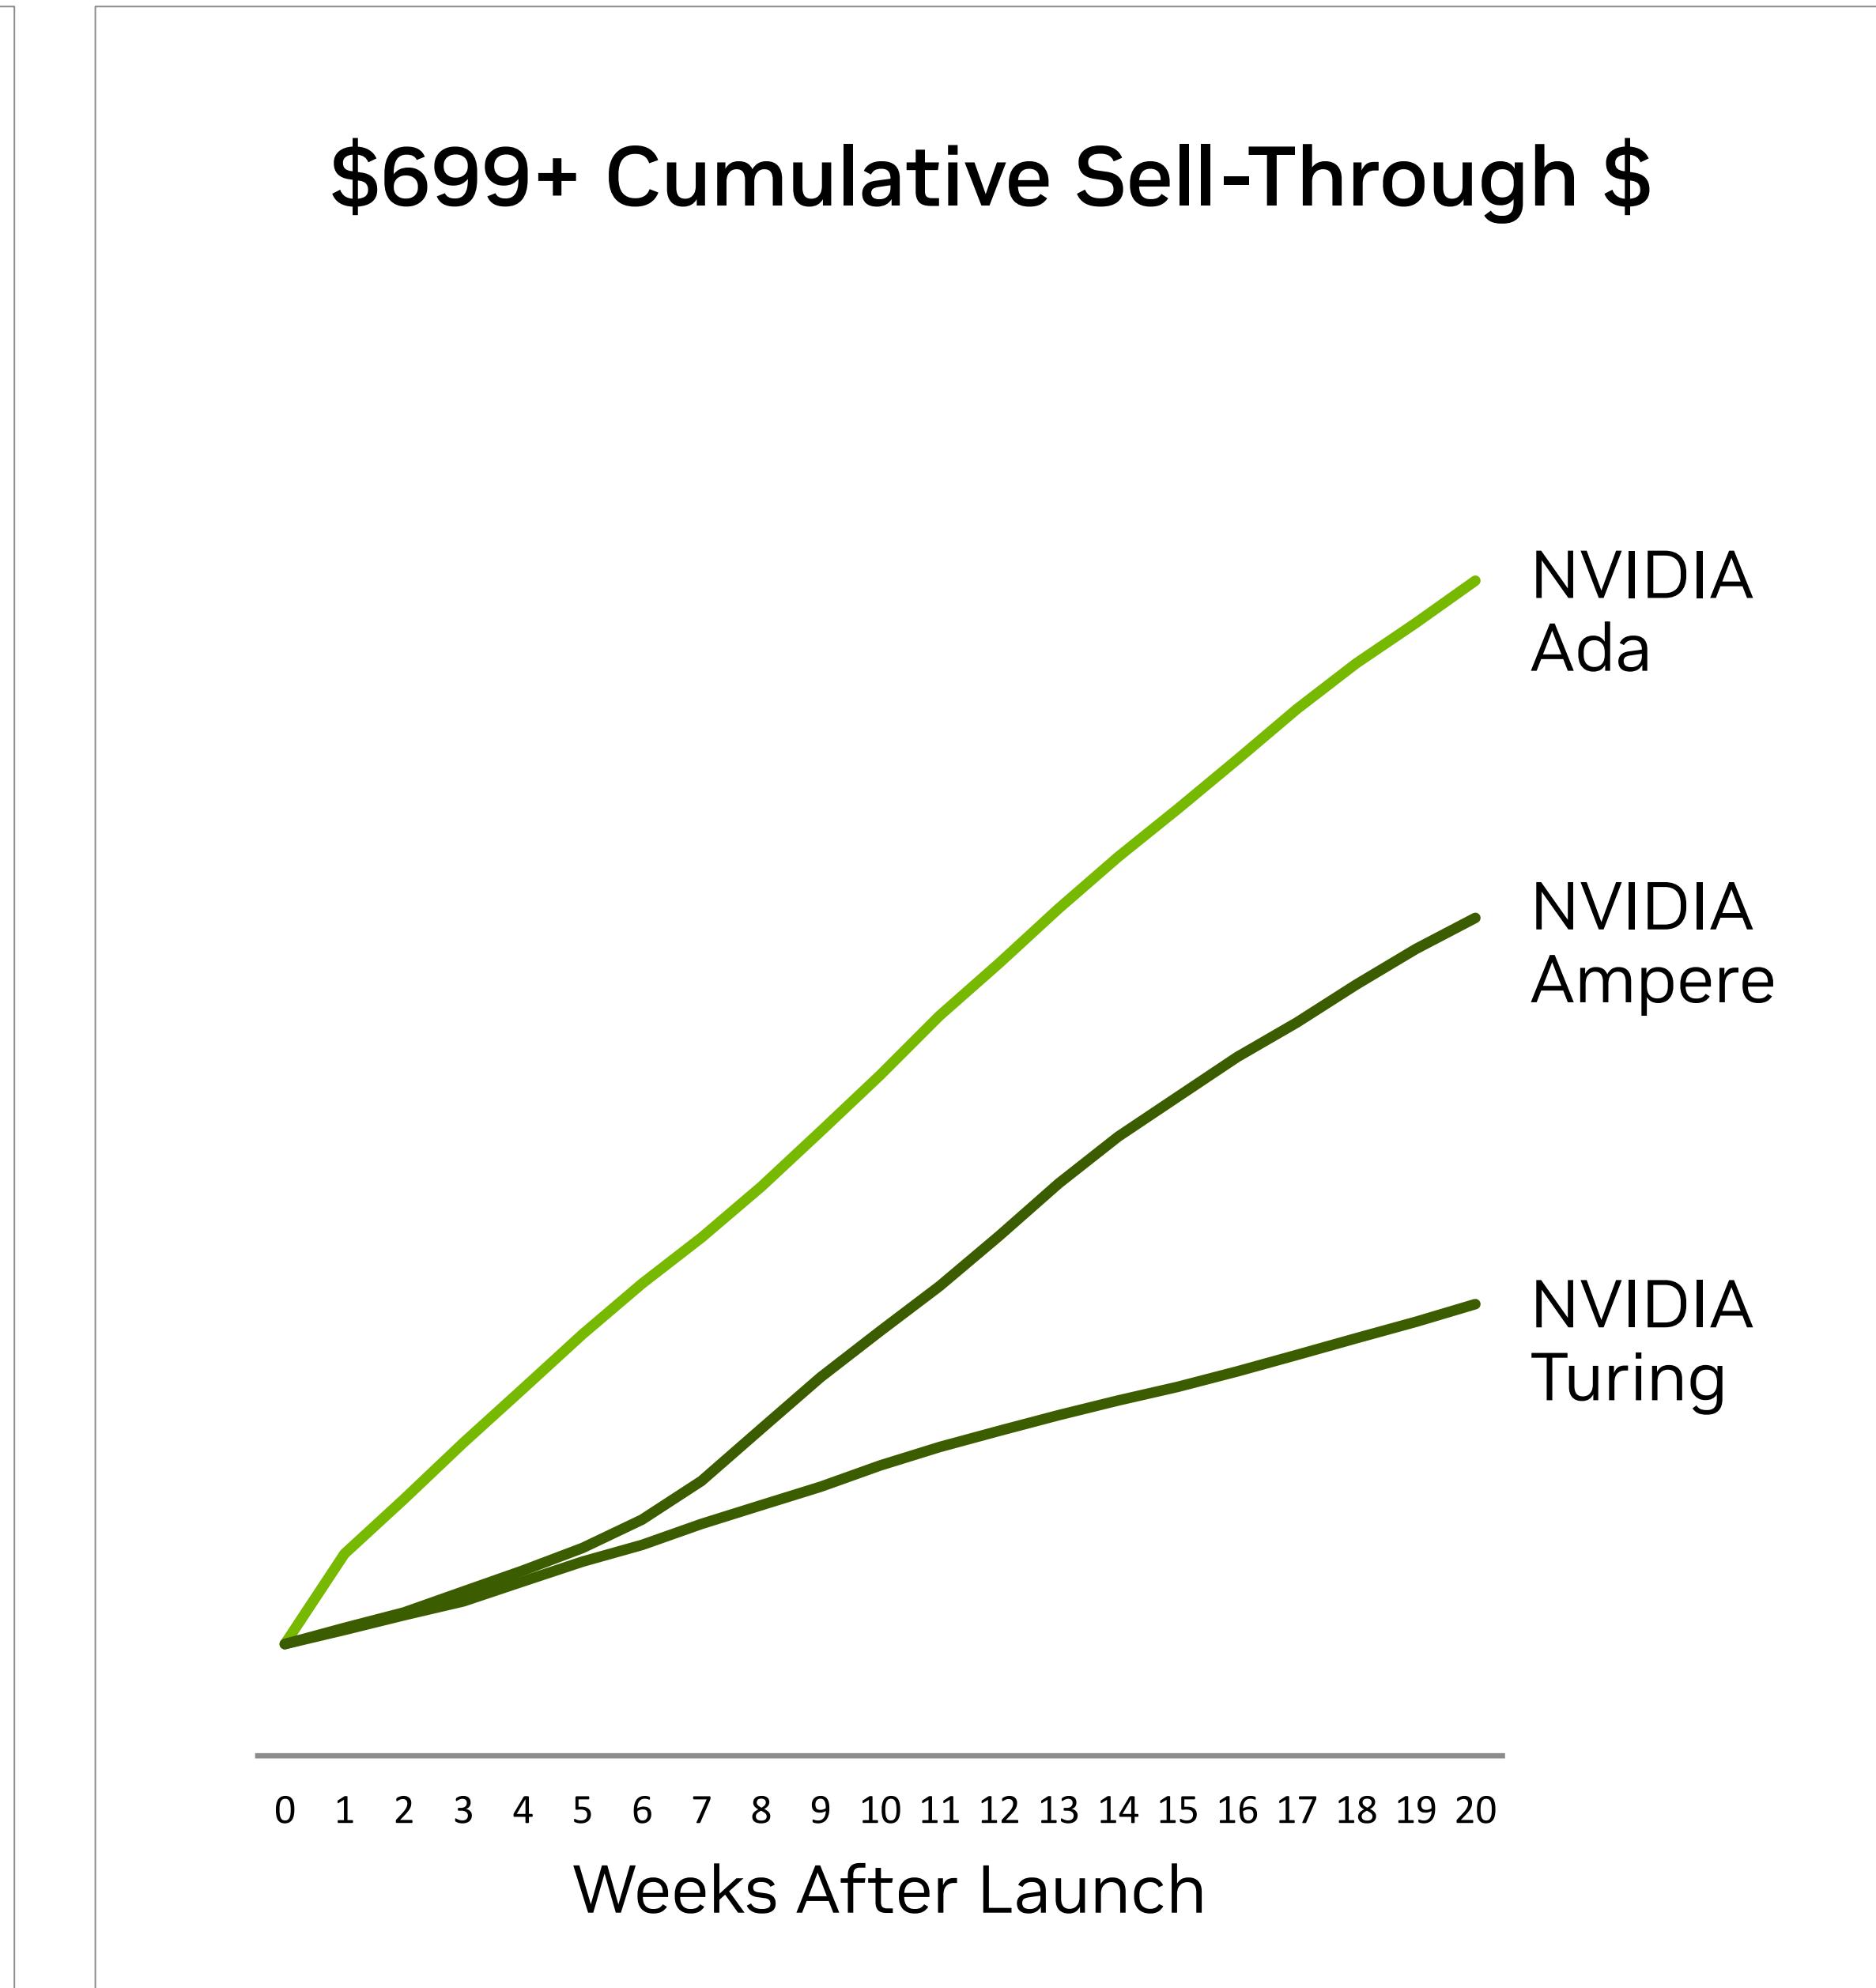
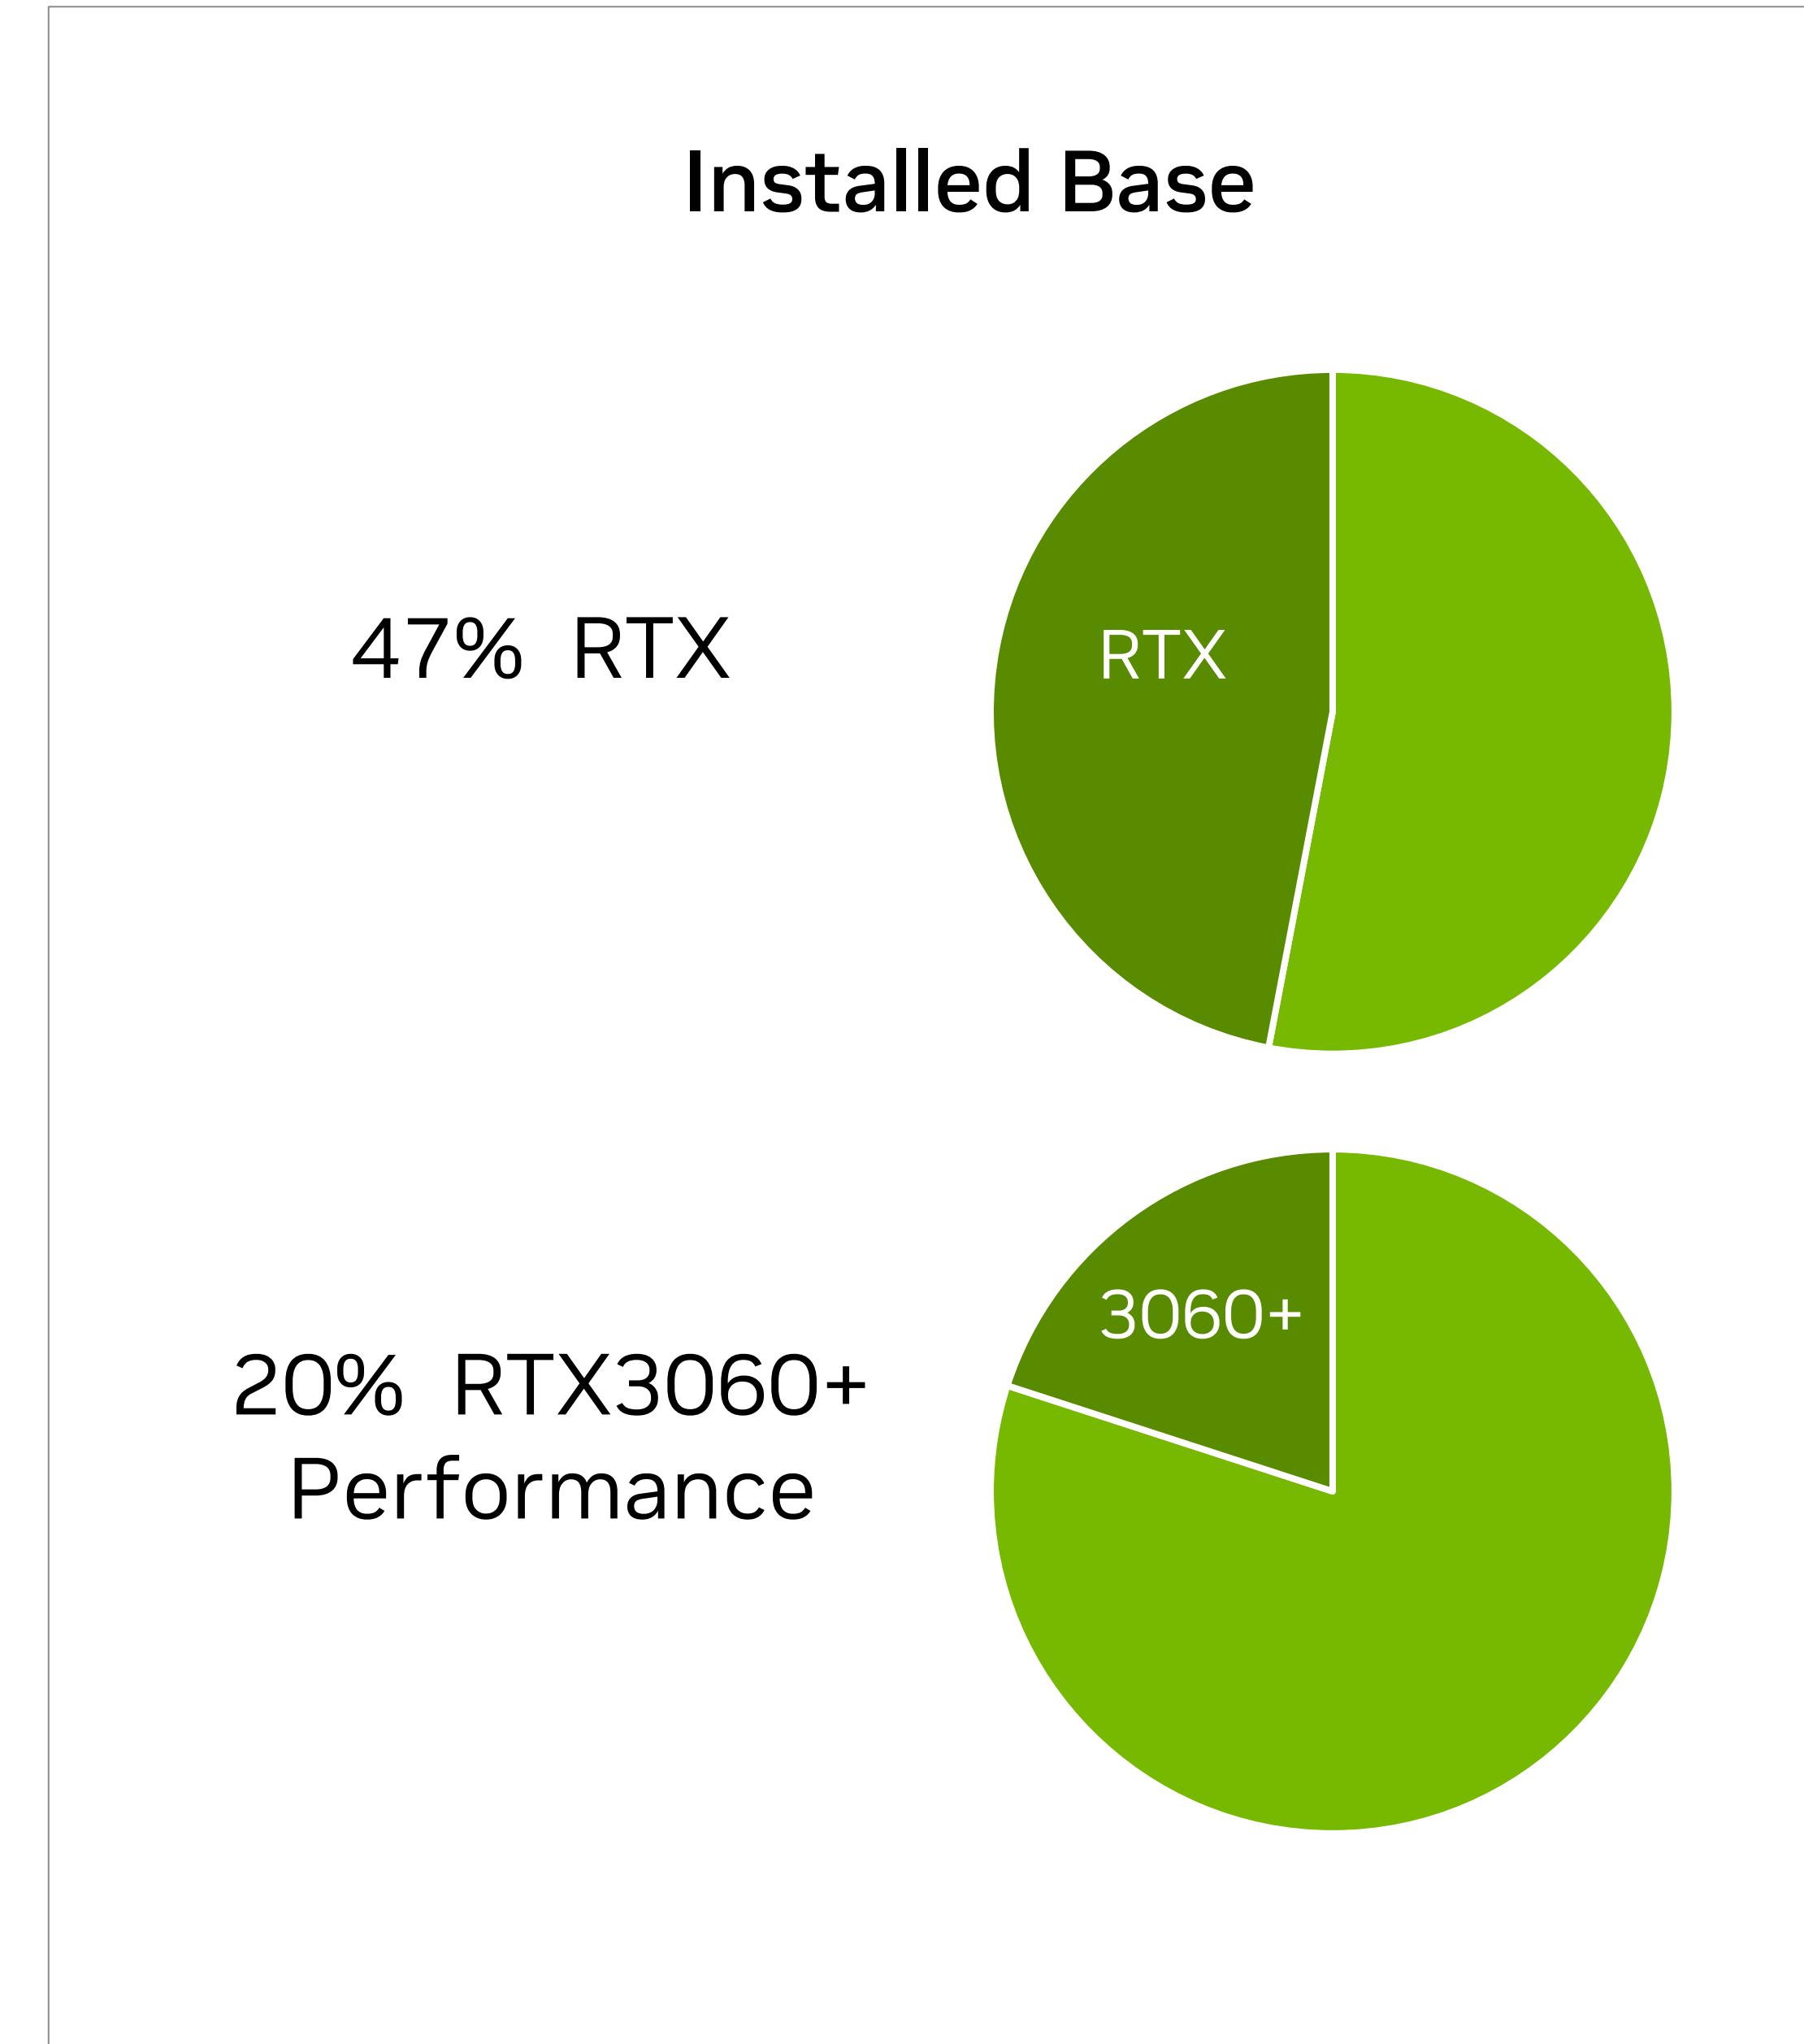
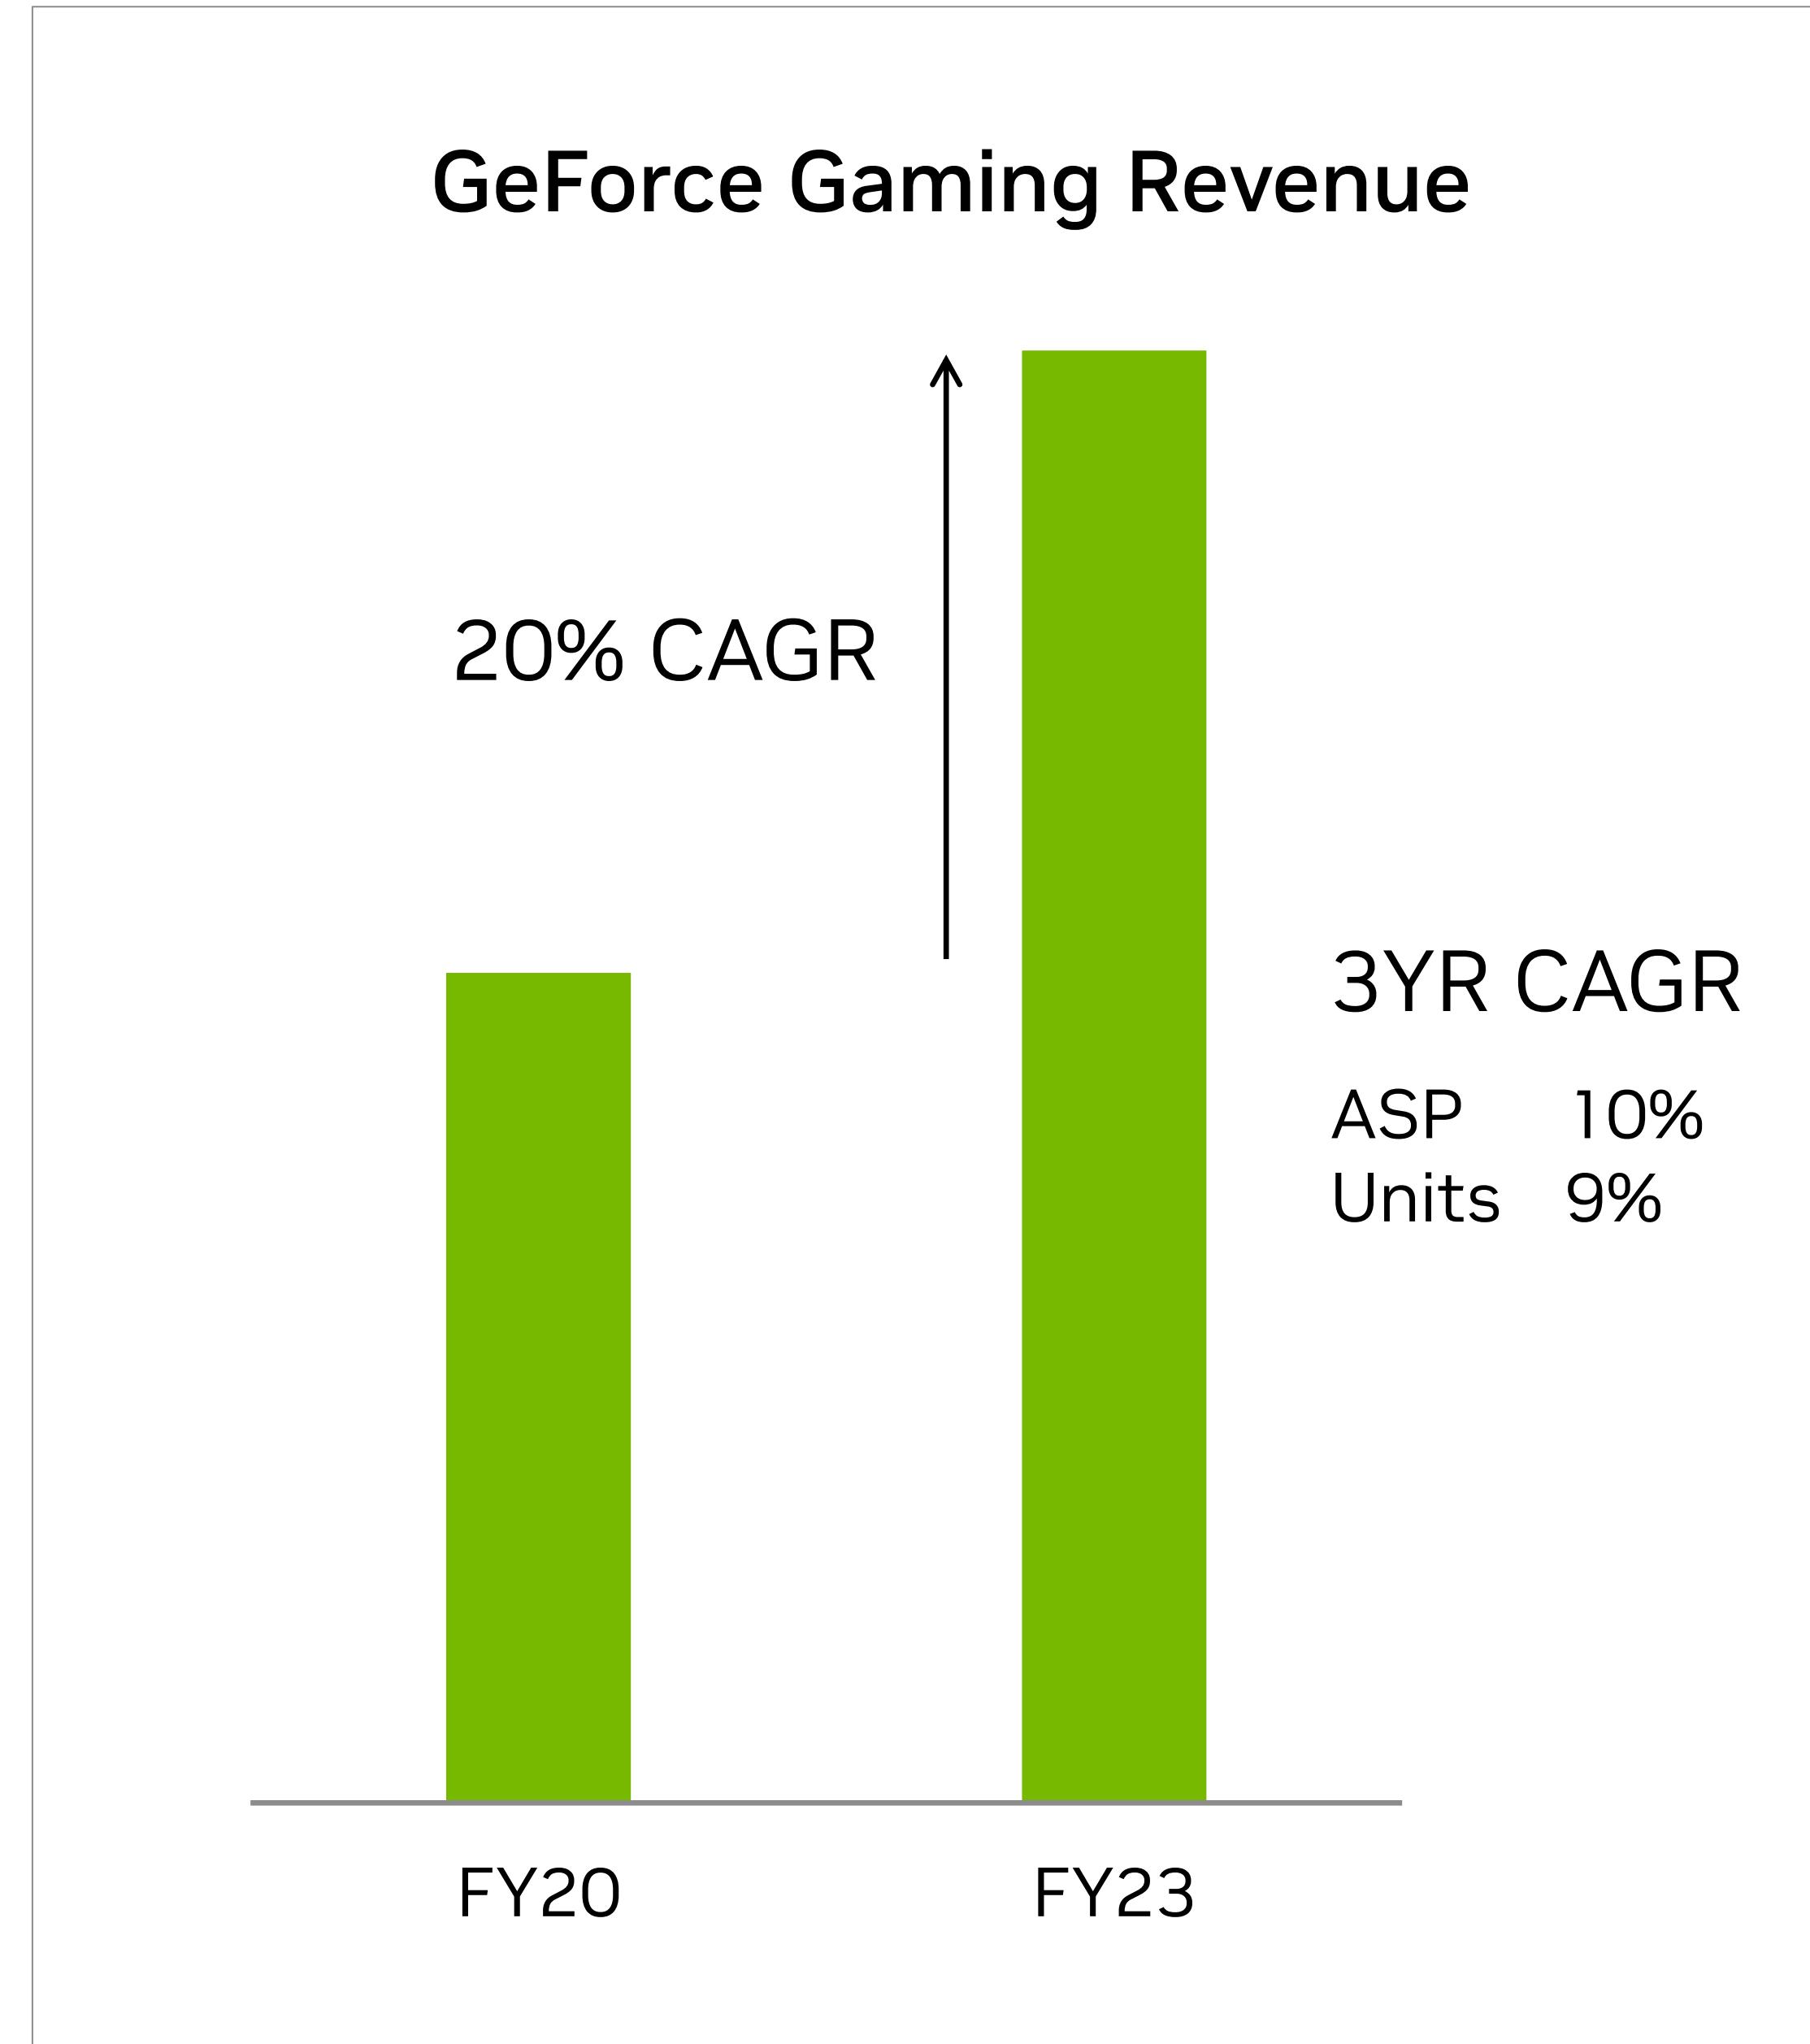
Source: Mercury Research, Dell'Oro
Assumes NVIDIA Fiscal Year aligns to Calendar Year (e.g. FY23 = CY22)

Gaming

GeForce — the world's largest gaming platform



GeForce Extends Growth, Large Upgrade Opportunity



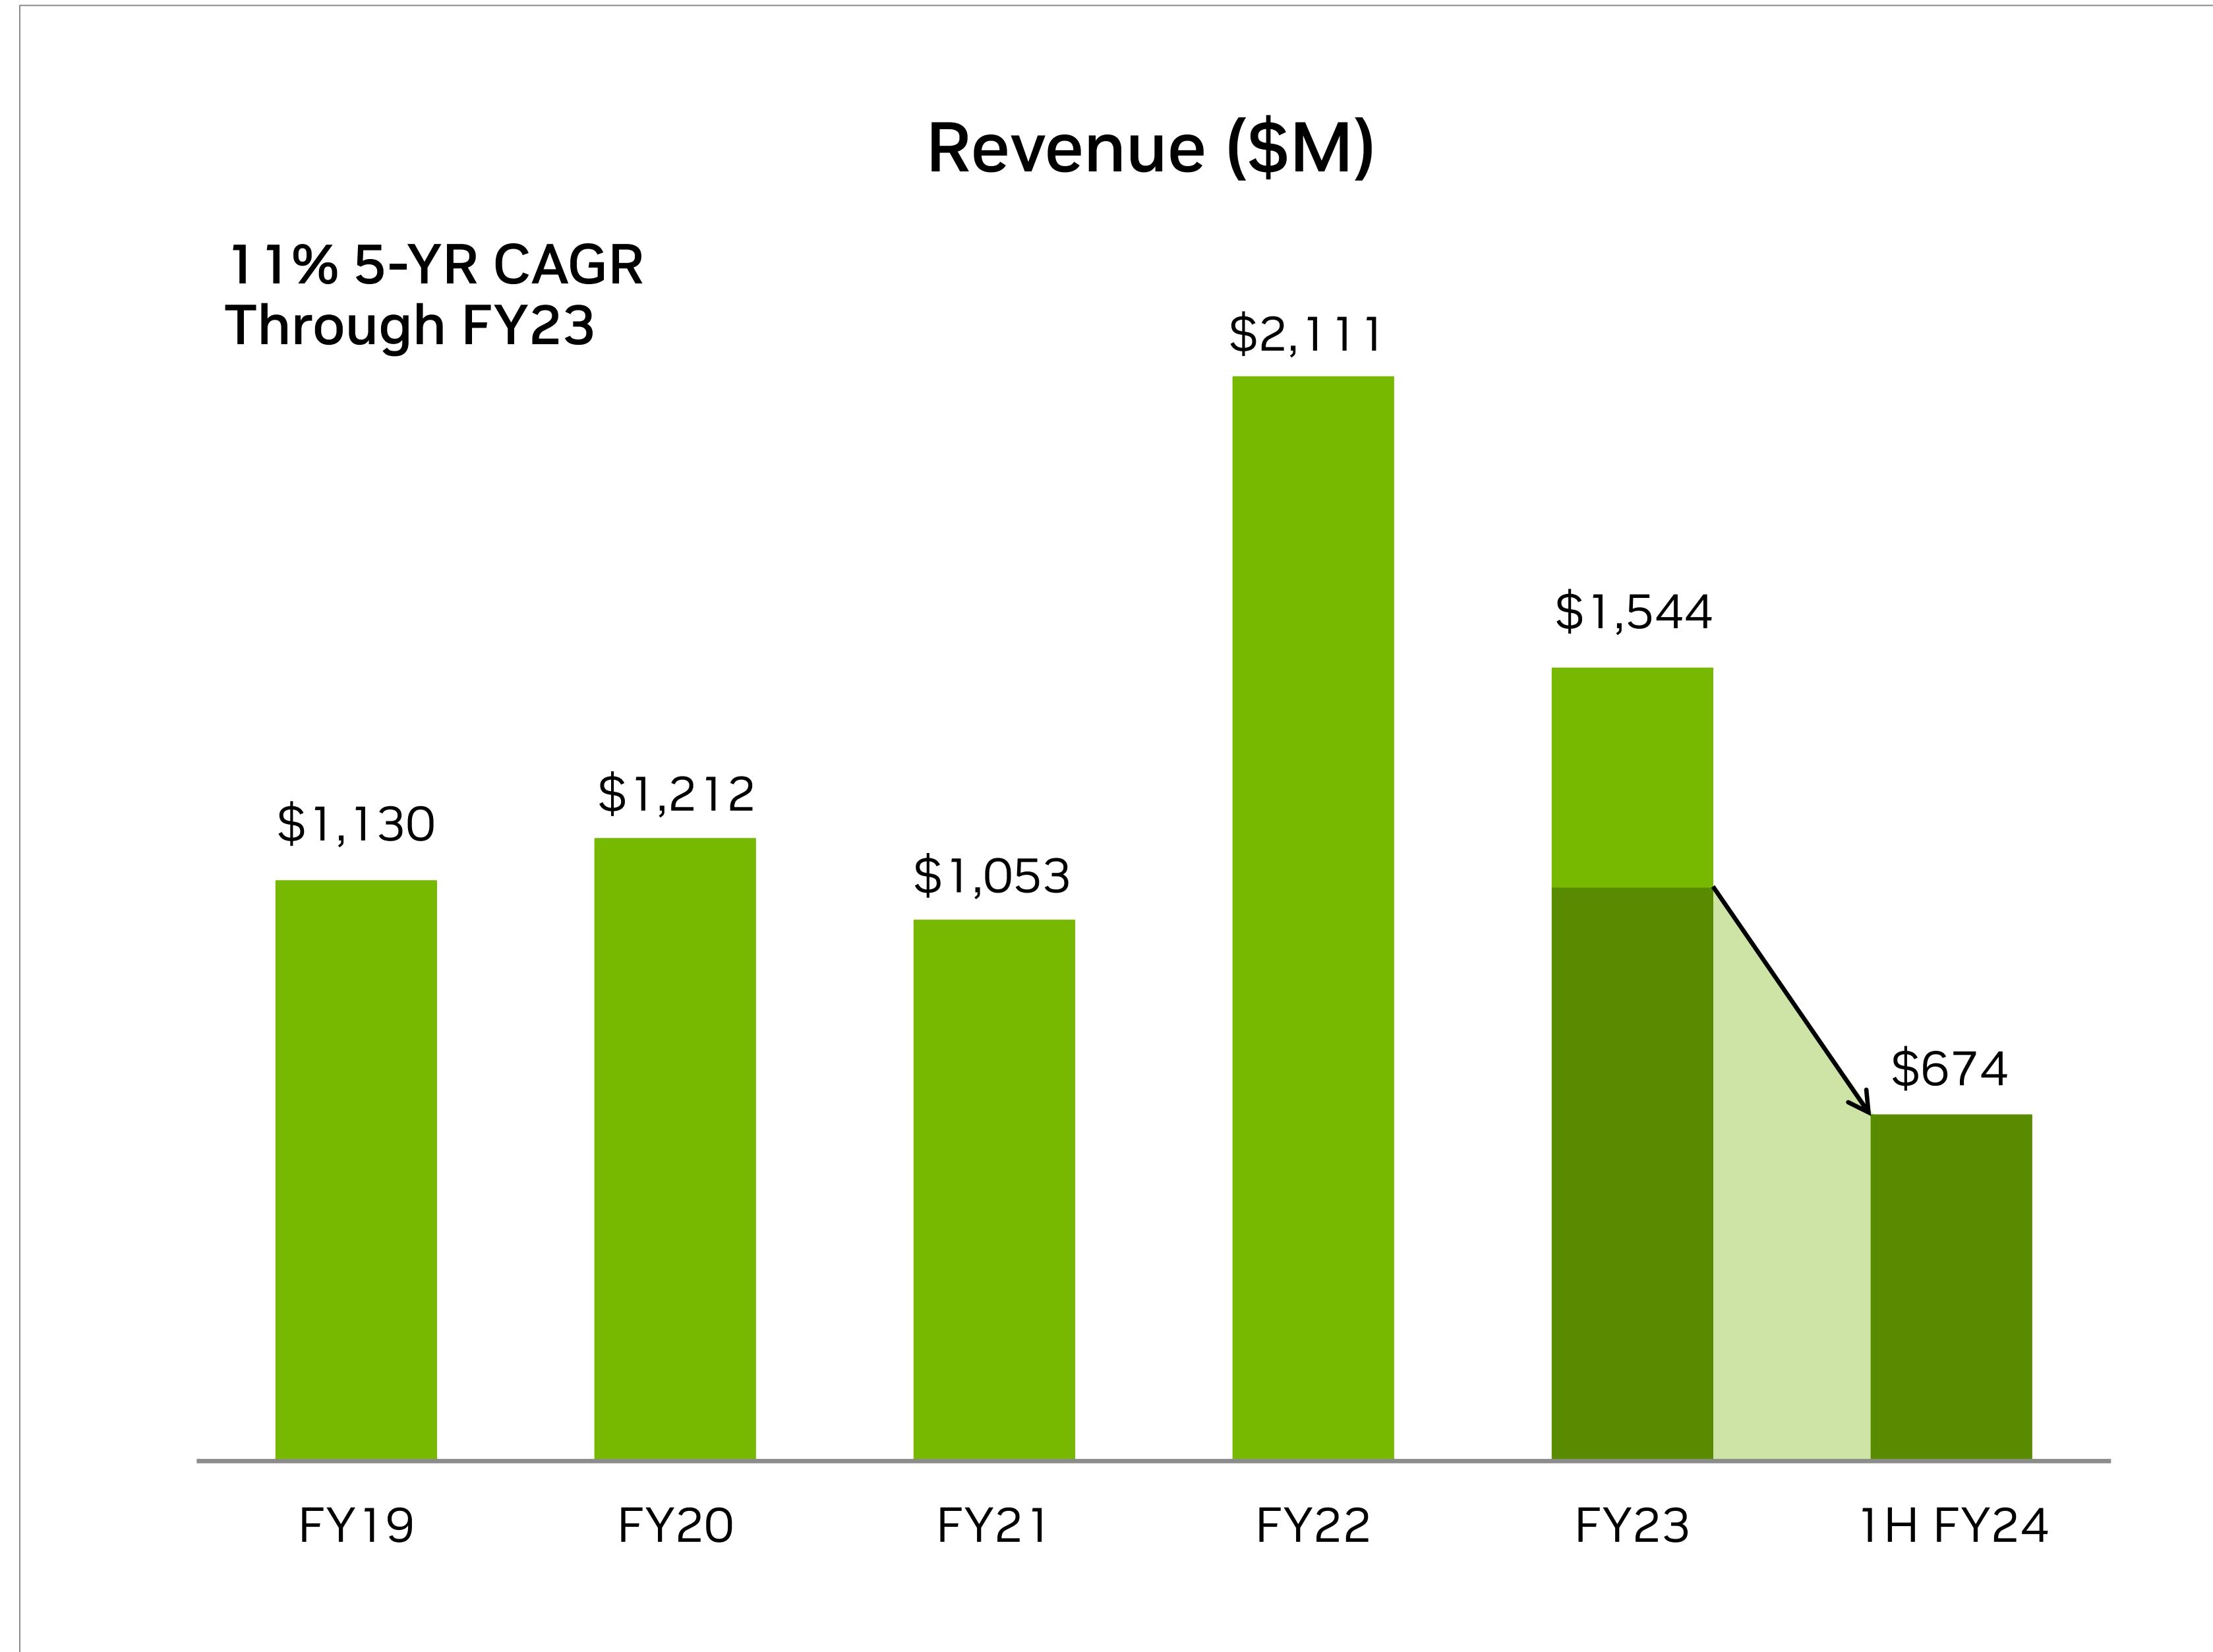
More Gamers, Richer Mix

Installed Base Needs Upgrade

Ada: 3X Turing Ramp at \$699+

Professional Visualization

Workstation graphics



Leader in Workstation Graphics

95%+ market share in graphics for workstations

45M Designers and Creators

Strong software ecosystem with over 100 RTX accelerated and supported applications

Growth Drivers

Ray Tracing and generative AI revolutionizing design and content creation

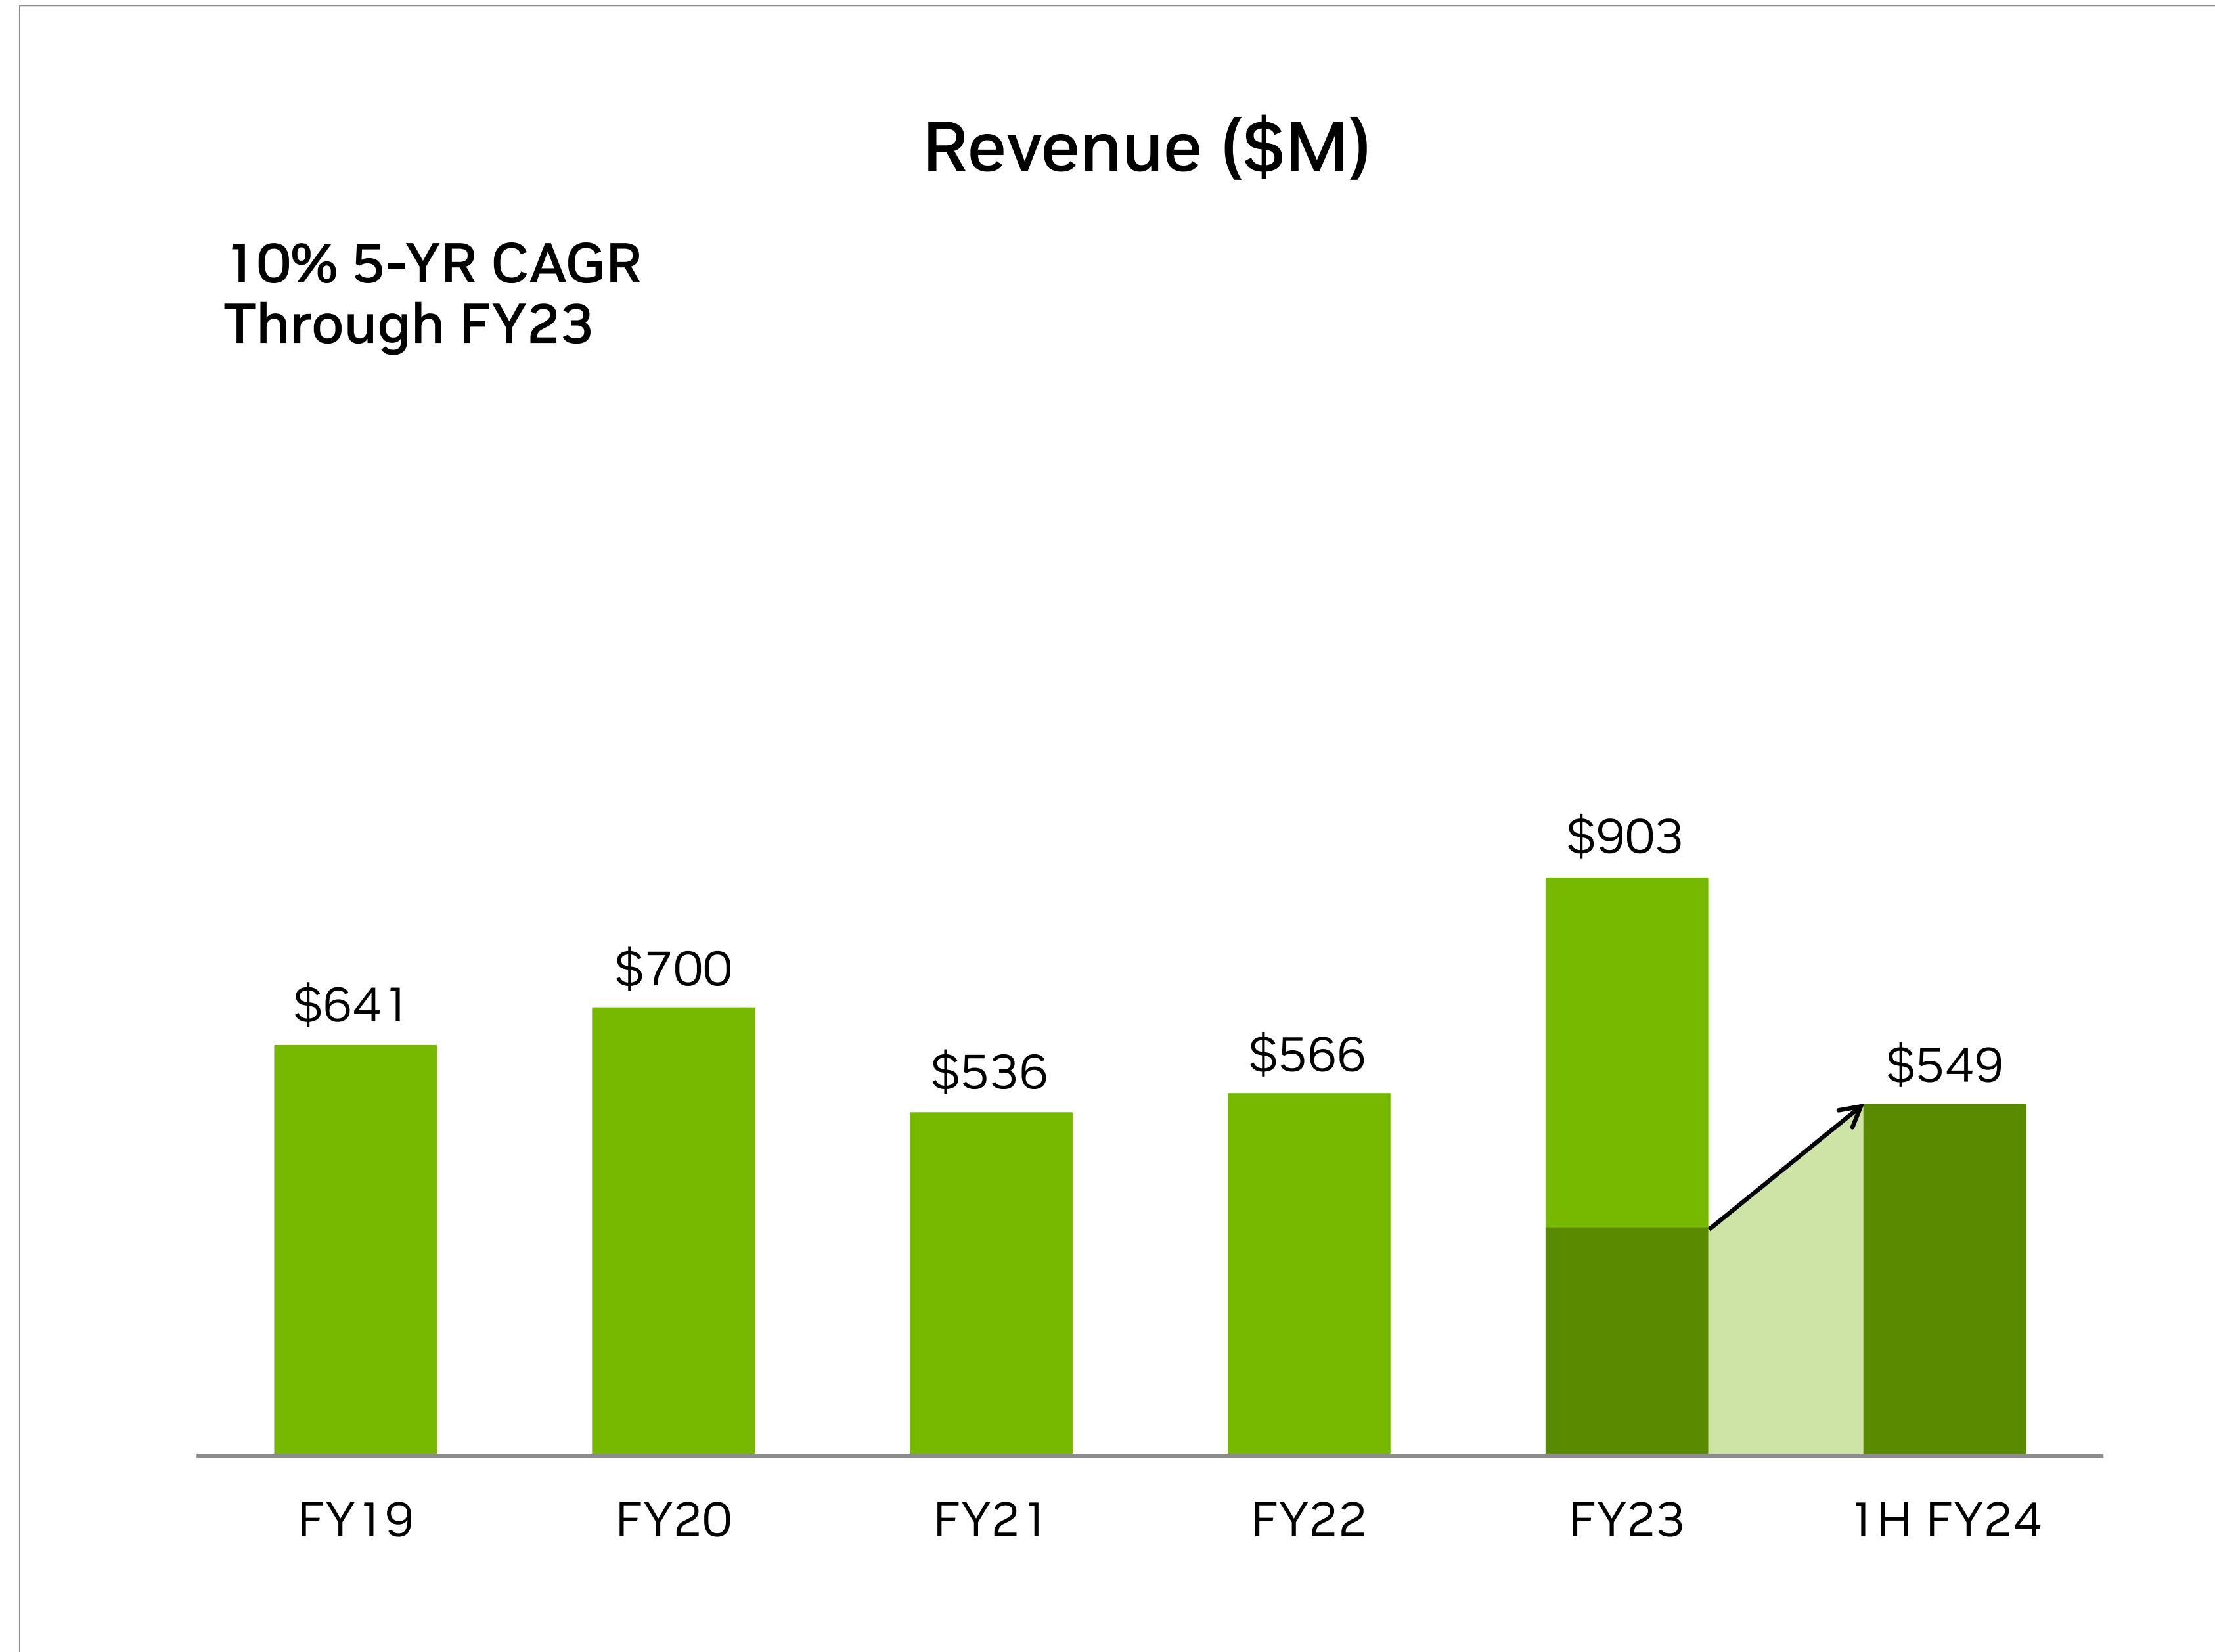
Expanding universe of designers and creators

Collaborative 3D design / Omniverse

Hybrid work environments

Automotive

Autonomous Vehicles (AV) & AI Cockpit



Leader in Autonomous Driving

Revenue growth primarily fueled by NVIDIA DRIVE, our AV and AI cockpit platform with full software stack

Inflection in FY23 driven by AV as DRIVE Orin SoC began to ramp

Next-generation DRIVE Thor SoC to ramp in FY26

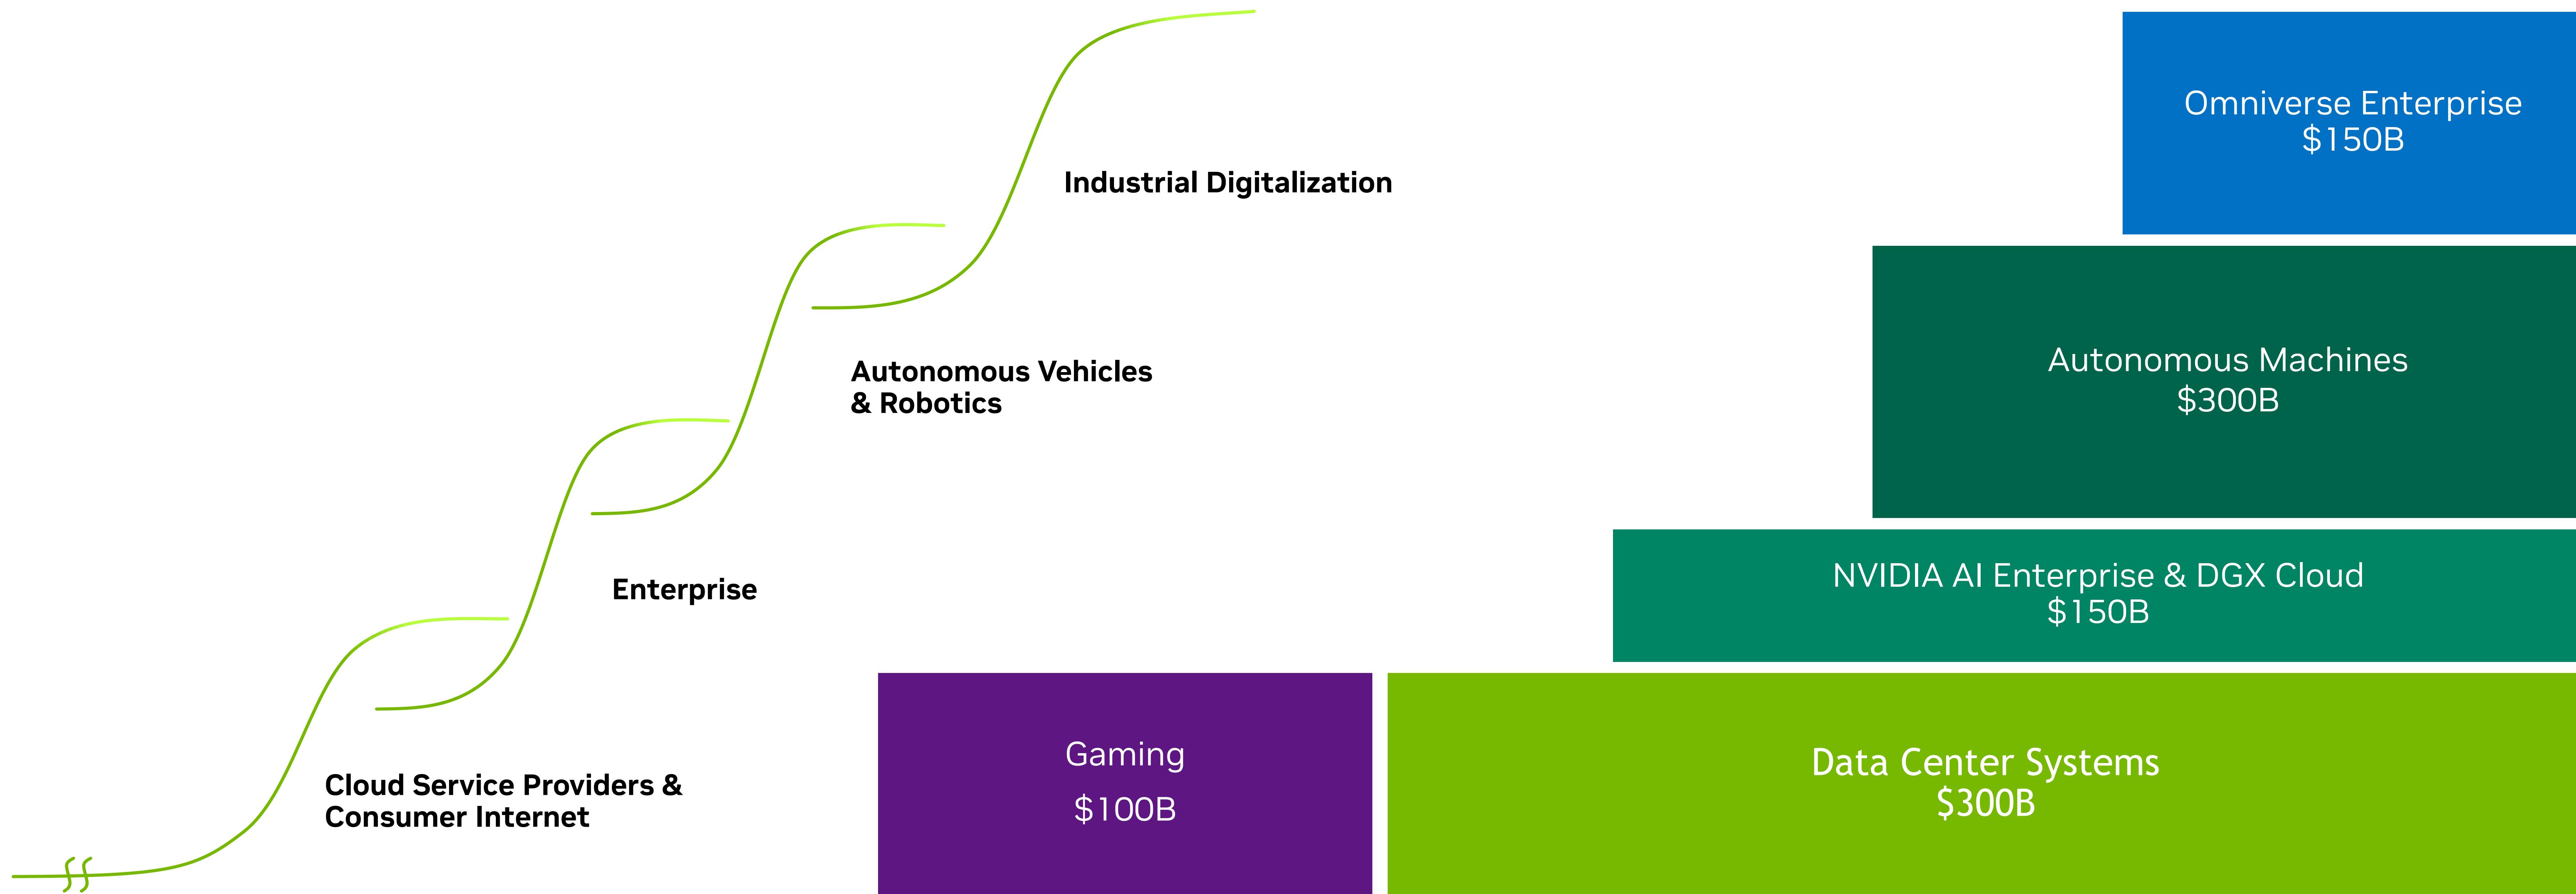
Growth Drivers

Adoption of centralized car computing and software-defined vehicle architectures

AV software and services:
Mercedes Benz
Jaguar Land Rover

\$14B Design Win Pipeline Through FY29

\$1 Trillion Long-Term Annual Market Opportunity



Summary

Gen AI is the tipping point for the new computing era

AI is the new software and Accelerated Computing the new hardware

Huge ROI from Gen AI – from new revenue or dramatically lower costs - is driving a powerful new investment cycle

NVIDIA's accelerated computing platform delivers unmatched performance and TCO savings

Strong revenue, operating profit, and cash flow growth

\$1T market opportunity

Reconciliation of Non-GAAP to GAAP Financial Measures

Reconciliation of Non-GAAP to GAAP Financial Measures

Gross Margin (\$ in Millions & Margin Percentage)	Non-GAAP	Acquisition-Related and Other Costs (A)	Stock-Based Compensation (B)	IP-Related Costs	GAAP
FY 2019	\$7,233	—	(27)	(35)	\$7,171
	61.7%	—	(0.2)	(0.3)	61.2%
FY 2020	\$6,821	—	(39)	(14)	\$6,768
	62.5%	—	(0.4)	(0.1)	62.0%
FY 2021	\$10,947	(425)	(88)	(38)	\$10,396
	65.6%	(2.6)	(0.5)	(0.2)	62.3%
FY 2022	\$17,969	(344)	(141)	(9)	\$17,475
	66.8%	(1.4)	(0.5)	—	64.9%
FY 2023	\$15,965	(455)	(138)	(16)	\$15,356
	59.2%	(1.7)	(0.5)	(0.1)	56.9%
1H FY 2023	\$8,636	(214)	(76)	—	\$8,346
	57.6%	(1.4)	(0.5)	—	55.7%
1H FY 2024	\$14,417	(239)	(58)	(10)	\$14,110
	69.7%	(1.2)	(0.3)	—	68.2%

A. Consists of amortization of intangible assets and inventory step-up

B. Stock-based compensation charge was allocated to cost of goods sold

Reconciliation of Non-GAAP to GAAP Financial Measures (contd.)

Operating Income and Margin (\$ in Millions & Margin Percentage)	Non-GAAP	Acquisition Termination Cost	Acquisition-Related and Other Costs (A)	Stock-Based Compensation (B)	IP-Related Costs	Other (C)	GAAP
FY 2019	\$4,407	—	(2)	(557)	(35)	(9)	\$3,804
	37.6%	—	—	(4.7)	(0.3)	(0.1)	32.5%
FY 2020	\$3,735	—	(31)	(844)	(14)	—	\$2,846
	34.2%	—	(0.3)	(7.7)	(0.1)	—	26.1%
FY 2021	\$6,803	—	(836)	(1,397)	(38)	—	\$4,532
	40.8%	—	(5.0)	(8.4)	(0.2)	—	27.2%
FY 2022	\$12,690	—	(636)	(2,004)	(9)	—	\$10,041
	47.2%	—	(2.5)	(7.4)	—	—	37.3%
FY 2023	\$9,040	(1,353)	(674)	(2,710)	(16)	(63)	\$4,224
	33.5%	(5.0)	(2.5)	(10.0)	(0.1)	(0.2)	15.7%
1H FY 2023	\$5,280	(1,353)	(324)	(1,227)	—	(9)	\$2,367
	35.2%	(9.0)	(2.2)	(8.2)	—	—	15.8%
1H FY 2024	\$10,828	—	(311)	(1,576)	(10)	10	\$8,941
	52.3%	—	(1.5)	(7.6)	—	—	43.2%

A. Consists of amortization of acquisition-related intangible assets, inventory step-up, transaction costs, compensation charges, and other costs

B. Stock-based compensation charge was allocated to cost of goods sold, research and development expense, and sales, general and administrative expense

C. Comprises of legal settlement costs, contributions, restructuring costs and assets held for sale related adjustments

Reconciliation of Non-GAAP to GAAP Financial Measures

(\$ in Millions)	Free Cash Flow	Purchases Related to Property and Equipment and Intangible Assets	Principal Payments on Property and Equipment and Intangible Assets	Net Cash Provided by Operating Activities
FY 2019	\$3,143	600	—	\$3,743
FY 2020	\$4,272	489	—	\$4,761
FY 2021	\$4,677	1,128	17	\$5,822
FY 2022	\$8,049	976	83	\$9,108
FY 2023	\$3,750	1,833	58	\$5,641
1H FY 2023	\$2,171	794	36	\$3,001
1H FY 2024	\$8,691	537	31	\$9,259

